

## Assembly Contents

- [Registers](#)
- [Instruction format](#)
- [Aside: Directives](#)
- [Address modes](#)
- [Address computations](#)
- [%rip-relative addressing](#)
- [Arithmetic instructions](#)
- [Calling convention](#)
- [Argument passing and stack frames](#)
- [Stack](#)
- [Return address and entry and exit sequence](#)
- [Callee-saved registers and caller-saved registers](#)
- [Base pointer \(frame pointer\)](#)
- [Stack size and red zone](#)
- [Branches](#)
- [Aside: C++ data structures](#)
- [Compiler optimizations](#)

## Registers

Registers are the fastest kind of memory available in the machine. x86-64 has 14 general-purpose registers and several special-purpose registers. This table gives all the basic registers, with special-purpose registers highlighted in yellow. You'll notice different naming conventions, a side effect of the long history of the x86 architecture (the 8086 was first released in 1978).

| Full register<br>(bits 0-63) | 32-bit<br>(bits 0-31) | 16-bit<br>(bits 0-15) | 8-bit low<br>(bits 0-7) | 8-bit high<br>(bits 8-15) | Use in   | <a href="#">calling convention</a> | <a href="#">Callee-saved?</a> |
|------------------------------|-----------------------|-----------------------|-------------------------|---------------------------|--|------------------------------------|-------------------------------|
| General-purpose registers:   |                       |                       |                         |                           |  |                                    |                               |
| %rax                         | %eax                  | %ax                   | %al                     | %ah                       | Return value (accumulator)                               | No                                 |                               |
| %rbx                         | %ebx                  | %bx                   | %bl                     | %bh                       | -  | Yes                                |                               |
| %rcx                         | %ecx                  | %cx                   | %cl                     | %ch                       | 4th function argument                                    | No                                 |                               |
| %rdx                         | %edx                  | %dx                   | %dl                     | %dh                       | 3rd function argument                                    | No                                 |                               |
| %rsi                         | %esi                  | %si                   | %sil                    | -                         | 2nd function argument                                    | No                                 |                               |
| %rdi                         | %edi                  | %di                   | %dil                    | -                         | 1st function argument                                    | No                                 |                               |
| %r8                          | %r8d                  | %r8w                  | %r8b                    | -                         | 5th function argument                                    | No                                 |                               |
| %r9                          | %r9d                  | %r9w                  | %r9b                    | -                         | 6th function argument                                    | No                                 |                               |
| %r10                         | %r10d                 | %r10w                 | %r10b                   | -                         | -  | No                                 |                               |
| %r11                         | %r11d                 | %r11w                 | %r11b                   | -                         | -  | No                                 |                               |
| %r12                         | %r12d                 | %r12w                 | %r12b                   | -                         | -  | Yes                                |                               |
| %r13                         | %r13d                 | %r13w                 | %r13b                   | -                         | -  | Yes                                |                               |
| %r14                         | %r14d                 | %r14w                 | %r14b                   | -                         | -  | Yes                                |                               |
| %r15                         | %r15d                 | %r15w                 | %r15b                   | -                         | -  | Yes                                |                               |
| Special-purpose registers:   |                       |                       |                         |                           |  |                                    |                               |
| %rsp                         | %esp                  | %sp                   | %spl                    | -                         | Stack pointer  | Yes                                |                               |
| %rbp                         | %ebp                  | %bp                   | %bpl                    | -                         | Base pointer<br>(general-purpose in some compiler modes) | Yes                                |                               |

| Full register<br>(bits 0-63) | 32-bit<br>(bits 0-31) | 16-bit<br>(bits 0-15) | 8-bit low<br>(bits 0-7) | 8-bit high<br>(bits 8-15) | Use in  | <a href="#">calling convention</a> | <a href="#">Callee-saved?</a> |
|------------------------------|-----------------------|-----------------------|-------------------------|---------------------------|---|------------------------------------|-------------------------------|
| %rip                         | %eip                  | %ip                   | -                       | -                         | Instruction pointer<br>(Program counter; called %pc in GDB) | *                                  |                               |
| %rflags                      | %eflags               | %flags                | -                       | -                         | Flags and condition codes                                   | No                                 |                               |

Note that unlike *primary* memory (which is what we think of when we discuss memory in a C/C++ program), registers have no addresses! There is no address value that, if cast to a pointer and dereferenced, would return the contents of the **%rax** register. Registers live in a separate world from the memory whose contents are partially prescribed by the C abstract machine.

The %rbp register has a special purpose: it points to the bottom of the current function's stack frame, and local variables are often accessed relative to its value. However, when optimization is on, the compiler may determine that all local variables can be stored in registers. This frees up %rbp for use as another general-purpose register.

The relationship between different register bit widths is a little weird.

1. Modifying a 32-bit register name sets the *upper* 32 bits of the register to zero. Thus, after **movl \$-1, %eax**, the **%rax** register has value 0x0000'0000'FFFF'FFFF. The same is true after **movq \$-1, %rax; addl \$0, %eax!** (The **movq** sets **%rax** to 0xFFFF'FFFF'FFFF'FFFF; the **addl** sets its upper 32 bits to zero.)
2. Modifying a 16- or 8-bit register name *leaves all other bits of the register unchanged*.

There are special instructions for loading signed and unsigned 8-, 16-, and 32-bit quantities into registers, recognizable by instruction suffixes. For instance, **movzbl** moves an 8-bit quantity (a **byte**) into a 32-bit register (a longword) with zero extension; **movslq** moves a 32-bit quantity (longword) into a 64-bit register (**quadword**) with sign extension. There's no need for **movz1q** (why?).

## Instruction format

The basic kinds of assembly instructions are:

1. **Arithmetic.** These instructions perform computation on values, typically values stored in registers. Most have zero or one *source operands* and one *source/destination operand*. The source operand is listed first in the *instruction*, but the source/destination operand comes first in the *computation* (this matters for non-commutative operators like subtraction). For example, the instruction **addq %rax, %rbx** performs the computation **%rbx := %rbx + %rax**.
2. **Data movement.** These instructions move data between registers and memory. Almost all have one source operand and one destination operand; the source operand comes first.
3. **Control flow.** Normally the CPU executes instructions in sequence. Control flow instructions change the instruction pointer in other ways. There are unconditional branches (the instruction pointer is set to a new value), conditional branches (the instruction pointer is set to a new value if a condition is true), and function call and return instructions.

(We use the "AT&T syntax" for x86-64 assembly. For the "Intel syntax," which you can find in online documentation from Intel, see the Aside in CS:APP3e §3.3, p177, or [Wikipedia](#), or other online resources. AT&T syntax is distinguished by several features, but especially by the use of percent signs for registers. Sadly, the Intel syntax puts destination registers *before* source registers.)

Some instructions appear to combine arithmetic and data movement. For example, given the C code **int\* ip; ... ++(\*ip);** the compiler might generate **incl (%rax)** rather than **movl (%rax), %ebx; incl %ebx; movl %ebx, (%rax)**. However, the processor actually divides

these complex instructions into tiny, simpler, invisible instructions called [microcode](#), because the simpler instructions can be made to execute faster. The complex `incl` instruction actually runs in three phases: data movement, then arithmetic, then data movement. This matters when we introduce parallelism.

## Directives

Assembly generated by a compiler contains instructions as well as *labels* and *directives*. Labels look like `labelname:` or `labelnumber::`; directives look like `.directivename arguments`. Labels are markers in the generated assembly, used to compute addresses. We usually see them used in control flow instructions, as in `jmp L3` ("jump to L3"). Directives are instructions to the assembler; for instance, the `.globl L` instruction says "label `L` is globally visible in the executable", `.align` sets the alignment of the following data, `.long` puts a number in the output, and `.text` and `.data` define the current segment.

We also frequently look at assembly that is *disassembled* from executable instructions by GDB, `objdump -d`, or `objdump -S`. This output looks different from compiler-generated assembly: in disassembled instructions, there are no intermediate labels or directives. This is because the labels and directives disappear during the process of generating executable instructions.

For instance, here is some compiler-generated assembly:

```
.globl _Z1fiii
.type _Z1fiii, @function
_Z1fiii:
.LFB0:
    cmpl %edx, %esi
    je .L3
    movl %esi, %eax
    ret
.L3:
    movl %edi, %eax
    ret
.LFE0:
.size _Z1fiii, .-_Z1fiii
```

And a disassembly of the same function, from an object file:

```
0000000000000000 <_Z1fiii>:
0: 39 d6          cmp    %edx,%esi
2: 74 03          je     7 <_Z1fiii+0x7>
4: 89 f0          mov    %esi,%eax
6: c3             retq
7: 89 f8          mov    %edi,%eax
9: c3             retq
```

Everything but the instructions is removed, and the helpful `.L3` label has been replaced with an actual address. The function appears to be located at address 0. This is just a placeholder; the final address is assigned by the linking process, when a final executable is created.

Finally, here is some disassembly from an executable:

```
0000000000400517 <_Z1fiii>:
400517: 39 d6          cmp    %edx,%esi
400519: 74 03          je     40051e <_Z1fiii+0x7>
40051b: 89 f0          mov    %esi,%eax
40051d: c3            retq
40051e: 89 f8          mov    %edi,%eax
400520: c3            retq
```

The instructions are the same, but the addresses are different. (Other compiler flags would generate different addresses.)

## Address modes

Most instruction operands use the following syntax for values. (See also CS:APP3e Figure 3.3 in §3.4.1, p181.)

| Type      | Example syntax                   | Value used  |
|-----------|----------------------------------|---|
| Register  | <code>%rbp</code>                | Contents of <code>%rbp</code>   |
| Immediate | <code>\$0x4</code>               | 0x4   |
| Memory    | <code>0x4</code>                 | Value stored at address 0x4   |
|           | <code>symbol_name</code>         | Value stored in global <code>symbol_name</code><br>(the compiler resolves the symbol name to an address when creating the executable) |
|           | <code>symbol_name(%rip)</code>   | <a href="#">%rip-relative addressing</a> for global   |
|           | <code>symbol_name+4(%rip)</code> | Simple arithmetic on symbols are allowed<br>(the compiler resolves the arithmetic when creating the executable)                       |
|           | <code>(%rax)</code>              | Value stored at address in <code>%rax</code>  |
|           | <code>0x4(%rax)</code>           | Value stored at address <code>%rax + 4</code>   |
|           | <code>(%rax,%rbx)</code>         | Value stored at address <code>%rax + %rbx</code>  |
|           | <code>(%rax,%rbx,4)</code>       | Value stored at address <code>%rax + %rbx*4</code>  |
|           | <code>0x18(%rax,%rbx,4)</code>   | Value stored at address <code>%rax + 0x18 + %rbx*4</code>   |

The full form of a memory operand is `offset(base,index,scale)`, which refers to the address `offset + base + index*scale`. In `0x18(%rax,%rbx,4)`, `%rax` is the base, `0x18` the offset, `%rbx` the index, and `4` the scale. The offset (if used) must be a constant and the base and index (if used) must be registers; the scale must be either 1, 2, 4, or 8. The default offset, base, and index are 0, and the default scale is 1.

`symbol_names` are found only in compiler-generated assembly; disassembly uses raw addresses (`0x601030`) or `%rip`-relative offsets (`0x200bf2(%rip)`).

Jumps and function call instructions use different syntax `Ⓜ: *`, rather than `()`, represents indirection.

| Type      | Example syntax                               | Address used  |
|-----------|--|---|
| Register  | <code>*%rax</code>                           | Contents of <code>%rax</code>                             |
| Immediate | <code>.L3</code>                             | Address of <code>.L3</code> (compiler-generated assembly) |
|           | <code>400410</code> or <code>0x400410</code> | Given address   |
| Memory    | <code>*0x200b96(%rip)</code>                 | Value stored at address <code>%rip + 0x200b96</code>      |
|           | <code>*(%r12,%rbp,8)</code>                  | Other address modes accepted                              |

## Address arithmetic

The `base(offset,index,scale)` form compactly expresses many array-style address computations. It's typically used with a `mov`-type instruction to dereference memory. However, the compiler can use that form to compute addresses, thanks to the `lea` (Load Effective Address) instruction.

For instance, in `movl 0x18(%rax,%rbx,4), %ecx`, the address `%rax + 0x18 + %rbx*4` is computed, then immediately dereferenced: the 4-byte value located there is loaded into `%ecx`. In `leaq 0x18(%rax,%rbx,4), %rcx`, the same address is computed, but it is *not* dereferenced. Instead, the *computed address* is moved into register `%rcx`.

Thanks to `lea`, the compiler will also prefer the `base(offset,index,scale)` form over `add` and `mov` for certain arithmetic computations on integers. For example, this instruction: `leaq (%rax,%rbx,2), %rcx` performs the function `%rcx := %rax + 2 * %rbx`, but in one instruction, rather than three (`movq %rax, %rcx; addq %rbx, %rcx; addq %rbx, %rcx`).

## %rip-relative addressing

x86-64 code often refers to globals using **%rip-relative** addressing: a global variable named `a` is referenced as `a(%rip)` rather than `a`.

This style of reference supports [position-independent code](#) (PIC), a security feature. It specifically supports *position-independent executables* (PIEs), which are programs that work independently of where their code is loaded into memory.

To run a conventional program, the operating system loads the program’s instructions into memory *at a fixed address that’s the same every time*, then starts executing the program at its first instruction. This works great, and runs the program in a predictable execution environment (the addresses of functions and global variables are the same every time). Unfortunately, the very predictability of this environment makes the program easier to attack.

In a position-independent executable, the operating system loads the program at *varying* locations: every time it runs, the program’s functions and global variables have different addresses. This makes the program harder to attack (though [not impossible](#)). Program startup performance matters, so the operating system doesn’t recompile the program with different addresses each time. Instead, the compiler does most of the work in advance by using *relative addressing*.

When the operating system loads a PIE, it picks a starting point and loads all instructions and globals relative to that starting point. The PIE’s instructions never refer to global variables using direct addressing: you’ll never see `movl global_int, %eax`. Globals are referenced *relatively* instead, using deltas relative to the next `%rip`: `movl global_int(%rip), %eax`. These relative addresses work great independent of starting point! For instance, consider an instruction located at (starting-point + 0x80) that loads a variable `g` located at (starting-point + 0x1000) into `%rax`. In a non-PIE, the instruction might be written `movq 0x400080, %rax` (in compiler output, `movq g, %rax`); but this relies on `g` having a fixed address. In a PIE, the instruction might be written `movq 0xf79(%rip), %rax` (in compiler output, `movq g(%rip), %rax`), which works out beautifully no matter the starting point.

| At starting point... | The <code>mov</code> instruction is at... | The next instruction is at... | And <code>g</code> is at... | So the delta ( <code>g - next %rip</code> ) is... |
|----------------------|---|-------------------------------|-----------------------------|---|
| 0x400000             | 0x400080                                  | 0x400087                      | 0x401000                    | 0xF79   |
| 0x404000             | 0x404080                                  | 0x404087                      | 0x405000                    | 0xF79   |
| 0x4003F0             | 0x400470                                  | 0x400477                      | 0x4013F0                    | 0xF79   |

Arithmetic instructions

The operations of many x86-64 arithmetic instructions are easy enough to guess from their names. Then there are some arithmetic instructions, particularly those associated with [Streaming SIMD Extensions](#) and its follow-ons, that are hard to guess (`phminposuw?`). The basic arithmetic instructions on 64-bit quantities (“quadwords”) are:

| Instruction  | Operation             | Type   | Expansion  |
|--|-----------------------|--------|--|
| <code>addq SRC, DST</code>                                 | Addition              | Normal | <code>DST := DST + SRC</code>                          |
| <code>subq SRC, DST</code>                                 | Subtraction           | Normal | <code>DST := DST - SRC</code>                          |
| <code>incq DST</code>                                      | Increment             | Normal | <code>DST := DST + 1</code>                            |
| <code>decq DST</code>                                      | Decrement             | Normal | <code>DST := DST - 1</code>                            |
| <code>imulq SRC, DST</code>                                | Signed multiplication | Normal | <code>DST := DST * SRC</code>                          |
| <code>negq DST</code>                                      | Negation              | Normal | <code>DST := -DST (DST := ~DST + 1)</code>             |
| <code>andq SRC, DST</code>                                 | Bitwise and           | Normal | <code>DST := DST &amp; SRC</code>                      |
| <code>orq SRC, DST</code>                                  | Bitwise or            | Normal | <code>DST := DST   SRC</code>                          |
| <code>xorq SRC, DST</code>                                 | Bitwise exclusive or  | Normal | <code>DST := DST ^ SRC</code>                          |
| <code>notq DST</code>                                      | Complement            | Normal | <code>DST := ~DST</code>                               |
| <code>sal SRC, DST</code><br>( <code>shl SRC, DST</code> ) | Left shift            | Normal | <code>DST := DST &lt;&lt; SRC</code>                   |
| <code>sar SRC, DST</code>                                  | Signed right shift    | Normal | <code>DST:= DST&gt;&gt;SRC,shifting in sign bit</code> |
| <code>shr SRC, DST</code>                                  | Unsigned right shift  | Normal | <code>DST:= DST&gt;&gt;SRC,shifting in zeros</code>    |

| Instruction                 | Operation             | Type       | Expansion  |
|-----------------------------|-----------------------|------------|--|
| <code>cmpq SRC, DST</code>  | Subtraction for flags | Flags-only | <code>DST - SRC</code> ; <a href="#">see below</a>     |
| <code>testq SRC, DST</code> | Bitwise and for flags | Flags-only | <code>DST &amp; SRC</code> ; <a href="#">see below</a> |

There are also compact multiplication and division instructions that modify multiple registers at once and take fixed registers for some arguments. These instructions treat the combination of `%rax` and `%rdx` as a single 128-bit value where the most significant bits (bits 64–127) of the value are stored in `%rdx` and the least significant bits (0–63) are stored in `%rax`. The division instructions compute both a quotient and a remainder. (In the below, `TMP` is a 128-bit number.)

| Instruction            | Operation               | Type    | Expansion   |
|------------------------|-------------------------|---------|---|
| <code>imulq SRC</code> | Signed multiplication   | Mul/div | <code>TMP := %rax * SRC; %rdx := TMP&gt;&gt;64; %rax := TMP</code>                |
| <code>mulq SRC</code>  | Unsigned multiplication | Mul/div | <code>TMP := %rax * SRC; %rdx := TMP&gt;&gt;64; %rax := TMP</code>                |
| <code>idivq SRC</code> | Signed division         | Mul/div | <code>TMP := (%rdx&lt;&lt;64)   %rax; %rax := TMP / SRC; %rdx := TMP % SRC</code> |
| <code>divq SRC</code>  | Unsigned division       | Mul/div | <code>TMP := (%rdx&lt;&lt;64)   %rax; %rax := TMP / SRC; %rdx := TMP % SRC</code> |

Calling convention

A [calling convention](#) governs how functions on a particular architecture and operating system interact. This includes rules about includes how function arguments are placed, where return values go, what registers functions may use, how they may allocate local variables, and so forth. Calling conventions ensure that functions compiled by different compilers can interoperate, and they ensure that operating systems can run code from different programming languages and compilers. Some aspects of a calling convention are derived from the instruction set itself, but some are conventional, meaning decided upon by people (for instance, at a convention). Calling conventions constrain both *callers* and *callees*. A caller is a function that calls another function; a callee is a function that was called. The currently-executing function is a callee, but not a caller.

For concreteness, we learn the [x86-64 calling conventions for Linux](#). These conventions are shared by many OSes, including MacOS (but not Windows), and are officially called the “System V AMD64 ABI.”

The official specification: [AMD64 ABI](#)

Argument passing and stack frames

One set of calling convention rules governs how function arguments and return values are passed. On x86-64 Linux, the first six function arguments are passed in registers `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8`, and `%r9`, respectively. The seventh and subsequent arguments are passed on the stack, about which more below. The return value is passed in register `%rax`.

The full rules more complex than this. You can read them in [the AMD64 ABI](#), section 3.2.3, but they’re quite detailed. Some highlights:

- 1. A structure argument that fits in a single machine word (64 bits/8 bytes) is passed in a single register.  
Example: `struct small { char a1, a2; }`
- 2. A structure that fits in two to four machine words (16–32 bytes) is passed in sequential registers, as if it were multiple arguments.  
Example: `struct medium { long a1, a2; }`
- 3. A structure that’s larger than four machine words is always passed on the stack.  
Example: `struct large { long a, b, c, d, e, f, g; }`

- Floating point arguments are generally passed in special registers, the “SSE registers,” that we don’t discuss further.
- Return values of up to one machine word are passed in `%rax`. Many return values that fit in two machine words—for instance, a pair of `longs`—are passed in `%rax` (for the first 8 bytes) and `%rdx` (for the second 8 bytes). For return values that fit only in three or more machine words, the *caller* reserves space for the return value, and passes the *address* of that space as the first argument of the function. The callee will fill in that space when it returns.

Writing small programs to demonstrate these rules is a pleasant exercise; for example:

```
struct small { char a1, a2; };
int f(small s) {
    return s.a1 + 2 * s.a2;
}
```

compiles to:

```
movl %edi, %eax          # copy argument to %eax
movsbl %dil, %edi        # %edi := sign-extension of lowest byte of argument
                          # (s.a1)
movsbl %ah, %eax         # %eax := sign-extension of 2nd byte of argument
                          # (s.a2)
movsbl %al, %eax
leal (%rdi,%rax,2), %eax  # %eax := %edi + 2 * %eax
ret
```

## Stack

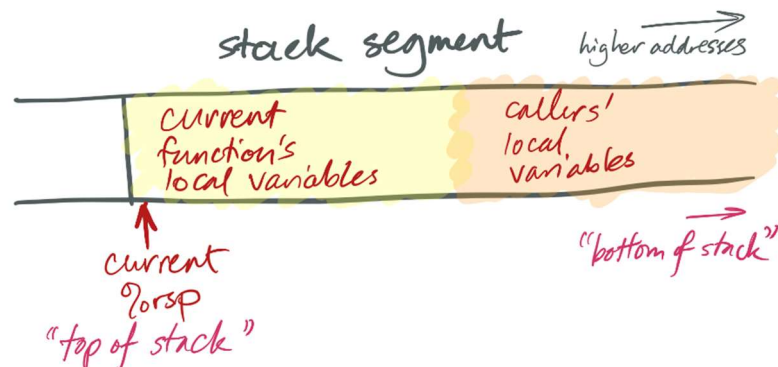
[Recall](#) that the stack is a segment of memory used to store objects with automatic lifetime.

Typical stack addresses on x86-64 look like `0x7ffd'9f10'4f58`—that is, close to  $2^{47}$ .

The stack is named after a data structure, which was sort of named after pancakes. Stack data structures support at least three operations: **push** adds a new element to the “top” of the stack; **pop** removes the top element, showing whatever was underneath; and **top** accesses the top element. Note what’s missing: the data structure does not allow access to elements other than the top. (Which is sort of how stacks of pancakes work.) This restriction can speed up stack implementations.

Like a stack data structure, the stack memory segment is only accessed from the top. The currently running function accesses *its* local variables; the function’s caller, grand-caller, great-grand-caller, and so forth are dormant until the currently running function returns.

x86-64 stacks look like this:



The x86-64 `%rsp` register is a special-purpose register that defines the current “stack pointer.” This holds the address of the current top of the stack. On x86-64, as on many architectures, stacks grow *down*: a “push” operation adds space for more automatic-lifetime objects by moving the stack pointer left, to a numerically-smaller address, and a “pop” operation recycles

space by moving the stack pointer right, to a numerically-larger address. This means that, considered numerically, the “top” of the stack has a smaller address than the “bottom.” This is built in to the architecture by the operation of instructions like `pushq`, `popq`, `call`, and `ret`. A `push` instruction pushes a value onto the stack. This both modifies the stack pointer (making it smaller) and modifies the stack segment (by moving data there). For instance, the instruction `pushq X` means:

```
subq $8, %rsp
movq X, (%rsp)
```

And `popq X` undoes the effect of `pushq X`. It means:

```
movq (%rsp), X
addq $8, %rsp
```

`X` can be a register or a memory reference.

The portion of the stack reserved for a function is called that function’s **stack frame**. Stack frames are aligned: x86-64 requires that each stack frame be a multiple of 16 bytes, and when a `callq` instruction begins execution, the `%rsp` register must be 16-byte aligned. This means that every function’s entry `%rsp` address will be 8 bytes off a multiple of 16.

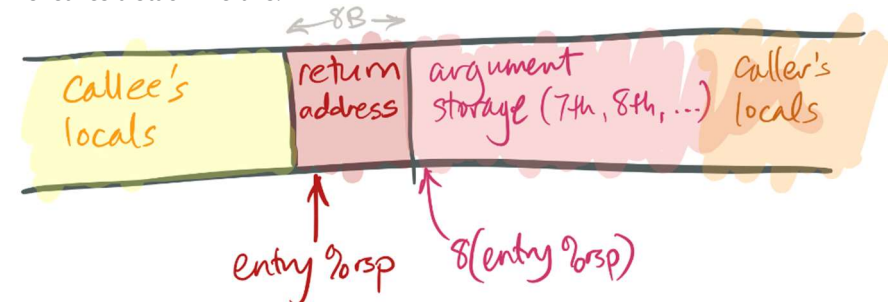
## Return address and entry and exit sequence

The steps required to call a function are sometimes called the *entry sequence* and the steps required to return are called the *exit sequence*. Both caller and callee have responsibilities in each sequence.

To prepare for a function call, the caller performs the following tasks in its entry sequence.

- The caller stores the first six arguments in the corresponding registers.
- If the callee takes more than six arguments, or if some of its arguments are large, the caller must store the surplus arguments on its stack frame. It stores these in increasing order, so that the 7th argument has a smaller address than the 8th argument, and so forth. The 7th argument must be stored at `(%rsp)` (that is, the top of the stack) when the caller executes its `callq` instruction.
- The caller saves any caller-saved registers (see below).
- The caller executes `callq FUNCTION`. This has an effect like `pushq $NEXT_INSTRUCTION; jmp FUNCTION` (or, equivalently, `subq $8, %rsp; movq $NEXT_INSTRUCTION, (%rsp); jmp FUNCTION`), where `NEXT_INSTRUCTION` is the address of the instruction immediately following `callq`.

This leaves a stack like this:



To return from a function:

- The callee places its return value in `%rax`.
- The callee restores the stack pointer to its value at entry (“entry `%rsp`”), if necessary.
- The callee executes the `retq` instruction. This has an effect like `popq %rip`, which removes the return address from the stack and jumps to that address.
- The caller then cleans up any space it prepared for arguments and restores caller-saved registers if necessary.



Particularly simple callees don't need to do much more than return, but most callees will perform more tasks, such as allocating space for local variables and calling functions themselves.

### Callee-saved registers and caller-saved registers

The calling convention gives callers and callees certain guarantees and responsibilities about the values of registers across function calls. Function implementations may expect these guarantees to hold, and must work to fulfill their responsibilities.

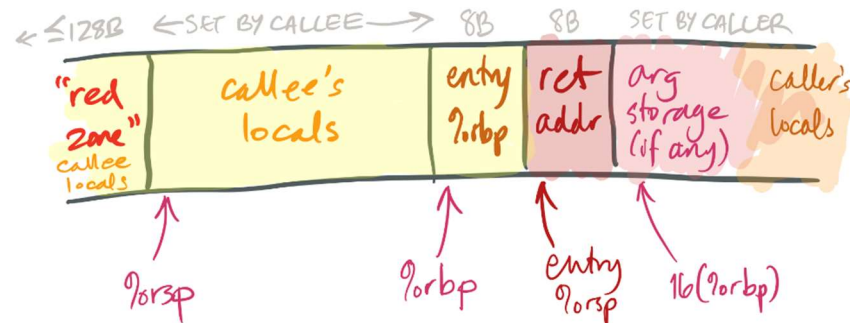
The most important responsibility is that certain registers' values *must be preserved across function calls*. A callee may use these registers, but if it changes them, it must restore them to their original values before returning. These registers are called **callee-saved registers**. All other registers are **caller-saved**.

Callers can simply use callee-saved registers across function calls; in this sense they behave like C++ local variables. Caller-saved registers behave differently: if a caller wants to preserve the value of a caller-saved register across a function call, the caller must explicitly save it before the `callq` and restore it when the function resumes.

On x86-64 Linux, `%rbp`, `%rbx`, `%r12`, `%r13`, `%r14`, and `%r15` are callee-saved, as (sort of) are `%rsp` and `%rip`. The other registers are caller-saved.

### Base pointer (frame pointer)

The `%rbp` register is called the *base pointer* (and sometimes the *frame pointer*). For simple functions, an optimizing compiler generally treats this like any other callee-saved general-purpose register. However, for more complex functions, `%rbp` is used in a specific pattern that facilitates debugging. It works like this:



1. The first instruction executed on function entry is `pushq %rbp`. This saves the caller's value for `%rbp` into the callee's stack. (Since `%rbp` is callee-saved, the callee must save it.)
2. The second instruction is `movq %rsp, %rbp`. This saves the current stack pointer in `%rbp` (so `%rbp = entry %rsp - 8`). This adjusted value of `%rbp` is the callee's "frame pointer." The callee will not change this value until it returns. The frame pointer provides a stable reference point for local variables and caller arguments. (Complex functions may need a stable reference point because they reserve varying amounts of space for calling different functions.) Note, also, that the value stored at `(%rbp)` is the caller's `%rbp`, and the value stored at `8(%rbp)` is the return address. This information can be used to trace backwards through callers' stack frames by functions such as debuggers.
3. The function ends with `movq %rbp, %rsp; popq %rbp; retq`, or, equivalently, `leave; retq`. This sequence restores the caller's `%rbp` and entry `%rsp` before returning.

### Stack size and red zone

Functions execute fast because allocating space within a function is simply a matter of decrementing `%rsp`. This is much cheaper than a call to `malloc` or `new`! But making this work

takes a lot of machinery. We'll see this in more detail later; but in brief: The operating system knows that `%rsp` points to the stack, so if a function accesses nonexistent memory near `%rsp`, the OS assumes it's for the stack and transparently allocates new memory there.

So how can a program "run out of stack"? The operating system puts a limit on each function's stack, and if `%rsp` gets too low, the program segmentation faults.

The diagram above also shows a nice feature of the x86-64 architecture, namely the **red zone**. This is a small area *above* the stack pointer (that is, at lower addresses than `%rsp`) that can be used by the currently-running function for local variables. The red zone is nice because it can be used without mucking around with the stack pointer; for small functions `push` and `pop` instructions end up taking time.

### Branches

The processor typically executes instructions in sequence, incrementing `%rip` each time.

Deviations from sequential instruction execution, such as function calls, are called **control flow transfers**.

Function calls aren't the only kind of control flow transfer. A *branch* instruction jumps to a new instruction without saving a return address on the stack.

Branches come in two flavors, unconditional and conditional. The `jmp` or `j` instruction executes an unconditional branch (like a `goto`). All other branch instructions are conditional: they only branch if some condition holds. That condition is represented by **condition flags** that are set as a side effect of every arithmetic operation.

Arithmetic instructions change part of the `%rflags` register as a side effect of their operation. The most often used flags are:

- **ZF** (zero flag): set iff the result was zero.
- **SF** (sign flag): set iff the most significant bit (the sign bit) of the result was one (i.e., the result was negative if considered as a signed integer).
- **CF** (carry flag): set iff the result overflowed when considered as unsigned (i.e., the result was greater than  $2^W-1$ ).
- **OF** (overflow flag): set iff the result overflowed when considered as signed (i.e., the result was greater than  $2^{W-1}-1$  or less than  $-2^{W-1}$ ).

Flags are most often accessed via conditional jump or conditional move instructions. The conditional branch instructions are:

| Instruction            | Mnemonic                    | C example                             | Flags             |
|------------------------|-----------------------------|---------------------------------------|-------------------|
| <code>j (jmp)</code>   | Jump                        | <code>break;</code>                   | (Unconditional)   |
| <code>je (jz)</code>   | Jump if equal (zero)        | <code>if (x == y)</code>              | ZF                |
| <code>jne (jnz)</code> | Jump if not equal (nonzero) | <code>if (x != y)</code>              | !ZF               |
| <code>jg (jnl)</code>  | Jump if greater             | <code>if (x &gt; y), signed</code>    | !ZF && !(SF ^ OF) |
| <code>jge (jnl)</code> | Jump if greater or equal    | <code>if (x &gt;= y), signed</code>   | !(SF ^ OF)        |
| <code>jl (jnge)</code> | Jump if less                | <code>if (x &lt; y), signed</code>    | SF ^ OF           |
| <code>jle (jng)</code> | Jump if less or equal       | <code>if (x &lt;= y), signed</code>   | (SF ^ OF)    ZF   |
| <code>ja (jnbe)</code> | Jump if above               | <code>if (x &gt; y), unsigned</code>  | !CF && !ZF        |
| <code>jae (jnb)</code> | Jump if above or equal      | <code>if (x &gt;= y), unsigned</code> | !CF               |
| <code>jb (jnae)</code> | Jump if below               | <code>if (x &lt; y), unsigned</code>  | CF                |
| <code>jbe (jna)</code> | Jump if below or equal      | <code>if (x &lt;= y), unsigned</code> | CF    ZF          |
| <code>js</code>        | Jump if sign bit            | <code>if (x &lt; 0), signed</code>    | SF                |
| <code>jns</code>       | Jump if not sign bit        | <code>if (x &gt;= 0), signed</code>   | !SF               |
| <code>jc</code>        | Jump if carry bit           | N/A                                   | CF                |
| <code>jnc</code>       | Jump if not carry bit       | N/A                                   | !CF               |
| <code>jo</code>        | Jump if overflow bit        | N/A                                   | OF                |
| <code>jno</code>       | Jump if not overflow bit    | N/A                                   | !OF               |

The `test` and `cmp` instructions are frequently seen before a conditional branch. These operations perform arithmetic but throw away the result, except for condition codes. `test` performs binary-and, `cmp` performs subtraction.

`cmp` is hard to grasp: remember that `subq %rax, %rbx` performs `%rbx := %rbx - %rax`—the source/destination operand is on the left. So `cmpq %rax, %rbx` evaluates `%rbx - %rax`. The sequence `cmpq %rax, %rbx; jg L` will jump to label `L` if and only if `%rbx` is greater than `%rax` (signed).

The weird-looking instruction `testq %rax, %rax`, or more generally `testq REG, SAMEREG`, is used to load the condition flags appropriately for a single register. For example, the bitwise-and of `%rax` and `%rax` is zero if and only if `%rax` is zero, so `testq %rax, %rax; je L` jumps to `L` if and only if `%rax` is zero.

You will occasionally see instructions named `setFLAG` that load the binary value of a condition (0 or 1) into an 8-bit register. For example, `setz %al` sets `%al` to 1 if the zero flag ZF is on, and 0 otherwise. There are `set` constructions for all conditions. `set` instructions are often followed by zero-extending instructions: `setz %al; movzbl %al, %eax` sets *all of* `%rax` to 1 if the zero flag ZF is on, and 0 otherwise.

Data-movement and control-flow instructions do not modify flags. Oddly, for example, `lea` does not modify flags (it counts as data movement), though `add` does (it counts as arithmetic).

#### Sidebar: C++ data structures

C++ compilers and data structure implementations have been designed to avoid the so-called *abstraction penalty*, which is when convenient data structures compile to more and more-expensive instructions than simple, raw memory accesses. When this works, it works quite well; for example, this:

```
long f(std::vector<int>& v) {
    long sum = 0;
    for (auto& i : v) {
        sum += i;
    }
    return sum;
}
```

compiles to this, a very tight loop similar to the C version:

```
    movq    (%rdi), %rax
    movq    8(%rdi), %rcx
    cmpq    %rcx, %rax
    je      .L4
    movq    %rax, %rdx
    addq    $4, %rax
    subq    %rax, %rcx
    andq    $-4, %rcx
    addq    %rax, %rcx
    movl    $0, %eax
.L3:
    movslq  (%rdx), %rsi
    addq    %rsi, %rax
    addq    $4, %rdx
    cmpq    %rcx, %rdx
    jne     .L3
    rep ret
.L4:
    movl    $0, %eax
    ret
```

We can also use this output to infer some aspects of `std::vector`'s implementation. It looks like:

- The first element of a `std::vector` structure is a pointer to the first element of the vector;
- The elements are stored in memory in a simple array;
- The second element of a `std::vector` structure is a pointer to *one-past-the-end* of the elements of the vector (i.e., if the vector is empty, the first and second elements of the structure have the same value).

#### Compiler optimizations

##### Argument elision

A compiler may decide to elide (or remove) certain operations setting up function call arguments, if it can decide that the registers containing these arguments will hold the correct value before the function call takes place. Let's see an example of a function disassembled

function `f` in `f31.s`:

```
subq    $8, %rsp
call    _Z1gi@PLT
addq    $8, %rsp
addl    $1, %eax
ret
```

This function calls another function `g`, adds 1 to `g`'s return value, and returns that value.

It is possible that the function has the following definition in C++:

```
int f() {
    return 1 + g();
}
```

However, the actual definition of `f` in `f31.cc` is:

```
int f(x) {
    return 1 + g(x);
}
```

The compiler realizes that the argument to function `g`, which is passed via register `%rdi`, already has the right value when `g` is called, so it doesn't bother doing anything about it. This is one example of numerous optimizations a compiler can perform to reduce the size of generated code.

##### Inlining

A compiler may also copy the body of function to its call site, instead of doing an explicit function call, when it decides that the overhead of performing a function call outweighs the overhead of doing this copy. For example, if we have a function `g` defined as `g(x) = 2 + x`, and `f` is defined as `f(x) = 1 + g(x)`, then the compiler may actually generate `f(x)` as simply `3 + x`, without inserting any `call` instructions. In assembly terms, function `g` will look like

```
leal    2(%rdi), %eax
ret
```

and `f` will simply be

```
leal    3(%rdi), %eax
ret
```

##### Tail call elimination

Let's look at another example in `f32.s`:

```
addl    $1, %edi
jmp     _Z1gi@PLT
```

This function doesn't even contain a `ret` instruction! What is going on? Let's take a look at the actual definition of `f`, in `f32.cc`:

```
int f(int x) {
    return g(x + 1);
}
```

Note that the call to function `g` is the last operation in function `f`, and the return value of `f` is just the return value of the invocation of `g`. In this case the compiler can perform a *tail call*

*elimination*: instead of calling `g` explicitly, it can simply jump to `g` and have `g` return to the same address that `f` would have returned to.

A tail call elimination may occur if a function (caller) ends with another function call (callee) and performs no cleanup once the callee returns. In this case the caller and simply jump to the callee, instead of doing an explicit call.

Loop unrolling

Before we jump into loop unrolling, let's take a small excursion into an aspect of calling conventions called caller/callee-saved registers. This will help us under the sample program in `f33.s` better.

*Calling conventions: caller/callee-saved registers*

Let's look at the function definition in `f33.s`:

```
pushq    %r12
pushq    %rbp
pushq    %rbx
testl    %edi, %edi
je       .L4
movl     %edi, %r12d
movl     $0, %ebx
movl     $0, %ebp
.L3:
movl     %ebx, %edi
call     _Z1gj@PLT
addl     %eax, %ebp
addl     $1, %ebx
cmpl     %ebx, %r12d
jne      .L3
.L1:
movl     %ebp, %eax
popq     %rbx
popq     %rbp
popq     %r12
ret
.L4:
movl     %edi, %ebp
jmp      .L1
```

From the assembly we can tell that the backwards jump to `.L3` is likely a loop. The loop index is in `%ebx` and the loop bound is in `%r12d`. Note that upon entry to the function we first moved the value `%rdi` to `%r12d`. This is necessary because in the loop `f` calls `g`, and `%rdi` is used to pass arguments to `g`, so we must move its value to a different register to use it as the loop bound (this case `%r12`). But there is more to this: the compiler also needs to ensure that this register's value is preserved across function calls. Calling conventions dictate that certain registers always exhibit this property, and they are called **callee-saved registers**. If a register is callee-saved, then the caller doesn't have to save its value before entering a function call. We note that upon entry to the function, `f` saved a bunch of registers by pushing them to the stack: `%r12, %rbp, %rbx`. It is because all these registers are callee-saved registers, and `f` uses them during the function call. In general, the following registers in x86\_64 are callee-saved:

`%rbx, %r12-%r15, %rbp, %rsp (%rip)`

All the other registers are **caller-saved**, which means the callee doesn't have to preserve their values. If the caller wants to reuse values in these registers across function calls, it will have to explicitly save and restore these registers. In general, the following registers in x86\_64 are caller-saved:

`%rax, %rcx, %rdx, %r8-%r11`

Now let's get back to loop unrolling. Let us look at the program in `f34.s`:

```
testl    %edi, %edi
je       .L7
```

```
leal     -1(%rdi), %eax
cmpl     $7, %eax
jbe      .L8
pxor     %xmm0, %xmm0
movl     %edi, %edx
xorl     %eax, %eax
movdqa   .LC0(%rip), %xmm1
shrl     $2, %edx
movdqa   .LC1(%rip), %xmm2
.L5:
addl     $1, %eax
padd     %xmm1, %xmm0
padd     %xmm2, %xmm1
cmpl     %edx, %eax
jb       .L5
movdqa   %xmm0, %xmm1
movl     %edi, %edx
andl     $-4, %edx
psrldq   $8, %xmm1
padd     %xmm1, %xmm0
movdqa   %xmm0, %xmm1
cmpl     %edx, %edi
psrldq   $4, %xmm1
padd     %xmm1, %xmm0
movd     %xmm0, %eax
je       .L10
.L3:
leal     1(%rdx), %ecx
addl     %edx, %eax
cmpl     %ecx, %edi
je       .L1
addl     %ecx, %eax
leal     2(%rdx), %ecx
cmpl     %ecx, %edi
je       .L1
addl     %ecx, %eax
leal     3(%rdx), %ecx
cmpl     %ecx, %edi
je       .L1
addl     %ecx, %eax
leal     4(%rdx), %ecx
cmpl     %ecx, %edi
je       .L1
addl     %ecx, %eax
leal     5(%rdx), %ecx
cmpl     %ecx, %edi
je       .L1
addl     %ecx, %eax
leal     6(%rdx), %ecx
cmpl     %ecx, %edi
je       .L1
addl     %ecx, %eax
addl     $7, %edx
leal     (%rax,%rdx), %ecx
cmpl     %edx, %edi
cmovne   %ecx, %eax
ret
.L7:
xorl     %eax, %eax
```

```

.L1:
    rep ret
.L10:
    rep ret
.L8:
    xorl    %edx, %edx
    xorl    %eax, %eax
    jmp     .L3

```

Wow this looks long and repetitive! Especially the section under label `.L3`! If we take a look at the original function definition in `f34.cc`, we will find that it's almost the same as `f33.cc`, except that in `f34.cc` we know the definition of `g` as well. With knowledge of what `g` does the compiler's optimizer decides that unrolling the loop into 7-increment batches results in faster code.

Code like this can become difficult to understand, especially when the compiler begins to use more advanced registers reserved for vector operations. We can fine-tune the optimizer to disable certain optimizations. For example, we can use the `-mno-sse -fno-unroll-loops` compiler options to disable the use of SSE registers and loop unrolling. The resulting code, in `f35.s`, for the same function definitions in `f34.cc`, becomes much easier to understand:

```

    testl   %edi, %edi
    je      .L4
    xorl    %edx, %edx
    xorl    %eax, %eax
.L3:
    addl    %edx, %eax
    addl    $1, %edx
    cmpl    %edx, %edi
    jne     .L3
    rep ret
.L4:
    xorl    %eax, %eax
    ret

```

Note that the compiler still performed inlining to eliminate function `g`.

Optimizing recursive functions

Let's look at the following recursive function in `f36.cc`:

```

int f(int x) {
    if (x > 0) {
        return x * f(x - 1);
    } else {
        return 0;
    }
}

```

At the first glance it may seem that the function returns factorial of `x`. But it actually returns 0. Despite it doing a series of multiplications, in the end it always multiplies the whole result with 0, which produces 0.

When we compile this function to assembly without much optimization, we see the expensive computation occurring:

```

    movl    $0, %eax
    testl   %edi, %edi
    jg      .L8
    rep ret
.L8:
    pushq   %rbx
    movl    %edi, %ebx
    leal    -1(%rdi), %edi
    call    _Zlfi

```

```

    imull    %ebx, %eax
    popq     %rbx
    ret

```

In `f37.cc` there is an actual factorial function:

```

int f(int x) {
    if (x > 0) {
        return x * f(x - 1);
    } else {
        return 1;
    }
}

```

If we compile this function using level-2 optimization (`-O2`), we get the following assembly:

```

    testl   %edi, %edi
    movl    $1, %eax
    jle     .L4
.L3:
    imull    %edi, %eax
    subl    $1, %edi
    jne     .L3
    rep ret
.L4:
    rep ret

```

There is no `call` instructions again! The compiler has transformed the recursive function into a loop.

If we revisit our "fake" factorial function that always returns 0, and compile it with `-O2`, we see yet more evidence of compiler's deep understanding of our program:

```

xorl    %eax, %eax
ret

```

Optimizing arithmetic operations

The assembly code in `f39.s` looks like this:

```

leal    (%rdi,%rdi,2), %eax
leal    (%rdi,%rax,4), %eax
ret

```

It looks like some rather complex address computations! The first `leal` instruction basically loads `%eax` with value `3*%rdi` (or `%rdi + 2*%rdi`). The second `leal` multiplies the previous result by another 4, and adds another `%rdi` to it. So what it actually does is `3*%rdi*4 + %rdi`, or simply `13*%rdi`. This is also revealed in the function name in `f39.s`.

The compiler choose to use `leal` instructions instead of an explicit multiply because the two `leal` instructions actually take less space.

## 關於 GNU Inline Assembly

DEC 10TH, 2015 10:16 PM | [COMMENTS](#)

以前稍微接觸過 GNU Inline Assembly，對於那些奇怪的符號總是覺得匪夷所思。這次找時間把他整理一下。雖然釐清了一些觀念，不過卻產生更多的疑惑，也許以後有機會看到範例會慢慢有感覺吧。

### 目錄

- [前言](#)
- [測試環境](#)
- [語法](#)
  - [Output operands](#)
  - [Input operands](#)
  - [Clobbered registers list](#)
  - [Constraints](#)
- [參考資料](#)

### 前言

我自己對於 GNU Inline Assembly 的看法。

- 編譯器 夠聰明，所以暫存器分配可以安心交給編譯器處理。也就是說語法上面要處理這塊。



- 暫存器、變數有些資訊仍然要讓編譯器知道，讓編譯器產生 object binary 遵守這樣的規則，如
  - 這個 operand 是一個暫存器
  - 這個 operand 是一塊記憶體
  - 這個 operand 是浮點常數
  - ...
- 不想讓編譯器幫你安排暫存器，而是在 Inline Assembly 指定暫存器的話，就要明確的列出來。讓編譯器知道這些暫存器有被改過資料，進而針對這些暫存器做適當的處理。

測試環境

我使用 ARMv7 為主的 Banana Pi 開發版加上 Ubuntu 14.04 作為測試環境。

```
$ lsb_release -a
No LSB modules are available.
Distributor ID: Ubuntu
Description: Ubuntu 14.04.3 LTS
Release: 14.04
Codename: trusty

$ dmesg
[0.000000] Linux version 3.4.90 (bananapi@lemaker) (gcc version 4.6.3 (Ubuntu/Linaro 4.6.3-1ubuntu5) ) #2 SMP PREEMPT Tue Aug 5 14:11:40 CST 2014
[0.000000] CPU: ARMv7 Processor [410fc074] revision 4 (ARMv7), cr=10c5387d
...
$ gcc -v
Target: arm-linux-gnueabi
gcc version 4.8.4 (Ubuntu/Linaro 4.8.4-2ubuntu1~14.04)
```

語法

inline assembler 關鍵字是 asm，不過\_\_asm\_\_也可以使用(註)。

根據目前(Dec/2015)的 gcc 手冊，inline assembler 有分為 basic 和 extended 兩種。雖然我使用的平台是 gcc 4.8.4，而且 gcc 4.8.5 手冊(官方網站上沒有 4.8.4 手冊)並沒有提到這個部份。但是目前語法上測試的確沒有問題，但是有些說明上面卻很難驗證是否可以套用到 4.8.5 上(例如最佳化的說明、需要注意常犯的錯誤)，請自行斟酌。

以下是整理自最新的手冊說明，請自行斟酌您使用的 gcc 版本是否有符合。

Basic inline assembler

1 [volatile] asm("Assembler Template");

以下是整理自最新(Dec/2015)的手冊說明節錄，請自行斟酌您使用的 gcc 版本是否有符合。

- basic inline assembler 預設就是 volatile
- 基本上編譯器只是把引號內的東西抄錄，所以只要組譯器支援的語法，就可以寫入 Assembler Template 內
- 和 extended inline assembler 的差異
  - extended inline assembler 只允許在函數內使用
  - 有 naked 屬性的函數必須使用 basic inline assembler(見註解)
  - basic inline assembler 就是把 template 內的字串作為組合語言組譯。而%字元在 extended inline assembler 有特別意義，然而有些組合語言如 x86 中%是暫存器語法的一部份。以至於%字元要在 extended inline assembler 中改為\_\_才是真正的意思，舉個例子%eax->\_\_eax
- 有要使用 C 語言的資料，使用 extended inline assembler 比較妥當
- GCC 最佳化時是有可能把你的 inline assembler 幹掉或是和你想的不一樣，請注意
- 你不可以從一個 asm(..)裏面跳到另外一個 asm(..)的 label

最簡單的廢話範例如下

```
1 asm("nop"); /* 啥事都不要做 */
在沒有使用 C 語言的變數下，就和一般的組合語言沒有差太多。 更複雜一點的例子可以看 rtenv 裏面的使用方式：
1 size_t strlen(const char *s) __attribute__((naked));
2 size_t strlen(const char *s)
```

```
3 {
4     asm(
5         "    sub    r3, r0, #1          \n"
6         "strlen_loop:                \n"
7         "    ldrb   r2, [r3, #1]!      \n"
8         "    cmp    r2, #0              \n"
9         "    bne    strlen_loop        \n"
10        "    sub    r0, r3, r0          \n"
11        "    bx     lr                  \n"
12        :::
13    );
14 }
```

要注意\_\_attribute\_\_((naked));是有意義的。這是為何這段範例沒有直接指名用到 C 語言函式變數名稱的關鍵點。有興趣請看這邊，請直接找字串 naked。

Extended inline assembler

```
1 asm [volatile] ( AssemblerTemplate
2                  : OutputOperands // optional
3                  [ : InputOperands // optional
4                  [ : Clobbers ]    // optional
5                  )
```

Assembler Template 基本上就是你要寫的組語加上 Inline Assembler 專用的符號。要注意的是，在編譯的過程中，你寫的 inline assembler 可能由於最佳化考慮不會被組譯。如果你確認你 inline assembler 一定要被組譯，請加上 volatile keyword。

Assembler 專用的符號節錄如下：

| 符號 | 說明                                  |
|----|-------------------------------------|
| %% | 單一%字元                               |
| %{ | 單一{字元                               |
| %} | 單一}字元                               |
| {  | 單一 字元                               |
| %= | 只知道並驗證過會產生唯一的數字。用途部份看不懂，英文真是奧妙的東西啊。 |

AssemblerTemplate

由於前言提到的三項個人猜測，造成 inline assembler 要使用 C 語言變數時語法會出現很多令人眼花撩亂的符號。

由於編譯器提供協助分配暫存器和記憶體，也就是說需要有對應的語法指定目前指令的 operand 是什麼。GCC 有兩種方式指定，分別是

- 編號指定，從零開始編號
- Symbolic name 指定: GCC 3.1 以後支援出處

分別給個範例讓各位感受一下

編號指定，從零開始編號

這邊%0,%1 就是編號。後面 operand 可以看到就是指定變數、以及變數的限制。這邊簡單解釋一下=表示這是一個輸出、而 r 表示變數要放在暫存器中、m 表示變數是放在記憶體中。有興趣比對編譯出來的 binary 反組譯時的組合語言請看這邊。不過編號和指令中的 operand 似乎很隨意，我沒有看到特殊規範。只能交叉比對 assembler template 和 input/output operands 才能看出端倪。我猜更複雜的情況你還要比對反組譯出來的結果。

```
1 #include <stdio.h>
2
3 int main(void)
4 {
5     int var1 = 12;
6     int var2 = 10;
7
8     asm("mov %0, %1 \n
9         add %1, %0, $1" : "=r"(var1), "=r"(var2) : "r"(var2), "r"(var1));
10    printf("var1 = %d, var2 = %d\n", var1, var2);
11
12    asm("ldr r5, %0 \n": : "m"(var1): "r5");
13    asm("str r4, %0" : "=m"(var2): : "r4");
14    return 0;
15 }
```

Symbolic name 指定

編號的缺點就是可讀性比較差，所以 gcc 3.1 出現使用 symbolic name 的方式。至於那一個比較好，看你自己習慣。  
直接把上面的範例更改一下。[GCC 4.8.5 手冊](#)上面說 symbolic name 隨便取，甚至和變數同名稱都可以，只要單一 `asm(...)`內的 symbolic name 不要重複就好。有興趣比對編譯出來的 binary 反組譯時的組合語言請看[這邊](#)。

```
1#include <stdio.h>
2
3int main(void)
4{
5    int var1 = 12;
6    int var2 = 10;
7
8    asm("mov %[my_var1], %[my_var2] \n \
9        add %[my_var3], %[my_var4], $1" :
10        [my_var1] "=r" (var1), [my_var3] "=r" (var2) :
11        [my_var2] "r" (var2), [my_var4] "r" (var1) :);
12    printf("var1 = %d, var2 = %d\n", var1, var2);
13
14    asm("ld r5, %[my_var1] \n :: [my_var1] \"m\"(var1): \"r5\"");
15    asm("str r4, %[my_var1]: [my_var1] \"=m\" (var2):: \"r4\"");
16    return 0;
17}
```

接下來來看每個欄位吧。

Output operands

1. `[asmSymbolicName] [constraint (cvariablename)]`  
[asmSymbolicName] 是 GCC 3.1 以後支援語法，如前所述，不用 Symbolic Name 就用編號方式對應 assembler template operand。  
指定結果要存在 C 語言中的那個變數。要注意的除了要設定對的資訊（constraints，[下面](#)會節錄）以外，operand 的 prefix 一定要是 `=`或`+`這兩個 constraint。  
隨便舉幾個範例

- `=r(var1)`：變數請寫入並放在暫存器中
- `=m(var1)`：變數請寫入並存到記憶體中

Input operands

1. `[asmSymbolicName] [constraint (cvariablename)]`  
[asmSymbolicName] 是 GCC 3.1 以後支援語法，如前所述，不用 Symbolic Name 就用編號方式對應 assembler template operand。  
指定要從在 C 語言中的那個變數取出資料。主要是要設定對的資訊（constraints，[下面](#)會節錄）。

- `r(var1)`：變數請放在暫存器中
- `m(var1)`：變數是在記憶體中

Clobbered registers list

先講結論，在 `asm("語法")`中明確地指定暫存器名稱的話，要在這邊列出。  
現在我會習慣查單字。Clobbered 查英文單字會發現就是把東西用力地砸毀。所以翻譯成中文就是「砸爛的暫存器列表」。什麼是爛掉的暫存器？就是本節前面的結論囉。  
另外從 Dec/2015 的 gcc 手冊還有找到下面[語法](#)，一樣請注意版本問題

| 符號       | 說明  |
|----------|---|
| "cc"     | 和狀態有關的 flag 暫存器會被修改                       |
| "memory" | 這段組合語言會讀寫列出 operand 以外的記憶體內容，因此編譯器會視情況備份暫 |

Constraints

```
1<Constraints> ::= <Constraint Modifier> <Other Constraints> | <Other Constraints>
2<Other Constraints> ::= <Simple Constraints> | <Machine Constraints>
3
4/* 以上 BNF 是我整理的，terminal symbol 請自行看手冊 */
```

節錄整理我[看得懂](#)感興趣的部份。

Simple Constraints

| 符號 | 說明 |
|----|----|
|----|----|

|    |   |
|----|---|
| 空白 | 會被忽略，排版用  |
| m  | operand 存放在記憶體中   |
| r  | operand 將被放在暫存器中  |
| i  | operand 是一個整數常數，該常數包含下面的情形(symbolic name):<br>`#define MAX_LINE (32)` |
| n  | operand 是一個整數常數，只允許填入數字   |
| E  | operand 是一個浮點數常數，不清楚和 F 的差異   |
| F  | operand 是一個浮點數常數，不清楚和 E 的差異   |
| g  | operand 存在暫存器(r)或是記憶體內(m)，或是這是一個整數常數                                  |
| X  | 不用檢查 operand  |

你可以使用組合技如 `"rim"`，如果這樣寫的話，意思是要編譯器幫你挑一個最適合的方式處理對應於 assembler template 內的 operand。

Constraint Modifier

| 符號 | 說明  |
|----|---|
| =  | 表示這是一個 write only 的 operand，必須為 constraint 開始字元。  |
| +  | 表示這個 operand 在指令中是同時被讀寫的，必須為 constraint 開始字元。<br>該 operand 為 earlyclobber。earlyclobber 就是在 instruction 讀取該 operand 前，該 operand 會被寫入。雖然如此，到底是多久前？是和 data hazard 有關嘛？還是跟資料一致性有關？或者是和編譯器 最佳化造成非預期結果有關？真是一團謎 <b>完全搞不懂做啥用，也不清楚使用時機</b> 。 <a href="#">這邊有範例</a> ，一樣搞不懂為什麼要有 +, & 的 modifier |
| %  | 該 operand 可以讓編譯器 決定這個 operand 是否和後面的 operand 交換(commutative)， <b>完全搞不懂做啥用</b>   |

ARM 專用的 Constraint

我參考的是 gcc 4.8.5 手冊(因為和測試環境的 gcc 版本最接近)，可能有版本的問題，這些我都沒有做實驗測試，請自行斟酌。

| 符號 | 說明(一般模式)  |
|----|---|
| w  | VFP 浮點運算  |
| G  | 浮點運算的 0.0   |
| I  | 8 bit 正整數   |
| K  | I constraint 的 invert (一的補數)，<br>Wen: 不知道為什麼要扯到 I constraint? |
| L  | I constraint 的負數 (二的補數)，<br>Wen: 不知道為什麼要扯到 I constraint?      |
| M  | 0 ~ 32 的正整數   |
| Q  | 要參考的記憶體位址存放在一個暫存器內  |
| R  | operand 是一個 const pool 內的東西，不要問我 const pool 是啥，估狗到都和 Java 有關  |
| S  | operand 目前檔案中.text 內的一個 symbol                                |
| Uv | VFP load/store 指令可存取的記憶體                                      |
| Uy | iWMMXt load/store 指令可存取的記憶體                                   |
| Uq | ARMv4 ldrsb 指令可存取的記憶體   |

完整列表在這邊，要注意的是 2015 年 12 月的手冊又多了些新的 contstraint。請自行參考。

- [gcc 4.8.5: Constraints for Particular Machines](#)
  - 請自行參考你的硬體平台

- [Dec/2015 手冊: 6.44.4.4 Constraints for Particular Machines](#)

參考資料

- 中文
  - [\(BIG5\)用 Open Source 工具開發軟體: 新軟體開發觀念: Chapter 4. GNU Compiler Collection](#)
    - 題外話，寫這位文件的作者個人非常佩服，但是網路上似乎關於這位作者只有這份文件。真是神祕的人物

- [Nano 雞排: Inline Assembly](#)
- 英文
  - [gcc: How to Use Inline Assembly Language in C Code](#)
    - 手冊上寫是給 gcc 6.0，我目前從 GGG release 網站上看到[最新版本](#)是 5.3，怪。
    - 為什麼列出這個，因為我原本找的 gcc-4.8.5 對於 assembler template 說明沒有特別列出 gcc 支援的 inline assembler 符號。另外這份的文件結構的確比 4.8.5 清楚。
  - [gcc-4.8.5: Assembler Instructions with C Expression Operands](#)
  - [gcc-4.8.5: Constraints for asm Operands](#)
  - [ARM GCC Inline Assembler Cookbook](#)
    - 相當推荐，不論是給的說明和範例，更厲害的是關於 constraint 部份寫的比手冊清楚，一樣我沒測過就是了。
  - [A Tiny Guide to GCC Inline Assembly](#)
  - [GCC-Inline-Assembly-HOWTO](#)
  - [OSDev: Inline Assembly](#)
  - [Introduction to GCC inline assembler](#)
    - 似乎有點古老。

## 附錄

- C 語言標準有提到編譯器可以使用 `asm` keyword，而且沒有定義語法。有興趣可以找 `C11`、`C99`、`C89` 的標準，直接搜尋 `asm` 就可以看到了。

- `naked` 使用 basic inline assembler 和 extended inline assembler 比較  
下面兩個函數，`strcmp1` 沒有任何 extended inline assembler 而 `strcmp2` 硬塞了一個下去：

```

1 int strcmp1(const char *a, const char *b) __attribute__((naked));
2 int strcmp1(const char *a, const char *b)
3 {
4     asm(
5         "strcmp_lop1: \n"
6         "    ldrb    r2, [r0], #1    \n"
7         "    ldrb    r3, [r1], #1    \n"
8         "    cmp     r2, #1          \n"
9         "    it      hi              \n"
10        "    cmphi   r2, r3          \n"
11        "    beq     strcmp_lop1     \n"
12        "    sub     r0, r2, r3      \n"
13        "    bx      lr              \n"
14        :::
15    );
16 }
17
18 int strcmp2(const char *a, const char *b) __attribute__((naked));
19 int strcmp2(const char *a, const char *b)
20 {
21     int i;
22     asm(
23         "strcmp_lop2: \n"
24         "    ldrb    r2, [r0], #1    \n"
25         "    ldrb    r3, [r1], #1    \n"
26         "    cmp     r2, #1          \n"
27         "    it      hi              \n"
28         "    cmphi   r2, r3          \n"
29         "    mov     %1, $1 \n"
30         "    beq     strcmp_lop2     \n"
31         "    sub     r0, r2, r3      \n"
32         "    bx      lr              \n"
33         : "=r"(i)::
34     );
35 }
```

我們可以比較一下下面兩個函數最後編譯出來的指令，`strcmp2` 顯然和我們預期的差很多。

```

1 000083f4 <strcmp1>:
2
3
4 int strcmp1(const char *a, const char *b) __attribute__((naked));
5 int strcmp1(const char *a, const char *b)
6 {
7     asm(
8         83f4: f810 2b01    ldrb.w r2, [r0], #1
9         83f8: f811 3b01    ldrb.w r3, [r1], #1
10        83fc: 2a01      cmp    r2, #1
11        83fe: bf88      it     hi
12        8400: 429a      cmphi   r2, r3
13        8402: d0f7      beq.n  83f4 <strcmp1>
14        8404: eba2 0003    sub.w  r0, r2, r3
15        8408: 4770      bx      lr
16        "    beq     strcmp_lop1     \n"
17        "    sub     r0, r2, r3      \n"
18        "    bx      lr              \n"
19        :::
20    );
21 }
22
23 840a: 4618      mov    r0, r3
24
25 0000840c <strcmp2>:
26
27     "    beq     strcmp_lop2     \n"
28     "    sub     r0, r2, r3      \n"
29     "    bx      lr              \n"
30     : "=r"(i)::
31 );
```

## 範例一的反組譯節錄

```

1 $ objdump -d -S asm
2
3 000083f4 <main>:
4 #include <stdio.h>
5
6 int main(void)
7 {
8
9     int var1 = 12;
10    83fa: 230c      movs   r3, #12
11    83fc: 603b      str    r3, [r7, #0]
12
13    int var2 = 10;
14    83fe: 230a      movs   r3, #10
15    8400: 607b      str    r3, [r7, #4]
16
17    asm("mov %0, %1 \n \n"
18        add %1, %0, $1 : "=r"(var1), "=r"(var2) : "r"(var2), "r"(var1););
19    8402: 687b      ldr    r3, [r7, #4]
20    8404: 683a      ldr    r2, [r7, #0]
21    8406: 461a      mov    r2, r3
22    8408: f102 0301 add.w   r3, r2, #1
23    840c: 603a      str    r2, [r7, #0]
24    840e: 607b      str    r3, [r7, #4]
25
26    asm("ldr r5, %0 \n": "m"(var1): "r5");
27    8424: 683d      ldr    r5, [r7, #0]
28
29    asm("str r4, %0" : "=m"(var2): : "r4");
```

```
30      8426:  607c          str    r4, [r7, #4]
31...
32}
33...
```

---

- 範例二的反組譯節錄

```
1$ objdump -d -S asm
2...
3000083f4 <main>;
4#include <stdio.h>
5
6int main(void)
7{
8...
9    int var1 = 12;
10   83fa:  230c          movs   r3, #12
11   83fc:  603b          str    r3, [r7, #0]
12
13   int var2 = 10;
14   83fe:  230a          movs   r3, #10
15   8400:  607b          str    r3, [r7, #4]
16
17   asm("mov %[my_var1], %[my_var2] \n \
18       add %[my_var3], %[my_var4], $1" :
19       [my_var1] "=r" (var1), [my_var3] "=r" (var2) :
20       [my_var2] "r" (var2), [my_var4] "r" (var1) :);
21   8402:  687b          ldr    r3, [r7, #4]
22   8404:  683a          ldr    r2, [r7, #0]
23   8406:  461a          mov    r2, r3
24   8408:  f102 0301     add.w  r3, r2, #1
25   840c:  603a          str    r2, [r7, #0]
26   840e:  607b          str    r3, [r7, #4]
27...
28   asm("ldr r5, %[my_var1] \n :: [my_var1] \"m\"(var1): \"r5\"");
29   8424:  683d          ldr    r5, [r7, #0]
30
31   asm("str r4, %[my_var1] : [my_var1] \"=m\"(var2):: \"r4\"");
32   8426:  607c          str    r4, [r7, #4]
33...
34}
```

---



```

// Set registers for destination
"movq  %[dst_col_stride_q], %%r12" "shlq $2, %%r12" "leaq (%%r12,%%r12,0x2), %%r13"
// Set accumulators to zero.
"pxor %%xmm4, %%xmm4 " "pxor %%xmm5, %%xmm5 " "pxor %%xmm6, %%xmm6 " "pxor %%xmm7, %%xmm7 "
"pxor %%xmm8, %%xmm8 " "pxor %%xmm9, %%xmm9 " "pxor %%xmm10, %%xmm10" "pxor %%xmm11, %%xmm11"
"pxor %%xmm12, %%xmm12" "pxor %%xmm13, %%xmm13" "pxor %%xmm14, %%xmm14" "pxor %%xmm15, %%xmm15"
"movq  %[run_depth_cells], %%r14"
"subq $2, %%r14"
"js outerLoop1%="
// Loop for K unrolled by 4
"outerLoop2%="
// K = 1,2 // RHS cell to xmm1
"pmovzxbw (%%rhs_ptr), %%xmm1"
// LHS cell
"pmovzxbw 0x00(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2 " "pshufd $0x55, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm4 " "padd %%xmm3, %%xmm5 "
"prefetcht0 0x80(%%lhs_ptr) "
"pshufd $0xaa, %%xmm1, %%xmm2 " "pmaddwd %%xmm0, %%xmm2 "
"pshufd $0xff, %%xmm1, %%xmm3 " "pmaddwd %%xmm0, %%xmm3 "
// next LHS cell
"pmovzxbw 0x08(%%lhs_ptr), %%xmm0"
"padd %%xmm2, %%xmm6 " "padd %%xmm3, %%xmm7 "

"pshufd $0x00, %%xmm1, %%xmm2 " "pshufd $0x55, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm8 " "padd %%xmm3, %%xmm9 "
"prefetcht0 0x80(%%rhs_ptr) "
"pshufd $0xaa, %%xmm1, %%xmm2 " "pshufd $0xff, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm10 " "padd %%xmm3, %%xmm11 "
// next LHS cell
"pmovzxbw 0x10(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2 " "pshufd $0x55, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm12 " "padd %%xmm3, %%xmm13 "
"pshufd $0xaa, %%xmm1, %%xmm2 " "pshufd $0xff, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm14 " "padd %%xmm3, %%xmm15 "
// K = 3,4 // RHS cell to xmm1
"pmovzxbw 0x08(%%rhs_ptr), %%xmm1"
// LHS cell
"pmovzxbw 0x18(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2 " "pshufd $0x55, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm4 " "padd %%xmm3, %%xmm5 "
"pshufd $0xaa, %%xmm1, %%xmm2 " "pshufd $0xff, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm6 " "padd %%xmm3, %%xmm7 "
// next LHS cell
"pmovzxbw 0x20(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2 " "pshufd $0x55, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm8 " "padd %%xmm3, %%xmm9 "
"pshufd $0xaa, %%xmm1, %%xmm2 " "pshufd $0xff, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm10 " "padd %%xmm3, %%xmm11 "
// next LHS cell
"pmovzxbw 0x28(%%lhs_ptr), %%xmm0"
"addq $0x30, %%lhs_ptr " "addq $0x10, %%rhs_ptr "

"pshufd $0x00, %%xmm1, %%xmm2 " "pshufd $0x55, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm12 " "padd %%xmm3, %%xmm13 "

"pshufd $0xaa, %%xmm1, %%xmm2 " "pshufd $0xff, %%xmm1, %%xmm3 "

```

```

"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm14 " "padd %%xmm3, %%xmm15 "
"subq $2, %%run_depth_cells"
"ja outerLoop2%="
"movq %%run_depth_cells, %%r14"
"decq %%r14"
"js finish%="
// Loop for K unrolled by 2
"outerLoop1%="
// RHS cell to xmm1
"pmovzxbw (%%rhs_ptr), %%xmm1"
// LHS cell
"pmovzxbw 0x00(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2 " "pshufd $0x55, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm4 " "padd %%xmm3, %%xmm5 "
"pshufd $0xaa, %%xmm1, %%xmm2 " "pshufd $0xff, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm6 " "padd %%xmm3, %%xmm7 "
// next LHS cell
"pmovzxbw 0x08(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2 " "pshufd $0x55, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm8 " "padd %%xmm3, %%xmm9 "
"pshufd $0xaa, %%xmm1, %%xmm2 " "pshufd $0xff, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm10 " "padd %%xmm3, %%xmm11 "
// next LHS cell
"pmovzxbw 0x10(%%lhs_ptr), %%xmm0"

"addq $0x18, %%lhs_ptr " "addq $0x08, %%rhs_ptr "
"pshufd $0x00, %%xmm1, %%xmm2 " "pshufd $0x55, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm12 " "padd %%xmm3, %%xmm13 "
"pshufd $0xaa, %%xmm1, %%xmm2 " "pshufd $0xff, %%xmm1, %%xmm3 "
"pmaddwd %%xmm0, %%xmm2 " "pmaddwd %%xmm0, %%xmm3 "
"padd %%xmm2, %%xmm14 " "padd %%xmm3, %%xmm15 "
"decq %%run_depth_cells"
"jnz outerLoop1%="
"finish%="
"test %%start_depth, %%start_depth"
"jz storeDst%="
"padd 0x00(%%dst_ptr), %%xmm4 " "padd 0x10(%%dst_ptr), %%xmm8 "
"padd 0x20(%%dst_ptr), %%xmm12 " "padd 0x00(%%dst_ptr, %%r12, 1), %%xmm5 "
"padd 0x10(%%dst_ptr, %%r12, 1), %%xmm9 " "padd 0x20(%%dst_ptr, %%r12, 1), %%xmm13 "
"padd 0x00(%%dst_ptr, %%r12, 2), %%xmm6 " "padd 0x10(%%dst_ptr, %%r12, 2), %%xmm10 "
"padd 0x20(%%dst_ptr, %%r12, 2), %%xmm14 " "padd 0x00(%%dst_ptr, %%r13, 1), %%xmm7 "
"padd 0x10(%%dst_ptr, %%r13, 1), %%xmm11 " "padd 0x20(%%dst_ptr, %%r13, 1), %%xmm15 "
"storeDst%="
"movdqu %%xmm4, 0x00(%%dst_ptr) " "movdqu %%xmm8, 0x10(%%dst_ptr) "
"movdqu %%xmm12, 0x20(%%dst_ptr) " "movdqu %%xmm5, 0x00(%%dst_ptr, %%r12, 1)"
"movdqu %%xmm9, 0x10(%%dst_ptr, %%r12, 1)" "movdqu %%xmm13, 0x20(%%dst_ptr, %%r12, 1)"
"movdqu %%xmm6, 0x00(%%dst_ptr, %%r12, 2)" "movdqu %%xmm10, 0x10(%%dst_ptr, %%r12, 2)"
"movdqu %%xmm14, 0x20(%%dst_ptr, %%r12, 2)" "movdqu %%xmm7, 0x00(%%dst_ptr, %%r13, 1)"
"movdqu %%xmm11, 0x10(%%dst_ptr, %%r13, 1)" "movdqu %%xmm15, 0x20(%%dst_ptr, %%r13, 1)"

```

```

: // outputs
[%%lhs_ptr] "+r"(%%lhs_ptr), [%%rhs_ptr] "+r"(%%rhs_ptr),
[%%dst_ptr] "+r"(%%dst_ptr)
: // inputs
[start_depth] "r"(start_depth),
[dst_col_stride_q] "r"(dst_col_stride_q),
[run_depth_cells] "r"(run_depth_cells)
: // clobbers
"cc", "memory", "xmm0", "xmm1", "xmm3", "xmm2", "xmm4", "xmm5",
"xmm6", "xmm7", "xmm8", "xmm9", "xmm10", "r12", "r13", "r14",
"xmm11", "xmm12", "xmm13", "xmm14", "xmm15");

```

```

// Set registers for destination
"movq %[dst_col_stride_q], %%r12"
"shlq $2, %%r12"
"leaq (%%r12,%%r12,0x2), %%r13"
// Set accumulators to zero.
"pxor %%xmm4, %%xmm4" "pxor %%xmm5, %%xmm5"
"pxor %%xmm6, %%xmm6" "pxor %%xmm7, %%xmm7"
"pxor %%xmm8, %%xmm8" "pxor %%xmm9, %%xmm9"
"pxor %%xmm10, %%xmm10" "pxor %%xmm11, %%xmm11"
"pxor %%xmm12, %%xmm12" "pxor %%xmm13, %%xmm13"
"pxor %%xmm14, %%xmm14" "pxor %%xmm15, %%xmm15"
"movq %[run_depth_cells], %%r14"
"subq $2, %%r14"
"js outerLoop1%"

// Loop for K unrolled by 4
"outerLoop2%="
// K = 1,2
// RHS cell to xmm1
"pmovzxbw (%%rhs_ptr), %%xmm1"
// LHS cell
"pmovzxbw 0x00(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2" "pshufd $0x55, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm4" "padd %%xmm3, %%xmm5"
"prefetcht0 0x80(%%lhs_ptr)"
"pshufd $0xaa, %%xmm1, %%xmm2" "pmaddwd %%xmm0, %%xmm2"
"pshufd $0xff, %%xmm1, %%xmm3" "pmaddwd %%xmm0, %%xmm3"
// next LHS cell
"pmovzxbw 0x08(%%lhs_ptr), %%xmm0"
"padd %%xmm2, %%xmm6" "padd %%xmm3, %%xmm7"

"pshufd $0x00, %%xmm1, %%xmm2" "pshufd $0x55, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm8" "padd %%xmm3, %%xmm9"
"prefetcht0 0x80(%%rhs_ptr)"
"pshufd $0xaa, %%xmm1, %%xmm2" "pshufd $0xff, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm10" "padd %%xmm3, %%xmm11"
// next LHS cell
"pmovzxbw 0x10(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2" "pshufd $0x55, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm12" "padd %%xmm3, %%xmm13"
"pshufd $0xaa, %%xmm1, %%xmm2" "pshufd $0xff, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm14" "padd %%xmm3, %%xmm15"
// K = 3,4
// RHS cell to xmm1
"pmovzxbw 0x08(%%rhs_ptr), %%xmm1"
// LHS cell
"pmovzxbw 0x18(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2" "pshufd $0x55, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm4" "padd %%xmm3, %%xmm5"
"pshufd $0xaa, %%xmm1, %%xmm2" "pshufd $0xff, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm6" "padd %%xmm3, %%xmm7"
// next LHS cell
"pmovzxbw 0x20(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2" "pshufd $0x55, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm8" "padd %%xmm3, %%xmm9"
"pshufd $0xaa, %%xmm1, %%xmm2" "pshufd $0xff, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm10" "padd %%xmm3, %%xmm11"
// next LHS cell
"pmovzxbw 0x28(%%lhs_ptr), %%xmm0"
"addq $0x30, %%lhs_ptr" "addq $0x10, %%rhs_ptr"

"pshufd $0x00, %%xmm1, %%xmm2" "pshufd $0x55, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm12" "padd %%xmm3, %%xmm13"

"pshufd $0xaa, %%xmm1, %%xmm2" "pshufd $0xff, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm14" "padd %%xmm3, %%xmm15"
"subq $2, %[run_depth_cells]"
"ja outerLoop2%"

// RHS cell to xmm1
"pmovzxbw (%%rhs_ptr), %%xmm1"
// LHS cell
"pmovzxbw 0x00(%%lhs_ptr), %%xmm0"
"pshufd $0x00, %%xmm1, %%xmm2" "pshufd $0x55, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm8" "padd %%xmm3, %%xmm9"
"pshufd $0xaa, %%xmm1, %%xmm2" "pshufd $0xff, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm10" "padd %%xmm3, %%xmm11"
// next LHS cell
"pmovzxbw 0x08(%%lhs_ptr), %%xmm0"
"addq $0x18, %%lhs_ptr" "addq $0x08, %%rhs_ptr"

"pshufd $0x00, %%xmm1, %%xmm2" "pshufd $0x55, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm12" "padd %%xmm3, %%xmm13"
"pshufd $0xaa, %%xmm1, %%xmm2" "pshufd $0xff, %%xmm1, %%xmm3"
"pmaddwd %%xmm0, %%xmm2" "pmaddwd %%xmm0, %%xmm3"
"padd %%xmm2, %%xmm14" "padd %%xmm3, %%xmm15"
"decq %[run_depth_cells]"
"jnz outerLoop1%"

"finish%="
"finish%="

"test %[start_depth], %[start_depth]"
"jz storeDst%"
"padd0x00(%%dst_ptr), %%xmm4" "padd0x10(%%dst_ptr), %%xmm8"
"padd0x20(%%dst_ptr), %%xmm12" "padd0x00(%%dst_ptr), %%r12, 1, %%xmm5"
"padd0x10(%%dst_ptr), %%r12, 1, %%xmm9" "padd0x20(%%dst_ptr), %%r12, 1, %%xmm13"
"padd0x00(%%dst_ptr), %%r12, 2, %%xmm6" "padd0x10(%%dst_ptr), %%r12, 2, %%xmm10"
"padd0x20(%%dst_ptr), %%r12, 2, %%xmm14" "padd0x00(%%dst_ptr), %%r13, 1, %%xmm7"
"padd0x10(%%dst_ptr), %%r13, 1, %%xmm11" "padd0x20(%%dst_ptr), %%r13, 1, %%xmm15"
"storeDst%="
"movdqu %%xmm4, 0x00(%%dst_ptr)" "movdqu %%xmm8, 0x10(%%dst_ptr)"
"movdqu %%xmm12, 0x20(%%dst_ptr)" "movdqu %%xmm5, 0x00(%%dst_ptr), %%r12, 1)"
"movdqu %%xmm9, 0x10(%%dst_ptr), %%r12, 1)" "movdqu %%xmm13, 0x20(%%dst_ptr), %%r12, 1)"
"movdqu %%xmm6, 0x00(%%dst_ptr), %%r12, 2)" "movdqu %%xmm10, 0x10(%%dst_ptr), %%r12, 2)"
"movdqu %%xmm14, 0x20(%%dst_ptr), %%r12, 2)" "movdqu %%xmm7, 0x00(%%dst_ptr), %%r13, 1)"
"movdqu %%xmm11, 0x10(%%dst_ptr), %%r13, 1)" "movdqu %%xmm15, 0x20(%%dst_ptr), %%r13, 1)"

: // outputs
[ lhs_ptr ] "+r"(lhs_ptr), [ rhs_ptr ] "+r"(rhs_ptr),
[ dst_ptr ] "+r"(dst_ptr)
: // inputs
[ start_depth ] "r"(start_depth),
[ dst_col_stride_q ] "r"(dst_col_stride_q),
[ run_depth_cells ] "r"(run_depth_cells)
: // clobbers
"cc", "memory", "%xmm0", "%xmm1", "%xmm3", "%xmm2", "%xmm4", "%xmm5",
"%xmm6", "%xmm7", "%xmm8", "%xmm9", "%xmm10", "%r12", "%r13", "%r14",
"%xmm11", "%xmm12", "%xmm13", "%xmm14", "%xmm15";

+-----+-----+-----+-----+-----+-----+
|xmm0| |xmm1| |xmm2| |xmm3| |xmm4| |xmm5| |xmm6| |xmm7| |
|xmm0| |(Iter1)| |xmm4| |xmm5| |xmm6| |xmm7| |
|xmm0| |xmm4| |xmm5| |xmm6| |xmm7| |
|xmm0| |xmm4| |xmm5| |xmm6| |xmm7| |
+-----+-----+-----+-----+-----+-----+
|xmm0| |xmm8| |xmm9| |xmm10| |xmm11| |
|xmm0| |(Iter2)| |xmm8| |xmm9| |xmm10| |xmm11| |
|xmm0| |xmm8| |xmm9| |xmm10| |xmm11| |
|xmm0| |xmm8| |xmm9| |xmm10| |xmm11| |
+-----+-----+-----+-----+-----+-----+
|xmm0| |xmm12| |xmm13| |xmm14| |xmm15| |
|xmm0| |(Iter3)| |xmm12| |xmm13| |xmm14| |xmm15| |
|xmm0| |xmm12| |xmm13| |xmm14| |xmm15| |
|xmm0| |xmm12| |xmm13| |xmm14| |xmm15| |
+-----+-----+-----+-----+-----+-----+

```

```

131 //void Run(std::int32_t* dst_ptr, std::size_t dst_row_stride,
132 // std::size_t dst_col_stride, const std::uint8_t* lhs_ptr,
133 // const std::uint8_t* rhs_ptr, std::size_t start_depth,
134 // std::size_t run_depth) const override
135 for (int r = 0; r < rows; r += block_params.l2_rows) {
136     int rs = std::min(block_params.l2_rows, rows - r);
137
138     PackLhs(&packed_lhs, lhs.block(r, 0, rs, depth));
139
140     for (int c = 0; c < cols; c += block_params.l2_cols) {
141         int cs = std::min(block_params.l2_cols, cols - c);
142
143         if (!pack_rhs_once) {
144             PackRhs(&packed_rhs, rhs.block(0, c, depth, cs));
145         }
146
147         Compute(kernel, block_params, &packed_result, packed_lhs, packed_rhs,
148             depth);
149
150         UnpackResult<KernelFormat>(
151             result, MatrixBlockBounds(r, c, rs, cs), packed_result, depth,
152             packed_lhs.sums_of_each_slice(), packed_rhs.sums_of_each_slice(),
153             lhs_offset.block(r, rs), rhs_offset.block(c, cs), output_pipeline);
154     }
155 }
156
157 allocator->Decommit();

```

|                      |   |
|----------------------|---|
| block_params.l2_rows | 4 |
| block_params.l2_cols | 4 |
| rs                   | 3 |
| cs                   | 2 |

```

105 template <typename PackedLhs, typename PackedRhs, typename PackedResult>
106 void Compute(const KernelBase& kernel, const BlockParams& block_params,
107             PackedResult* packed_result, const PackedLhs& packed_lhs,
108             const PackedRhs& packed_rhs, int depth) {
109     ScopedProfilingLabel label("compute");
110     ComputeImpl<PackedLhs, PackedRhs, PackedResult> impl(
111         kernel, block_params, packed_result, packed_lhs, packed_rhs);
112
113     impl.Compute(depth); 已用时间 <= 1ms
114 }

```

```

ComputeImpl(const KernelBase& _kernel, const BlockParams& _block_params,
            PackedResult* _packed_result, const PackedLhs& _packed_lhs,
            const PackedRhs& _packed_rhs)
: kernel_(_kernel),
  block_params_(_block_params),
  packed_result_(_packed_result),
  packed_lhs_(_packed_lhs),
  packed_rhs_(_packed_rhs) {}

void Compute(int depth) { 已用时间 <= 1ms
    depth = RoundUp<Format::kDepth>(depth);
    assert(depth <= block_params_.l2_depth);
    for (int d = 0; d < depth; d += block_params_.l1_depth) {
        int ds = std::min(block_params_.l1_depth, depth - d);

        for (int r = 0; r < block_params_.l2_rows; r += block_params_.l1_rows) {
            int rs = std::min(block_params_.l1_rows, block_params_.l2_rows - r);

            ComputeL1(r, rs, 0, block_params_.l2_cols, d, ds);
        }
    }

    void ComputeL1(int start_row, int rows, int start_col, int cols,
                    int start_depth, int depth) {
        assert(rows % Format::kRows == 0);
        assert(cols % Format::kCols == 0);
        assert(depth % Format::kDepth == 0);

        for (int c = 0; c < cols; c += Format::kCols) {
            for (int r = 0; r < rows; r += Format::kRows) {
                ComputeRun(start_row + r, start_col + c, start_depth, depth);
            }
        }
    }

    void ComputeRun(int start_row, int start_col, int start_depth,
                    int depth) GEMMLOWP_NOINLINE { 已用时间 <= 1ms
        packed_lhs_.seek_run(start_row, start_depth);
        packed_rhs_.seek_run(start_col, start_depth);
        auto packed_result_block = packed_result_>Map().block(
            start_row, start_col, Format::kRows, Format::kCols);
        kernel_.Run(packed_result_block.data(), packed_result_block.rows_stride(),
            packed_result_block.cols_stride(), packed_lhs_.current_data(),
            packed_rhs_.current_data(), start_depth, depth);
        MarkPackedResultBlockAsInitialized(packed_result_block);
    }
}

```

|             |    |
|-------------|----|
| start_row   | 0  |
| rows        | 4  |
| start_col   | 0  |
| cols        | 4  |
| start_depth | 0  |
| depth       | 16 |

```

void Run(std::int32_t* dst_ptr, std::size_t dst_row_stride,
        std::size_t dst_col_stride, const std::uint8_t* lhs_ptr,
        const std::uint8_t* rhs_ptr, std::size_t start_depth,
        std::size_t run_depth) const override { 已用时间<=1ms
    std::int32_t accumulator[Format::kRows * Format::kCols];
    memset(accumulator, 0, sizeof(accumulator));

    const int run_depth_cells = static_cast<int>(run_depth / Format::kDepth);

    // The outer loop is over the depth dimension.
    for (int dc = 0; dc < run_depth_cells; dc++) {
        // The next two loops are over cells of the Lhs (stacked vertically),
        // and over cells of the Rhs (stacked horizontally).
        for (int rc = 0; rc < Format::Lhs::kCells; rc++) {
            const std::uint8_t* lhs_cell_ptr =
                lhs_ptr + (dc * Format::Lhs::kCells + rc) *
                    Format::Lhs::Cell::kWidth * Format::kDepth;
            for (int cc = 0; cc < Format::Rhs::kCells; cc++) {
                const std::uint8_t* rhs_cell_ptr =
                    rhs_ptr + (dc * Format::Rhs::kCells + cc) *
                        Format::Rhs::Cell::kWidth * Format::kDepth;

                // Now we are inside one cell of the Lhs and inside one cell
                // of the Rhs, so the remaining inner loops are just
                // traditional three loops of matrix multiplication.
                for (int di = 0; di < Format::kDepth; di++) {
                    for (int ri = 0; ri < Format::Lhs::Cell::kWidth; ri++) {
                        for (int ci = 0; ci < Format::Rhs::Cell::kWidth; ci++) {
                            const std::uint8_t* lhs_coeff_ptr =
                                lhs_cell_ptr +
                                ▶ OffsetIntoCell<typename Format::Lhs::Cell>(ri, di);
                            const std::uint8_t* rhs_coeff_ptr =
                                rhs_cell_ptr +
                                OffsetIntoCell<typename Format::Rhs::Cell>(ci, di);
                            std::int32_t* accumulator_coeff_ptr =
                                accumulator + (ri + rc * Format::Lhs::Cell::kWidth) +
                                    (ci + cc * Format::Rhs::Cell::kWidth) * Format::kRows;
                            *accumulator_coeff_ptr +=
                                std::int32_t(*lhs_coeff_ptr) * std::int32_t(*rhs_coeff_ptr);
                        }
                    }
                }
            }
        }
    }
}

```

```

if (start_depth == 0) {
    // start_depth == 0 means we haven't accumulated anything yet, so we need
    // to overwrite the accumulator, as it hasn't been initialized to zero.
    for (int r = 0; r < Format::kRows; r++) {
        for (int c = 0; c < Format::kCols; c++) {
            dst_ptr[r * dst_row_stride + c * dst_col_stride] =
                accumulator[r + c * Format::kRows];
        }
    }
} else {
    // We have already accumulated stuff, so we need to continue accumulating
    // instead of just overwriting.
    for (int r = 0; r < Format::kRows; r++) {
        for (int c = 0; c < Format::kCols; c++) {
            dst_ptr[r * dst_row_stride + c * dst_col_stride] +=
                accumulator[r + c * Format::kRows];
        }
    }
}
};

```



```

;-----
:##### Ch02_01.asm
;-----
; extern "C" int IntegerAddSub_(int a, int b, int c, int d);

.code
IntegerAddSub_ proc
; Calculate a + b + c - d
    mov eax,ecx           ;eax = a
    add eax,edx           ;eax = a + b
    add eax,r8d           ;eax = a + b + c
    sub eax,r9d           ;eax = a + b + c - d

    ret                  ;return result to caller
IntegerAddSub_ endp
end

;-----
:##### Ch02_02.asm
;-----
; extern "C" unsigned int IntegerLogical_(unsigned int a, unsigned int b, unsigned int c,
unsigned int d);

extern g_Val1:dword       ;external doubleword (32-bit) value

.code
IntegerLogical_ proc
; Calculate ((a & b) | c) ^ d + g_Val1
    and ecx,edx           ;ecx = a & b
    or ecx,r8d            ;ecx = (a & b) | c
    xor ecx,r9d           ;ecx = ((a & b) | c) ^ d
    add ecx,[g_Val1]       ;ecx = ((a & b) | c) ^ d + g_Val1

    mov eax,ecx           ;eax = final result
    ret                  ;return to caller
IntegerLogical_ endp
end

;-----
:##### Ch02_03.asm
;-----
; extern "C" int IntegerShift_(unsigned int a, unsigned int count, unsigned int* a_shl, unsigned
int* a_shr);
; Returns:      0 = error (count >= 32), 1 = success

.code
IntegerShift_ proc
    xor eax,eax           ;set return code in case of error
    cmp edx,31            ;compare count against 31
    ja InvalidCount       ;jump if count > 31

    xchg ecx,edx           ;exchange contents of ecx & edx
    mov eax,edx           ;eax = a
    shl eax,cl            ;eax = a << count;
    mov [r8],eax          ;save result

    shr edx,cl            ;edx = a >> count
    mov [r9],edx          ;save result

    mov eax,1             ;set success return code

InvalidCount:
    ret                  ;return to caller

IntegerShift_ endp
end
;-----

```

```

:##### Ch02_04.asm
;-----
; extern "C" int IntegerMulDiv_(int a, int b, int* prod, int* quo, int* rem);
; Returns:      0 = error (divisor equals zero), 1 = success

.code
IntegerMulDiv_ proc
; Make sure the divisor is not zero
    mov eax,edx           ;eax = b
    or eax,eax            ;logical OR sets status flags
    jz InvalidDivisor     ;jump if b is zero

; Calculate product and save result
    imul eax,ecx          ;eax = a * b
    mov [r8],eax          ;save product

; Calculate quotient and remainder, save results
    mov r10d,edx          ;r10d = b
    mov eax,ecx           ;eax = a
    cdq                  ;edx:eax contains 64-bit dividend
    idiv r10d             ;eax = quotient, edx = remainder

    mov [r9],eax          ;save quotient
    mov rax,[rsp+40]       ;rax = 'rem'
    mov [rax],edx         ;save remainder
    mov eax,1             ;set success return code

InvalidDivisor:
    ret                  ;return to caller

IntegerMulDiv_ endp
end

;-----
:##### Ch02_05.asm
;-----
; extern "C" int64_t IntegerMul_(int8_t a, int16_t b, int32_t c, int64_t d, int8_t e, int16_t f,
int32_t g, int64_t h);

.code
IntegerMul_ proc
; Calculate a * b * c * d
    movsx rax,cl          ;rax = sign_extend(a)
    movsx rdx,dx          ;rdx = sign_extend(b)
    imul rax,rdx          ;rax = a * b
    movsxd rcx,r8d        ;rcx = sign_extend(c)
    imul rcx,r9           ;rcx = c * d
    imul rax,rcx          ;rax = a * b * c * d

; Calculate e * f * g * h
    movsx rcx,byte ptr [rsp+40] ;rcx = sign_extend(e)
    movsx rdx,word ptr [rsp+48] ;rdx = sign_extend(f)
    imul rcx,rdx          ;rcx = e * f
    movsxd rdx,dword ptr [rsp+56] ;rdx = sign_extend(g)
    imul rdx,qword ptr [rsp+64] ;rdx = g * h
    imul rcx,rdx          ;rcx = e * f * g * h

; Compute the final product
    imul rax,rcx          ;rax = final product

    ret
IntegerMul_ endp

; extern "C" int UnsignedIntegerDiv_(uint8_t a, uint16_t b, uint32_t c, uint64_t d, uint8_t e,
uint16_t f, uint32_t g, uint64_t h, uint64_t* quo, uint64_t* rem);

```

UnsignedIntegerDiv\_ proc

; Calculate a + b + c + d

```
movzx rax, cl      ;rax = zero_extend(a)
movzx rdx, dx      ;rdx = zero_extend(b)
add rax, rdx        ;rax = a + b
mov r8d, r8d       ;r8 = zero_extend(c)
add r8, r9          ;r8 = c + d
add rax, r8         ;rax = a + b + c + d
xor rdx, rdx        ;rdx:rax = a + b + c + d
```

; Calculate e + f + g + h

```
movzx r8, byte ptr [rsp+40] ;r8 = zero_extend(e)
movzx r9, word ptr [rsp+48] ;r9 = zero_extend(f)
add r8, r9          ;r8 = e + f
mov r10d, [rsp+56]   ;r10 = zero_extend(g)
add r10, [rsp+64]    ;r10 = g + h
add r8, r10          ;r8 = e + f + g + h
jnz DivOK           ;jump if divisor is not zero
```

```
xor eax, eax        ;set error return code
jmp done
```

; Calculate (a + b + c + d) / (e + f + g + h)

```
DivOK: div r8        ;unsigned divide rdx:rax / r8
mov rcx, [rsp+72]
mov [rcx], rax       ;save quotient
mov rcx, [rsp+80]
mov [rcx], rdx       ;save remainder

mov eax, 1           ;set success return code
```

Done: ret  
UnsignedIntegerDiv\_ endp  
end

##### Ch02\_06.asm

; Simple lookup table (.const section data is read only)

```
.const
FibVals    dword 0, 1, 1, 2, 3, 5, 8, 13
           dword 21, 34, 55, 89, 144, 233, 377, 610, 987, 1597
```

```
NumFibVals_ dword ($ - FibVals) / sizeof dword
public NumFibVals_
```

; Data section (data is read/write)

```
.data
FibValsSum_ dword ? ;value to demo RIP-relative addressing
public FibValsSum_
```

; extern "C" int MemoryAddressing\_(int i, int\* v1, int\* v2, int\* v3, int\* v4);  
; Returns: 0 = error (invalid table index), 1 = success

.code  
MemoryAddressing\_ proc

; Make sure 'i' is valid

```
cmp ecx, 0
jl InvalidIndex ;jump if i < 0
cmp ecx, [NumFibVals_]
jge InvalidIndex ;jump if i >= NumFibVals_
```

; Sign extend i for use in address calculations

```
movsxd rcx, ecx ;sign extend i
mov [rsp+8], rcx ;save copy of i (in rcx home area)
```

; Example #1 - base register

```
mov r11, offset FibVals ;r11 = FibVals
shl rcx, 2               ;rcx = i * 4
add r11, rcx             ;r11 = FibVals + i * 4
mov eax, [r11]           ;eax = FibVals[i]
mov [rdx], eax           ;save to v1
```

; Example #2 - base register + index register

```
mov r11, offset FibVals ;r11 = FibVals
mov rcx, [rsp+8]         ;rcx = i
shl rcx, 2               ;rcx = i * 4
mov eax, [r11+rcx]       ;eax = FibVals[i]
mov [r8], eax            ;save to v2
```

; Example #3 - base register + index register \* scale factor

```
mov r11, offset FibVals ;r11 = FibVals
mov rcx, [rsp+8]         ;rcx = i
mov eax, [r11+rcx*4]     ;eax = FibVals[i]
mov [r9], eax            ;save to v3
```

; Example #4 - base register + index register \* scale factor + disp

```
mov r11, offset FibVals-42 ;r11 = FibVals - 42
mov rcx, [rsp+8]           ;rcx = i
mov eax, [r11+rcx*4+42]    ;eax = FibVals[i]
mov r10, [rsp+40]          ;r10 = ptr to v4
mov [r10], eax             ;save to v4
```

; Example #5 - RIP relative

```
add [FibValsSum_], eax ;update sum
```

```
mov eax, 1 ;set success return code
ret
```

InvalidIndex:

```
xor eax, eax ;set error return code
ret
```

MemoryAddressing\_ endp

end

##### Ch02\_07.asm

; extern "C" int SignedMinA\_(int a, int b, int c);  
; Returns: min(a, b, c)

.code

```
SignedMinA_ proc
mov eax, ecx
cmp eax, edx ;compare a and b
jle @F       ;jump if a <= b
mov eax, edx ;eax = b
```

```
@@: cmp eax, r8d ;compare min(a, b) and c
     jle @F
     mov eax, r8d ;eax = min(a, b, c)
```

```
@@: ret
SignedMinA_ endp
```

; extern "C" int SignedMaxA\_(int a, int b, int c);  
; Returns: max(a, b, c)

```
SignedMaxA_ proc
mov eax, ecx
cmp eax, edx ;compare a and b
jge @F       ;jump if a >= b
mov eax, edx ;eax = b
```

```
@@: cmp eax, r8d ;compare max(a, b) and c
     jge @F
     mov eax, r8d ;eax = max(a, b, c)
```

```

@@:    ret
SignedMaxA_ endp

; extern "C" int SignedMinB_(int a, int b, int c);
; Returns:    min(a, b, c)
SignedMinB_ proc
    cmp ecx,edx
    cmovg ecx,edx           ;ecx = min(a, b)
    cmp ecx,r8d
    cmovg ecx,r8d           ;ecx = min(a, b, c)
    mov eax,ecx
    ret
SignedMinB_ endp

; extern "C" int SignedMaxB_(int a, int b, int c);
; Returns:    max(a, b, c)
SignedMaxB_ proc
    cmp ecx,edx
    cmovl ecx,edx           ;ecx = max(a, b)
    cmp ecx,r8d
    cmovl ecx,r8d           ;ecx = max(a, b, c)
    mov eax,ecx
    ret
SignedMaxB_ endp
end

;-----
;##### Ch03_01.asm
;-----
; extern "C" int CalcArraySum_(const int* x, int n)
; Returns:    Sum of elements in array x

    .code
CalcArraySum_ proc

; Initialize sum to zero
xor eax, eax                ;sum = 0

; Make sure 'n' is greater than zero
cmp edx, 0
jle InvalidCount           ;jump if n <= 0

; Sum the elements of the array
@@:    add eax,[rcx]         ;add next element to total (sum += *x)
    add rcx, 4              ;set pointer to next element (x++)
    dec edx                 ;adjust counter (n -= 1)
    jnz @B                 ;repeat if not done

InvalidCount:
    ret

CalcArraySum_ endp
end

;-----
;##### Ch03_02.asm
;-----
; extern "C" long long CalcArrayValues_(long long* y, const int* x, int a, short b, int n);
; Calculation: y[i] = x[i] * a + b
; Returns:    Sum of the elements in array y.

    .code
CalcArrayValues_ proc frame

; Function prolog
push rsi                    ;save volatile register rsi
    .pushreg rsi
push rdi                    ;save volatile register rdi
    .pushreg rdi

;-----

```

```

    .endprolog

; Initialize sum to zero and make sure 'n' is valid
xor rax,rax                 ;sum = 0
mov r11d,[rsp+56]           ;r11d = n
cmp r11d, 0
jle InvalidCount           ;jump if n <= 0

; Initialize source and destination pointers
mov rsi,rdx                 ;rsi = ptr to array x
mov rdi,rcx                 ;rdi = ptr to array y

; Load expression constants and array index
movsxd r8,r8d               ;r8 = a (sign extended)
movsx r9,r9w                ;r9 = b (sign extended)
xor edx,edx                 ;edx = array index i

; Repeat until done
@@:    movsxd rcx,dword ptr [rsi+rdx*4] ;rcx = x[i] (sign extended)
    imul rcx,r8              ;rcx = x[i] * a
    add rcx,r9               ;rcx = x[i] * a + b
    mov qword ptr [rdi+rdx*8],rcx ;y[i] = rcx

    add rax,rcx              ;update running sum

    inc edx                 ;edx = i + 1
    cmp edx,r11d            ;is i >= n?
    jl @B                   ;jump if i < n

InvalidCount:

; Function epilog
pop rdi                     ;restore caller's rdi
pop rsi                     ;restore caller's rsi
ret

CalcArrayValues_ endp
end

;-----
;##### Ch03_03.asm
;-----
; void CalcMatrixSquares_(int* y, const int* x, int nrows, int ncols);
; Calculates: y[i][j] = x[j][i] * x[j][i]

    .code
CalcMatrixSquares_ proc frame

; Function prolog
push rsi                    ;save caller's rsi
    .pushreg rsi
push rdi                    ;save caller's rdi
    .pushreg rdi
    .endprolog

; Make sure nrows and ncols are valid
cmp r8d, 0
jle InvalidCount           ;jump if nrows <= 0
cmp r9d, 0
jle InvalidCount           ;jump if ncols <= 0

; Initialize pointers to source and destination arrays
mov rsi,rdx                 ;rsi = x
mov rdi,rcx                 ;rdi = y
xor rcx,rcx                 ;rcx = i
movsxd r8,r8d               ;r8 = nrows sign extended
movsxd r9,r9d               ;r9 = ncols sign extended

; Perform the required calculations

```

```

Loop1:
xor rdx,rdx                ;rdx = j
Loop2:
mov rax,rdx                ;rax = j
imul rax,r9                ;rax = j * ncols
add rax,rcx                ;rax = j * ncols + i
mov r10d,dword ptr [rsi+rax*4] ;r10d = x[j][i]
imul r10d,r10d             ;r10d = x[j][i] * x[j][i]

mov rax,rcx                ;rax = i
imul rax,r9                ;rax = i * ncols
add rax,rdx                ;rax = i * ncols + j;
mov dword ptr [rdi+rax*4],r10d ;y[i][j] = r10d

inc rdx                    ;j += 1
cmp rdx,r9                 ;jump if j < ncols
jl Loop2

inc rcx                    ;i += 1
cmp rcx,r8                 ;jump if i < nrows
jl Loop1

InvalidCount:

; Function epilog
pop rdi                    ;restore caller's rdi
pop rsi                    ;restore caller's rsi
ret

CalcMatrixSquares_ endp
end
;-----
;##### Ch03_04.asm
;-----
; extern "C" int CalcMatrixRowColSums_(int* row_sums, int* col_sums, const int* x, int nrows,
int ncols)
; Returns:      0 = nrows <= 0 or ncols <= 0, 1 = success

.code
CalcMatrixRowColSums_ proc frame

; Function prolog
push rbx                    ;save caller's rbx
pushreg rbx
push rsi                    ;save caller's rsi
pushreg rsi
push rdi                    ;save caller's rdi
pushreg rdi
endprolog

; Make sure nrows and ncols are valid
xor eax,eax                ;set error return code

cmp r9d,0
jle InvalidArg             ;jump if nrows <= 0

mov r10d,[rsp+64]           ;r10d = ncols
cmp r10d,0

jle InvalidArg             ;jump if ncols <= 0

; Initialize elements of col_sums array to zero
mov rbx,rcx                ;temp save of row_sums
mov rdi,rdx                ;rdi = col_sums
mov ecx,r10d               ;rcx = ncols
xor eax,eax                ;eax = fill value
rep stosd                  ;fill array with zeros

```

```

; The code below uses the following registers:
; rcx = row_sums           rdx = col_sums
; r9d = nrows              r10d = ncols
; eax = i                  ebx = j
; edi = i * ncols          esi = i * ncols + j
; r8 = x                   r11d = x[i][j]

; Initialize outer loop variables.
mov rcx,rbx                ;rcx = row_sums
xor eax,eax                ;i = 0

Lp1: mov dword ptr [rcx+rax*4],0 ;row_sums[i] = 0
xor ebx,ebx                ;j = 0
mov edi,eax                ;edi = i
imul edi,r10d              ;edi = i * ncols

; Inner loop
Lp2: mov esi,edi            ;esi = i * ncols
add esi,ebx                ;esi = i * ncols + j
mov r11d,[r8+rsi*4]         ;r11d = x[i * ncols + j]
add [rcx+rax*4],r11d        ;row_sums[i] += x[i * ncols + j]
add [rdx+rbx*4],r11d        ;col_sums[j] += x[i * ncols + j]

; Is the inner loop finished?
inc ebx                    ;j += 1
cmp ebx,r10d               ;jump if j < ncols
jl Lp2

; Is the outer loop finished?
inc eax                    ;i += 1
cmp eax,r9d                ;jump if i < nrows
jl Lp1

mov eax,1                  ;set success return code

; Function epilog
InvalidArg:
pop rdi                    ;restore NV registers and return
pop rsi
pop rbx
ret

CalcMatrixRowColSums_ endp
end
;-----
;##### Ch03_05.asm
;-----
TestStruct struct
Val8 byte ?
Pad8 byte ?
Val16 word ?
Val32 dword ?
Val64 qword ?
TestStruct ends

; extern "C" int64_t CalcTestStructSum_(const TestStruct* ts);
; Returns:      Sum of structure's values as a 64-bit integer.

.code
CalcTestStructSum_ proc

; Compute ts->Val8 + ts->Val16, note sign extension to 32-bits
movsx eax,byte ptr [rcx+TestStruct.Val8]
movsx edx,word ptr [rcx+TestStruct.Val16]
add eax,edx

; Sign extend previous result to 64 bits

```



```

movsxd rax, eax

; Add ts->Val32 to sum
movsxd rdx, [rcx+TestStruct.Val32]
add rax, rdx

; Add ts->Val64 to sum
add rax, [rcx+TestStruct.Val64]
ret

CalcTestStructSum_ endp
end
-----
;##### Ch03_06.asm
;-----
; extern "C" unsigned long long CountChars_(const char* s, char c);
; Description: This function counts the number of occurrences
;              of a character in a string.
; Returns:     Number of occurrences found.

.code
CountChars_ proc frame

; Save non-volatile registers
push rsi                      ;save caller's rsi
.pushreg rsi
.endprolog

; Load parameters and initialize count registers
mov rsi, rcx                  ;rsi = s
mov cl, dl                    ;cl = c
xor edx, edx                  ;rdx = Number of occurrences
xor r8d, r8d                  ;r8 = 0 (required for add below)

; Repeat loop until the entire string has been scanned
@@: lodsb                     ;load next char into register al
or al, al                     ;test for end-of-string
jz @F                          ;jump if end-of-string found
cmp al, cl                    ;test current char
sete r8b                      ;r8b = 1 if match, 0 otherwise
add rdx, r8                   ;update occurrence count
jmp @B

@@: mov rax, rdx               ;rax = number of occurrences

; Restore non-volatile registers and return
pop rsi
ret

CountChars_ endp
end
-----
;##### Ch03_07.asm
;-----
; extern "C" size_t ConcatStrings_(char* des, size_t des_size, const char* const* src, size_t
; src_n);
; Returns:     -1 Invalid 'des_size'
;              n >= 0 Length of concatenated string

.code
ConcatStrings_ proc frame

; Save non-volatile registers
push rbx
.pushreg rbx
push rsi
.pushreg rsi
push rdi
.pushreg rdi

```

```

.endprolog

; Make sure des_size and src_n are valid
mov rax, -1                   ;set error code

test rdx, rdx                 ;test des_size
jz InvalidArg                 ;jump if des_size is 0

test r9, r9                   ;test src_n
jz InvalidArg                 ;jump if src_n is 0

; Registers used processing loop below
; rax = des                    rdx = des_size
; r8 = src                     r9 = src_n
; r10 = des_index              r11 = i
; rcx = string length
; rsi, rdi = pointers for scasb & movsb instructions

; Perform required initializations
xor r10, r10                  ;des_index = 0
xor r11, r11                  ;i = 0
mov rbx, rcx                  ;rbx = des
mov byte ptr [rbx], 0         ;*des = '\0'

; Repeat loop until concatenation is finished
Loop1: mov rax, r8              ;rax = 'src'
mov rdi, [rax+r11*8]          ;rdi = src[i]
mov rsi, rdi                  ;rsi = src[i]

; Compute length of s[i]
xor eax, eax
mov rcx, -1
repne scasb                   ;find '\0'
not rcx                       ;rcx = len(src[i])

; Compute des_index + src_len
mov rax, r10                  ;rax = des_index
add rax, rcx                   ;des_index + len(src[i])
cmp rax, rdx                   ;is des_index + src_len >= des_size?
jge Done                      ;jump if des is too small

; Update des_index
mov rax, r10                  ;des_index_old = des_index
add r10, rcx                  ;des_index += len(src[i])

; Copy src[i] to &des[des_index] (rsi already contains src[i])
inc rcx                       ;rcx = len(src[i]) + 1
lea rdi, [rbx+rax]            ;rdi = &des[des_index_old]
rep movsb                     ;perform string move

; Update i and repeat if not done
inc r11                       ;i += 1
cmp r11, r9                   ;jump if i < src_n
jl Loop1

; Return length of concatenated string

Done: mov rax, r10              ;rax = des_index (final length)
; Restore non-volatile registers and return

InvalidArg:
pop rdi
pop rsi
pop rbx
ret

ConcatStrings_ endp
end
-----

```

```

:##### Ch03_08.asm
:-----
: extern "C" long CompareArrays_(const int* x, const int* y, long long n)
: Returns      -1      Value of 'n' is invalid
:              0 <= i < n  Index of first non-matching element
:              n          All elements match

```

```

.code
CompareArrays_ proc frame

```

```

; Save non-volatile registers
    push rsi
    .pushreg rsi
    push rdi
    .pushreg rdi
    .endprolog

```

```

; Load arguments and validate 'n'
    mov rax, -1          ;rax = return code for invalid n
    test r8, r8
    jle @F              ;jump if n <= 0

```

```

; Compare the arrays for equality
    mov rsi, rcx          ;rsi = x
    mov rdi, rdx          ;rdi = y
    mov rcx, r8           ;rcx = n
    mov rax, r8           ;rax = n
    repe cmpsd
    je @F                ;arrays are equal

```

```

; Calculate index of first non-match
    sub rax, rcx          ;rax = index of mismatch + 1
    dec rax              ;rax = index of mismatch

```

```

; Restore non-volatile registers and return

```

```

@@:    pop rdi
        pop rsi
        ret

```

```

CompareArrays_ endp
end

```

```

:##### Ch03_09.asm
:-----

```

```

: extern "C" int ReverseArray_(int* y, const int* x, int n);
: Returns      0 = invalid n, 1 = success

```

```

.code
ReverseArray_ proc frame

```

```

; Save non-volatile registers
    push rsi
    .pushreg rsi
    push rdi
    .pushreg rdi
    .endprolog

```

```

; Make sure n is valid
    xor eax, eax          ;error return code
    test r8d, r8d         ;is n <= 0?
    jle InvalidArg        ;jump if n <= 0

```

```

; Initialize registers for reversal operation
    mov rsi, rdx          ;rsi = x
    mov rdi, rcx          ;rdi = y
    mov ecx, r8d          ;rcx = n
    lea rsi, [rsi+rcx*4-4] ;rsi = &x[n - 1]

```

```

; Save caller's RFLAGS.DF, then set RFLAGS.DF to 1
    pushfq                ;save caller's RFLAGS.DF
    std                   ;RFLAGS.DF = 1

```

```

; Repeat loop until array reversal is complete
@@:    lodsd                ;eax = *x--
        mov [rdi], eax      ;*y = eax
        add rdi, 4          ;y++
        dec rcx             ;n--
        jnz @B

```

```

; Restore caller's RFLAGS.DF and set return code
    popfq                 ;restore caller's RFLAGS.DF
    mov eax, 1            ;set success return code

```

```

; Restore non-volatile registers and return

```

```

InvalidArg:

```

```

    pop rdi
    pop rsi
    ret

```

```

ReverseArray_ endp
end

```

```

:##### Ch05_01.asm
:-----

```

```

    .const
r4_ScaleFtoC real4 0.55555556      ; 5 / 9
r4_ScaleCtoF real4 1.8             ; 9 / 5
r4_32p0      real4 32.0

```

```

; extern "C" float ConvertFtoC_(float deg_f)
; Returns: xmm0[31:0] = temperature in Celsius.

```

```

.code
ConvertFtoC_ proc
    vmovss xmm1, [r4_32p0]          ;xmm1 = 32
    vsubss xmm2, xmm0, xmm1         ;xmm2 = f - 32

    vmovss xmm1, [r4_ScaleFtoC]     ;xmm1 = 5 / 9
    vmulss xmm0, xmm2, xmm1         ;xmm0 = (f - 32) * 5 / 9
    ret
ConvertFtoC_ endp

```

```

; extern "C" float CtoF_(float deg_c)
; Returns: xmm0[31:0] = temperature in Fahrenheit.

```

```

ConvertCtoF_ proc
    vmulss xmm0, xmm0, [r4_ScaleCtoF] ;xmm0 = c * 9 / 5
    vaddss xmm0, xmm0, [r4_32p0]      ;xmm0 = c * 9 / 5 + 32
    ret
ConvertCtoF_ endp
end

```

```

:##### Ch05_02.asm
:-----

```

```

    .const
r8_PI real8 3.14159265358979323846
r8_4p0 real8 4.0
r8_3p0 real8 3.0

```

```

; extern "C" void CalcSphereAreaVolume_(double r, double* sa, double* vol);

```

```

.code
CalcSphereAreaVolume_ proc

```

```

; Calculate surface area = 4 * PI * r * r
    vmulsd xmm1, xmm0, xmm0         ;xmm1 = r * r

```

```

        vmulsd xmm2, xmm1, [r8_P1]      ;xmm2 = r * r * PI
        vmulsd xmm3, xmm2, [r8_4p0]    ;xmm3 = r * r * PI * 4

; Calculate volume = sa * r / 3
        vmulsd xmm4, xmm3, xmm0      ;xmm4 = r * r * r * PI * 4
        vdivsd xmm5, xmm4, [r8_3p0]   ;xmm5 = r * r * r * PI * 4 / 3

; Save results
        vmovsd real8 ptr [rdx], xmm3  ;save surface area
        vmovsd real8 ptr [r8], xmm5   ;save volume
        ret

CalcSphereAreaVolume_ endp
end

;-----
;##### Ch05_03.asm
;-----
; extern "C" double CalcDistance_(double x1, double y1, double z1, double x2, double y2, double
z2)
        .code
CalcDistance_ proc
; Load arguments from stack
        vmovsd xmm4, real8 ptr [rsp+40] ;xmm4 = y2
        vmovsd xmm5, real8 ptr [rsp+48] ;xmm5 = z2

; Calculate squares of coordinate distances
        vsubsd xmm0, xmm3, xmm0      ;xmm0 = x2 - x1
        vmulsd xmm0, xmm0, xmm0      ;xmm0 = (x2 - x1) * (x2 - x1)
        vsubsd xmm1, xmm4, xmm1      ;xmm1 = y2 - y1
        vmulsd xmm1, xmm1, xmm1      ;xmm1 = (y2 - y1) * (y2 - y1)

        vsubsd xmm2, xmm5, xmm2      ;xmm2 = z2 - z1
        vmulsd xmm2, xmm2, xmm2      ;xmm2 = (z2 - z1) * (z2 - z1)

; Calculate final distance
        vaddsd xmm3, xmm0, xmm1
        vaddsd xmm4, xmm2, xmm3      ;xmm4 = sum of squares
        vsqrtsd xmm0, xmm0, xmm4     ;xmm0 = final distance value
        ret
CalcDistance_ endp
end

;-----
;##### Ch05_04.asm
;-----
; extern "C" void CompareVCOMISS_(float a, float b, bool* results);

        .code
CompareVCOMISS_ proc

; Set result flags based on compare status
        vcomiss xmm0, xmm1
        setp byte ptr [r8]          ;RFLAGS.PF = 1 if unordered
        jnp @F
        xor al, al
        mov byte ptr [r8+1], al     ;Use default result values
        mov byte ptr [r8+2], al
        mov byte ptr [r8+3], al
        mov byte ptr [r8+4], al
        mov byte ptr [r8+5], al
        mov byte ptr [r8+6], al
        jmp Done

@@:    setb byte ptr [r8+1]          ;set byte if a < b
        setbe byte ptr [r8+2]       ;set byte if a <= b
        sete byte ptr [r8+3]        ;set byte if a == b
        setne byte ptr [r8+4]       ;set byte if a != b
        seta byte ptr [r8+5]        ;set byte if a > b
        setae byte ptr [r8+6]       ;set byte if a >= b

Done:  ret
CompareVCOMISS_ endp
end

```

```

Done:  ret
CompareVCOMISS_ endp

; extern "C" void CompareVCOMISD_(double a, double b, bool* results);

CompareVCOMISD_ proc

; Set result flags based on compare status
        vcomisd xmm0, xmm1
        setp byte ptr [r8]          ;RFLAGS.PF = 1 if unordered
        jnp @F
        xor al, al
        mov byte ptr [r8+1], al     ;Use default result values
        mov byte ptr [r8+2], al
        mov byte ptr [r8+3], al
        mov byte ptr [r8+4], al
        mov byte ptr [r8+5], al
        mov byte ptr [r8+6], al
        jmp Done

@@:    setb byte ptr [r8+1]          ;set byte if a < b
        setbe byte ptr [r8+2]       ;set byte if a <= b
        sete byte ptr [r8+3]        ;set byte if a == b
        setne byte ptr [r8+4]       ;set byte if a != b
        seta byte ptr [r8+5]        ;set byte if a > b
        setae byte ptr [r8+6]       ;set byte if a >= b

Done:  ret
CompareVCOMISD_ endp
end

;-----
;##### Ch05_05.asm
;-----
; include <cmpequ.asmh>

; extern "C" void CompareVCMPSD_(double a, double b, bool* results)
        .code
CompareVCMPSD_ proc

; Perform compare for equality
        vcmppsd xmm2, xmm0, xmm1, CMP_EQ ;perform compare operation
        vmovq rax, xmm2                 ;rax = compare result (all 1s or 0s)
        and al, 1                       ;mask out unneeded bits
        mov byte ptr [r8], al           ;save result as C++ bool

; Perform compare for inequality
        vcmppsd xmm2, xmm0, xmm1, CMP_NEQ
        vmovq rax, xmm2
        and al, 1
        mov byte ptr [r8+1], al

; Perform compare for less than
        vcmppsd xmm2, xmm0, xmm1, CMP_LT
        vmovq rax, xmm2
        and al, 1
        mov byte ptr [r8+2], al

; Perform compare for less than or equal
        vcmppsd xmm2, xmm0, xmm1, CMP_LE
        vmovq rax, xmm2
        and al, 1
        mov byte ptr [r8+3], al

; Perform compare for greater than
        vcmppsd xmm2, xmm0, xmm1, CMP_GT
        vmovq rax, xmm2
        and al, 1
        mov byte ptr [r8+4], al

```

```

; Perform compare for greater than or equal
vcmpsdxmm2, xmm0, xmm1, CMP_GE
vmovq rax, xmm2
and al, 1
mov byte ptr [r8+5], al

; Perform compare for ordered
vcmpsdxmm2, xmm0, xmm1, CMP_ORD
vmovq rax, xmm2
and al, 1
mov byte ptr [r8+6], al

; Perform compare for unordered
vcmpsdxmm2, xmm0, xmm1, CMP_UNORD
vmovq rax, xmm2
and al, 1
mov byte ptr [r8+7], al

ret
CompareVCMPD_ endp
end
;##### Ch05_06.asm
;-----
MxcscrRcMask equ 9ffffh ;bit pattern for MXCSR.RC
MxcscrRcShift equ 13 ;shift count for MXCSR.RC

; extern "C" RoundingMode GetMxcscrRoundingMode(void);
; Description: The following function obtains the current
; floating-point rounding mode from MXCSR.RC.
; Returns: Current MXCSR.RC rounding mode.

.code
GetMxcscrRoundingMode_ proc
vstmxcscr dword ptr [rsp+8] ;save mxcscr register
mov eax, [rsp+8]
shr eax, MxcscrRcShift ;eax[1:0] = MXCSR.RC bits
and eax, 3 ;masked out unwanted bits
ret
GetMxcscrRoundingMode_ endp

;extern "C" void SetMxcscrRoundingMode(RoundingMode rm);
; Description: The following function updates the rounding mode
; value in MXCSR.RC.

SetMxcscrRoundingMode_ proc
and ecx, 3 ;masked out unwanted bits
shl ecx, MxcscrRcShift ;ecx[14:13] = rm

vstmxcscr dword ptr [rsp+8] ;save current MXCSR
mov eax, [rsp+8]
and eax, MxcscrRcMask ;masked out old MXCSR.RC bits
or eax, ecx ;insert new MXCSR.RC bits
mov [rsp+8], eax
vldmxcscr dword ptr [rsp+8] ;load updated MXCSR
ret
SetMxcscrRoundingMode_ endp

; extern "C" bool ConvertScalar(Uval* des, const Uval* src, CvtOp cvt_op)
; Note: This function requires linker option /LARGEADDRESSAWARE:NO
; to be explicitly set.

ConvertScalar_ proc

; Make sure cvt_op is valid, then jump to target conversion code
mov eax, r8d ;eax = CvtOp

```

```

cmp eax, CvtOpTableCount
jae BadCvtOp ;jump if cvt_op is invalid
jmp [CvtOpTable+rax*8] ;jump to specified conversion

; Conversions between int32_t and float/double

I32_F32:
mov eax, [rdx] ;load integer value
vcvtsi2ss xmm0, xmm0, eax ;convert to float
vmovss real4 ptr [rcx], xmm0 ;save result
mov eax, 1
ret

F32_I32:
vmovss xmm0, real4 ptr [rdx] ;load float value
vcvtss2si eax, xmm0 ;convert to integer
mov [rcx], eax ;save result
mov eax, 1
ret

I32_F64:
mov eax, [rdx] ;load integer value
vcvtsi2sd xmm0, xmm0, eax ;convert to double
vmovsd real8 ptr [rcx], xmm0 ;save result
mov eax, 1
ret

F64_I32:
vmovsd xmm0, real8 ptr [rdx] ;load double value
vcvtsd2si eax, xmm0 ;convert to integer
mov [rcx], eax ;save result
mov eax, 1
ret

; Conversions between int64_t and float/double

I64_F32:
mov rax, [rdx] ;load integer value
vcvtsi2ss xmm0, xmm0, rax ;convert to float
vmovss real4 ptr [rcx], xmm0 ;save result
mov eax, 1
ret

F32_I64:
vmovss xmm0, real4 ptr [rdx] ;load float value
vcvtss2si rax, xmm0 ;convert to integer
mov [rcx], rax ;save result
mov eax, 1
ret

I64_F64:
mov rax, [rdx] ;load integer value
vcvtsi2sd xmm0, xmm0, rax ;convert to double
vmovsd real8 ptr [rcx], xmm0 ;save result
mov eax, 1
ret

F64_I64:
vmovsd xmm0, real8 ptr [rdx] ;load double value
vcvtsd2si rax, xmm0 ;convert to integer
mov [rcx], rax ;save result
mov eax, 1
ret

; Conversions between float and double

F32_F64:

```

```

        vmovss xmm0, real4 ptr [rdx]      ;load float value
        vcvts2sd xmm1, xmm1, xmm0        ;convert to double
        vmovsd real8 ptr [rcx], xmm1      ;save result
        mov eax, 1
        ret

F64_F32:
        vmovsd xmm0, real8 ptr [rdx]      ;load double value
        vcvtsd2ss xmm1, xmm1, xmm0        ;convert to float
        vmovss real4 ptr [rcx], xmm1      ;save result
        mov eax, 1
        ret

BadCvtOp:
        xor eax, eax                      ;set error return code
        ret

; The order of values in following table must match the enum CvtOp
; that's defined in the .cpp file.

        align 8
CvtOpTable equ $
        qword 132_F32, F32_I32
        qword 132_F64, F64_I32
        qword 164_F32, F32_I64
        qword 164_F64, F64_I64
        qword F32_F64, F64_F32
CvtOpTableCount equ ($ - CvtOpTable) / size qword

ConvertScalar_ endp
        end
;-----
;##### Ch05_07.asm
;-----
; extern "C" bool CalcMeanStdev(double* mean, double* stdev, const double* a, int n);
; Returns:      0 = invalid n, 1 = valid n

        .code
CalcMeanStdev_ proc

; Make sure 'n' is valid
        xor eax, eax                      ;set error return code (also i = 0)
        cmp r9d, 2
        jl InvalidArg                     ;jump if n < 2

; Compute sample mean
        vxorpd xmm0, xmm0, xmm0            ;sum = 0.0

@@:      vaddsd xmm0, xmm0, real8 ptr [r8+rax*8] ;sum += x[i]
        inc eax                            ;i += 1
        cmp eax, r9d
        jl @B                              ;jump if i < n

        vcvtsi2sd xmm1, xmm1, r9d          ;convert n to DPFP
        vdivsd xmm3, xmm0, xmm1            ;xmm3 = mean (sum / n)
        vmovsd real8 ptr [rcx], xmm3      ;save mean

; Compute sample stdev
        xor eax, eax                      ;i = 0
        vxorpd xmm0, xmm0, xmm0            ;sum2 = 0.0

@@:      vmovsd xmm1, real8 ptr [r8+rax*8] ;xmm1 = x[i]
        vsubsd xmm2, xmm1, xmm3            ;xmm2 = x[i] - mean
        vmulsd xmm2, xmm2, xmm2            ;xmm2 = (x[i] - mean) ** 2
        vaddsd xmm0, xmm0, xmm2            ;sum2 += (x[i] - mean) ** 2
        inc eax                            ;i += 1
        cmp eax, r9d

```

```

        jl @B                              ;jump if i < n

        dec r9d                            ;r9d = n - 1
        vcvtsi2sd xmm1, xmm1, r9d          ;convert n - 1 to DPFP
        vdivsd xmm0, xmm0, xmm1            ;xmm0 = sum2 / (n - 1)
        vsqrtsd xmm0, xmm0, xmm0           ;xmm0 = stdev
        vmovsd real8 ptr [rdx], xmm0      ;save stdev

        mov eax, 1                          ;set success return code

InvalidArg:
        ret
CalcMeanStdev_ endp
        end
;-----
;##### Ch05_08.asm
;-----
; void CalcMatrixSquaresF32_(float* y, const float* x, float offset, int nrows, int ncols);
; Calculates:   y[i][j] = x[j][i] * x[j][i] + offset

        .code
CalcMatrixSquaresF32_ proc frame

; Function prolog
        push rsi                          ;save caller's rsi
        .pushreg rsi
        push rdi                          ;save caller's rdi
        .pushreg rdi
        .endprolog

; Make sure nrows and ncols are valid
        movsxd r9, r9d                     ;r9 = nrows
        test r9, r9
        jle InvalidCount                  ;jump if nrows <= 0

        movsxd r10, dword ptr [rsp+56]     ;r10 = ncols
        test r10, r10
        jle InvalidCount                  ;jump if ncols <= 0

; Initialize pointers to source and destination arrays
        mov rsi, rdx                       ;rsi = x
        mov rdi, rcx                       ;rdi = y
        xor rcx, rcx                       ;rcx = i

; Perform the required calculations
Loop1:   xor rdx, rdx                       ;rdx = j

Loop2:   mov rax, rdx                       ;rax = j
        imul rax, r10                      ;rax = j * ncols
        add rax, rcx                       ;rax = j * ncols + i
        vmovss xmm0, real4 ptr [rsi+rax*4] ;xmm0 = x[j][i]
        vmulss xmm1, xmm0, xmm0            ;xmm1 = x[j][i] * x[j][i]
        vaddss xmm3, xmm1, xmm2            ;xmm2 = x[j][i] * x[j][i] + offset

        mov rax, rcx                       ;rax = i
        imul rax, r10                      ;rax = i * ncols
        add rax, rdx                       ;rax = i * ncols + j;
        vmovss real4 ptr [rdi+rax*4], xmm3 ;y[i][j] = x[j][i] * x[j][i] + offset

        inc rdx                            ;j += 1
        cmp rdx, r10
        jl Loop2                          ;jump if j < ncols

        inc rcx                            ;i += 1
        cmp rcx, r9
        jl Loop1                          ;jump if i < nrows

```

```

InvalidCount:

; Function epilog
    pop rdi                ;restore caller's rdi
    pop rsi                ;restore caller's rsi
    ret

CalcMatrixSquaresF32_ endp
    end

;-----
;##### Ch05_09.asm
;-----
; extern "C" Int64 Cc1_(int8_t a, int16_t b, int32_t c, int64_t d, int8_t e, int16_t f, int32_t g, int64_t h);

    .code
Cc1_  proc frame

; Function prolog
    push rbp                ;save caller's rbp register
    .pushreg rbp

    sub rsp,16              ;allocate local stack space
    .allocstack 16

    mov rbp, rsp            ;set frame pointer
    .setframe rbp,0

RBP_RA = 24                 ;offset from rbp to return addr
    .endprolog              ;mark end of prolog

; Save argument registers to home area (optional)
    mov [rbp+RBP_RA+8],rcx
    mov [rbp+RBP_RA+16],rdx
    mov [rbp+RBP_RA+24],r8
    mov [rbp+RBP_RA+32],r9

; Sum the argument values a, b, c, and d
    movsx rcx,cl            ;rcx = a
    movsx rdx,dx            ;rdx = b
    movsxd r8,r8d           ;r8 = c;
    add rcx,rdx             ;rcx = a + b
    add r8,r9               ;r8 = c + d
    add r8,rcx              ;r8 = a + b + c + d
    mov [rbp],r8            ;save a + b + c + d

; Sum the argument values e, f, g, and h
    movsx rcx,byte ptr [rbp+RBP_RA+40] ;rcx = e
    movsx rdx,word ptr [rbp+RBP_RA+48] ;rdx = f
    movsxd r8,dword ptr [rbp+RBP_RA+56] ;r8 = g
    add rcx,rdx             ;rcx = e + f
    add r8,qword ptr [rbp+RBP_RA+64] ;r8 = g + h
    add r8,rcx              ;r8 = e + f + g + h

; Compute the final sum
    mov rax,[rbp]           ;rax = a + b + c + d
    add rax,r8              ;rax = final sum

; Function epilog
    add rsp,16              ;release local stack space
    pop rbp                ;restore caller's rbp register
    ret

Cc1_ endp
    end

;-----
;##### Ch05_10.asm
;-----

```

```

; extern "C" void Cc2_(const int64_t* a, const int64_t* b, int32_t n, int64_t* sum_a, int64_t*
sum_b, int64_t* prod_a, int64_t* prod_b)
; Named expressions for constant values:
; NUM_PUSHREG = number of prolog non-volatile register pushes
; STK_LOCAL1 = size in bytes of STK_LOCAL1 area (see figure in text)
; STK_LOCAL2 = size in bytes of STK_LOCAL2 area (see figure in text)
; STK_PAD = extra bytes (0 or 8) needed to 16-byte align RSP
; STK_TOTAL = total size in bytes of local stack
; RBP_RA = number of bytes between RBP and ret addr on stack

NUM_PUSHREG = 4
STK_LOCAL1 = 32
STK_LOCAL2 = 16
STK_PAD = ((NUM_PUSHREG AND 1) XOR 1) * 8
STK_TOTAL = STK_LOCAL1 + STK_LOCAL2 + STK_PAD
RBP_RA = NUM_PUSHREG * 8 + STK_LOCAL1 + STK_PAD

    .const
TestVal db 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15

    .code
Cc2_  proc frame

; Save non-volatile GP registers on the stack
    push rbp
    .pushreg rbp
    push rbx
    .pushreg rbx
    push r12
    .pushreg r12
    push r13
    .pushreg r13

; Allocate local stack space and set frame pointer
    sub rsp,STK_TOTAL        ;allocate local stack space
    .allocstack STK_TOTAL

    lea rbp,[rsp+STK_LOCAL2] ;set frame pointer
    .setframe rbp,STK_LOCAL2

    .endprolog              ;end of prolog

; Initialize local variables on the stack (demonstration only)
    vmovdqu xmm5, xmmword ptr [TestVal]
    vmovdqa xmmword ptr [rbp-16],xmm5 ;save xmm5 to LocalVar2A/2B
    mov qword ptr [rbp],0aah        ;save 0xaa to LocalVar1A
    mov qword ptr [rbp+8],0bbh      ;save 0xbb to LocalVar1B
    mov qword ptr [rbp+16],0cch     ;save 0xcc to LocalVar1C
    mov qword ptr [rbp+24],0ddh     ;save 0xdd to LocalVar1D

; Save argument values to home area (optional)
    mov qword ptr [rbp+RBP_RA+8],rcx
    mov qword ptr [rbp+RBP_RA+16],rdx
    mov qword ptr [rbp+RBP_RA+24],r8
    mov qword ptr [rbp+RBP_RA+32],r9

; Perform required initializations for processing loop
    test r8d,r8d             ;is n <= 0?
    jle Error                ;jump if n <= 0

    xor rbx,rbx              ;rbx = current element offset
    xor r10,r10              ;r10 = sum_a
    xor r11,r11              ;r11 = sum_b
    mov r12,1                ;r12 = prod_a
    mov r13,1                ;r13 = prod_b

; Compute the array sums and products

```



```

@@:    mov rax,[rcx+rbx]          ;rax = a[i]
        add r10,rax              ;update sum_a
        imul r12,rax             ;update prod_a
        mov rax,[rdx+rbx]       ;rax = b[i]
        add r11,rax             ;update sum_b
        imul r13,rax            ;update prod_b

        add rbx,8               ;set ebx to next element
        dec r8d                 ;adjust count
        jnz @B                 ;repeat until done

; Save the final results
        mov [r9],r10            ;save sum_a
        mov rax,[rbp+RBP_RA+40] ;rax = ptr to sum_b
        mov [rax],r11           ;save sum_b
        mov rax,[rbp+RBP_RA+48] ;rax = ptr to prod_a
        mov [rax],r12          ;save prod_a
        mov rax,[rbp+RBP_RA+56] ;rax = ptr to prod_b
        mov [rax],r13          ;save prod_b
        mov eax,1              ;set return code to true

; Function epilog
Done:   lea rsp,[rbp+STK_LOCAL1+STK_PAD] ;restore rsp
        pop r13                ;restore non-volatile GP registers
        pop r12
        pop rbx
        pop rbp
        ret

Error:  xor eax,eax             ;set return code to false
        jmp Done
Cc2_   endp
        end

```

##### Ch05\_11.asm

```

; extern "C" bool Cc3(const double* r, const double* h, int n, double* sa_cone, double* vol_cone)
; Named expressions for constant values
; NUM_PUSHREG = number of prolog non-volatile register pushes
; STK_LOCAL1  = size in bytes of STK_LOCAL1 area (see figure in text)
; STK_LOCAL2  = size in bytes of STK_LOCAL2 area (see figure in text)
; STK_PAD     = extra bytes (0 or 8) needed to 16-byte align RSP
; STK_TOTAL   = total size in bytes of local stack
; RBP_RA      = number of bytes between RBP and ret addr on stack

```

```

NUM_PUSHREG = 7
STK_LOCAL1  = 16
STK_LOCAL2  = 64
STK_PAD     = ((NUM_PUSHREG AND 1) XOR 1) * 8
STK_TOTAL   = STK_LOCAL1 + STK_LOCAL2 + STK_PAD
RBP_RA      = NUM_PUSHREG * 8 + STK_LOCAL1 + STK_PAD

```

```

        .const
r8_3p0   real8 3.0
r8_pi    real8 3.14159265358979323846

```

```

        .code
Cc3_   proc frame

```

```

; Save non-volatile registers on the stack.
        push rbp
        .pushreg rbp
        push rbx
        .pushreg rbx
        push rsi
        .pushreg rsi
        push r12

```

```

        .pushreg r12
        push r13
        .pushreg r13
        push r14
        .pushreg r14
        push r15
        .pushreg r15

; Allocate local stack space and initialize frame pointer
        sub rsp,STK_TOTAL          ;allocate local stack space
        .allocstack STK_TOTAL
        lea rbp,[rsp+STK_LOCAL2]    ;rbp = stack frame pointer
        .setframe rbp,STK_LOCAL2

; Save non-volatile registers XMM12 - XMM15. Note that STK_LOCAL2 must
; be greater than or equal to the number of XMM register saves times 16.
        vmovdqa xmmword ptr [rbp-STK_LOCAL2+48],xmm12
        .savexmm128 xmm12,48
        vmovdqa xmmword ptr [rbp-STK_LOCAL2+32],xmm13
        .savexmm128 xmm13,32
        vmovdqa xmmword ptr [rbp-STK_LOCAL2+16],xmm14
        .savexmm128 xmm14,16
        vmovdqa xmmword ptr [rbp-STK_LOCAL2],xmm15
        .savexmm128 xmm15,0
        .endprolog

; Access local variables on the stack (demonstration only)
        mov qword ptr [rbp],-1      ;LocalVar1A = -1
        mov qword ptr [rbp+8],-2    ;LocalVar1B = -2

; Initialize the processing loop variables. Note that many of the
; register initializations below are performed merely to illustrate
; use of the non-volatile GP and XMM registers.
        mov esi,r8d                ;esi = n
        test esi,esi               ;is n > 0?
        jg @F                      ;jump if n > 0

        xor eax,eax                ;set error return code
        jmp done

@@:   xor rbx,rbx                   ;rbx = array element offset
        mov r12,rcx                ;r12 = ptr to r
        mov r13,rdx                ;r13 = ptr to h
        mov r14,r9                 ;r14 = ptr to sa_cone
        mov r15,[rbp+RBP_RA+40]    ;r15 = ptr to vol_cone
        vmovsd xmm14,real8 ptr [r8_pi] ;xmm14 = pi
        vmovsd xmm15,real8 ptr [r8_3p0] ;xmm15 = 3.0

; Calculate cone surface areas and volumes
; sa = pi * r * (r + sqrt(r * r + h * h))
; vol = pi * r * r * h / 3
@@:   vmovsd xmm0,real8 ptr [r12+rbx] ;xmm0 = r
        vmovsd xmm1,real8 ptr [r13+rbx] ;xmm1 = h
        vmovsd xmm12,xmm12,xmm0      ;xmm12 = r
        vmovsd xmm13,xmm13,xmm1      ;xmm13 = h

        vmulsd xmm0,xmm0,xmm0        ;xmm0 = r * r
        vmulsd xmm1,xmm1,xmm1        ;xmm1 = h * h
        vaddsd xmm0,xmm0,xmm1        ;xmm0 = r * r + h * h

        vsqrtsd xmm0,xmm0,xmm0       ;xmm0 = sqrt(r * r + h * h)
        vaddsd xmm0,xmm0,xmm12       ;xmm0 = r + sqrt(r * r + h * h)
        vmulsd xmm0,xmm0,xmm12       ;xmm0 = r * (r + sqrt(r * r + h * h))
        vmulsd xmm0,xmm0,xmm14       ;xmm0 = pi * r * (r + sqrt(r * r + h * h))
        vmulsd xmm12,xmm12,xmm12     ;xmm12 = r * r
        vmulsd xmm13,xmm13,xmm14     ;xmm13 = h * pi
        vmulsd xmm13,xmm13,xmm12     ;xmm13 = pi * r * r * h

```

```

        vdivsd xmm13,xmm13,xmm15    ;xmm13 = pi * r * r * h / 3

        vmovsd real8 ptr [r14+rbx],xmm0 ;save surface area
        vmovsd real8 ptr [r15+rbx],xmm13 ;save volume

        add rbx,8                    ;set rbx to next element
        dec esi                      ;update counter
        jnz @B                      ;repeat until done

        mov eax,1                    ;set success return code

; Restore non-volatile XMM registers
Done:   vmovdqa xmm12,xmmword ptr [rbp-STK_LOCAL2+48]
        vmovdqa xmm13,xmmword ptr [rbp-STK_LOCAL2+32]
        vmovdqa xmm14,xmmword ptr [rbp-STK_LOCAL2+16]
        vmovdqa xmm15,xmmword ptr [rbp-STK_LOCAL2]

; Function epilog
        lea rsp,[rbp+STK_LOCAL1+STK_PAD] ;restore rsp
        pop r15                      ;restore NV GP registers
        pop r14
        pop r13
        pop r12
        pop rsi
        pop rbx
        pop rbp
        ret

Cc3_    endp
end
;-----
;##### Ch05_12.asm
;-----
; extern "C" bool Cc4_(const double* ht, const double* wt, int n, double* bsa1, double* bsa2,
double* bsa3);
        include <MacrosX86-64-AVX.asmh>

        .const
r8_0p007184    real8 0.007184
r8_0p725      real8 0.725
r8_0p425      real8 0.425
r8_0p0235     real8 0.0235
r8_0p42246    real8 0.42246
r8_0p51456    real8 0.51456
r8_3600p0     real8 3600.0

        .code
extern pow:proc

Cc4_ proc frame
        _CreateFrame Cc4_,16,64,rbx,rsi,r12,r13,r14,r15
        _SaveXmmRegs xmm6,xmm7,xmm8,xmm9
        _EndProlog

; Save argument registers to home area (optional). Note that the home
; area can also be used to store other transient data values.
        mov qword ptr [rbp+Cc4_OffsetHomeRCX],rcx
        mov qword ptr [rbp+Cc4_OffsetHomeRDX],rdx
        mov qword ptr [rbp+Cc4_OffsetHomeR8],r8
        mov qword ptr [rbp+Cc4_OffsetHomeR9],r9

; Initialize processing loop pointers. Note that the pointers are
; maintained in non-volatile registers, which eliminates reloads
; after the calls to pow().
        test r8d,r8d                ;is n > 0?
        jg @F                      ;jump if n > 0

```

```

        xor eax,eax                  ;set error return code
        jmp Done

@@:   mov [rbp],r8d                  ;save n to local var
        mov r12,rcx                  ;r12 = ptr to ht
        mov r13,rdx                  ;r13 = ptr to wt
        mov r14,r9                   ;r14 = ptr to bsa1
        mov r15,[rbp+Cc4_OffsetStackArgs] ;r15 = ptr to bsa2
        mov rbx,[rbp+Cc4_OffsetStackArgs+8] ;rbx = ptr to bsa3
        xor rsi,rsi                  ;array element offset

; Allocate home space on stack for use by pow()
        sub rsp,32

; Calculate bsa1 = 0.007184 * pow(ht, 0.725) * pow(wt, 0.425);
@@:   vmovsd xmm0,real8 ptr [r12+rsi] ;xmm0 = height
        vmovsd xmm8,xmm8,xmm0
        vmovsd xmm1,real8 ptr [r8_0p725]
        call pow                      ;xmm0 = pow(ht, 0.725)
        vmovsd xmm6,xmm6,xmm0

        vmovsd xmm0,real8 ptr [r13+rsi] ;xmm0 = weight
        vmovsd xmm9,xmm9,xmm0
        vmovsd xmm1,real8 ptr [r8_0p425]
        call pow                      ;xmm0 = pow(wt, 0.425)
        vmulsd xmm6,xmm6,real8 ptr [r8_0p007184]
        vmulsd xmm6,xmm6,xmm0          ;xmm6 = bsa1

; Calculate bsa2 = 0.0235 * pow(ht, 0.42246) * pow(wt, 0.51456);
        vmovsd xmm0,xmm0,xmm8          ;xmm0 = height
        vmovsd xmm1,real8 ptr [r8_0p42246]
        call pow                      ;xmm0 = pow(ht, 0.42246)
        vmovsd xmm7,xmm7,xmm0

        vmovsd xmm0,xmm0,xmm9          ;xmm0 = weight
        vmovsd xmm1,real8 ptr [r8_0p51456]
        call pow                      ;xmm0 = pow(wt, 0.51456)
        vmulsd xmm7,xmm7,real8 ptr [r8_0p0235]
        vmulsd xmm7,xmm7,xmm0          ;xmm7 = bsa2

; Calculate bsa3 = sqrt(ht * wt / 60.0);
        vmulsd xmm8,xmm8,xmm9
        vdivsd xmm8,xmm8,real8 ptr [r8_3600p0]
        vsqrtsd xmm8,xmm8,xmm8          ;xmm8 = bsa3

; Save BSA results
        vmovsd real8 ptr [r14+rsi],xmm6 ;save bsa1 result
        vmovsd real8 ptr [r15+rsi],xmm7 ;save bsa2 result
        vmovsd real8 ptr [rbx+rsi],xmm8 ;save bsa3 result

        add rsi,8                    ;update array offset
        dec dword ptr [rbp]          ;n = n - 1
        jnz @B
        mov eax,1                    ;set success return code

Done:   _RestoreXmmRegs xmm6,xmm7,xmm8,xmm9
        _DeleteFrame rbx,rsi,r12,r13,r14,r15
        ret

Cc4_ endp
end
;-----
;##### Ch06_01.asm
;-----
        .const
        align 16
AbsMaskF32    dword 7fffffffh,7fffffffh,7fffffffh,7fffffffh ;Absolute value mask for SPFP

```

AbsMaskF64 qword 7fffffffffffffffh, 7fffffffffffffffh ;Absolute value mask for DPFP

; extern "C" void AvxPackedMathF32\_(const XmmVal& a, const XmmVal& b, XmmVal c[8]);

```
.code
AvxPackedMathF32_ proc
; Load packed SPFP values
vmovaps xmm0, xmmword ptr [rcx] ;xmm0 = a
vmovaps xmm1, xmmword ptr [rdx] ;xmm1 = b
```

```
; Packed SPFP addition
vaddps xmm2, xmm0, xmm1
vmovaps [r8+0], xmm2
```

```
; Packed SPFP subtraction
vsubps xmm2, xmm0, xmm1
vmovaps [r8+16], xmm2
```

```
; Packed SPFP multiplication
vmulps xmm2, xmm0, xmm1
vmovaps [r8+32], xmm2
```

```
; Packed SPFP division
vdivps xmm2, xmm0, xmm1
vmovaps [r8+48], xmm2
```

```
; Packed SPFP absolute value (b)
vandps xmm2, xmm1, xmmword ptr [AbsMaskF32]
vmovaps [r8+64], xmm2
```

```
; Packed SPFP square root (a)
vsqrtps xmm2, xmm0
vmovaps [r8+80], xmm2
```

```
; Packed SPFP minimum
vminps xmm2, xmm0, xmm1
vmovaps [r8+96], xmm2
```

```
; Packed SPFP maximum
vmaxps xmm2, xmm0, xmm1
vmovaps [r8+112], xmm2
ret
```

AvxPackedMathF32\_ endp

; extern "C" void AvxPackedMathF64\_(const XmmVal& a, const XmmVal& b, XmmVal c[8]);

```
AvxPackedMathF64_ proc
; Load packed DPFP values
vmovapd xmm0, xmmword ptr [rcx] ;xmm0 = a
vmovapd xmm1, xmmword ptr [rdx] ;xmm1 = b
```

```
; Packed DPFP addition
vaddpd xmm2, xmm0, xmm1
vmovapd [r8+0], xmm2
```

```
; Packed DPFP subtraction
vsubpd xmm2, xmm0, xmm1
vmovapd [r8+16], xmm2
```

```
; Packed DPFP multiplication
vmulpd xmm2, xmm0, xmm1
vmovapd [r8+32], xmm2
```

```
; Packed DPFP division
vdivpd xmm2, xmm0, xmm1
vmovapd [r8+48], xmm2
```

```
; Packed DPFP absolute value (b)
vandpd xmm2, xmm1, xmmword ptr [AbsMaskF64]
vmovapd [r8+64], xmm2
```

```
; Packed DPFP square root (a)
vsqrtpd xmm2, xmm0
vmovapd [r8+80], xmm2
```

```
; Packed DPFP minimum
vminpd xmm2, xmm0, xmm1
vmovapd [r8+96], xmm2
```

```
; Packed DPFP maximum
vmaxpd xmm2, xmm0, xmm1
vmovapd [r8+112], xmm2
ret
```

AvxPackedMathF64\_ endp

```
;-----
;##### Ch06_02.asm
include <cmpequ.asmh>
```

; extern "C" void AvxPackedCompareF32\_(const XmmVal& a, const XmmVal& b, XmmVal c[8]);

```
.code
AvxPackedCompareF32_ proc
vmovaps xmm0, [rcx] ;xmm0 = a
vmovaps xmm1, [rdx] ;xmm1 = b
```

```
; Perform packed EQUAL compare
vcmpsps xmm2, xmm0, xmm1, CMP_EQ
vmovdqa xmmword ptr [r8], xmm2
```

```
; Perform packed NOT EQUAL compare
vcmpsps xmm2, xmm0, xmm1, CMP_NEQ
vmovdqa xmmword ptr [r8+16], xmm2
```

```
; Perform packed LESS THAN compare
vcmpsps xmm2, xmm0, xmm1, CMP_LT
vmovdqa xmmword ptr [r8+32], xmm2
```

```
; Perform packed LESS THAN OR EQUAL compare
vcmpsps xmm2, xmm0, xmm1, CMP_LE
vmovdqa xmmword ptr [r8+48], xmm2
```

```
; Perform packed GREATER THAN compare
vcmpsps xmm2, xmm0, xmm1, CMP_GT
vmovdqa xmmword ptr [r8+64], xmm2
```

```
; Perform packed GREATER THAN OR EQUAL compare
vcmpsps xmm2, xmm0, xmm1, CMP_GE
vmovdqa xmmword ptr [r8+80], xmm2
```

```
; Perform packed ORDERED compare
vcmpsps xmm2, xmm0, xmm1, CMP_ORD
vmovdqa xmmword ptr [r8+96], xmm2
```

```
; Perform packed UNORDERED compare
vcmpsps xmm2, xmm0, xmm1, CMP_UNORD
vmovdqa xmmword ptr [r8+112], xmm2
ret
```

AvxPackedCompareF32\_ endp

; extern "C" void AvxPackedCompareF64\_(const XmmVal& a, const XmmVal& b, XmmVal c[8]);

```
AvxPackedCompareF64_ proc
vmovapd xmm0, [rcx] ;xmm0 = a
```

```

        vmovapd xmm1, [rdx]                ;xmm1 = b

; Perform packed EQUAL compare
vcmpdpd xmm2, xmm0, xmm1, CMP_EQ
vmovdqa xmmword ptr [r8], xmm2

; Perform packed NOT EQUAL compare
vcmpdpd xmm2, xmm0, xmm1, CMP_NEQ
vmovdqa xmmword ptr [r8+16], xmm2

; Perform packed LESS THAN compare
vcmpdpd xmm2, xmm0, xmm1, CMP_LT
vmovdqa xmmword ptr [r8+32], xmm2

; Perform packed LESS THAN OR EQUAL compare
vcmpdpd xmm2, xmm0, xmm1, CMP_LE
vmovdqa xmmword ptr [r8+48], xmm2

; Perform packed GREATER THAN compare
vcmpdpd xmm2, xmm0, xmm1, CMP_GT
vmovdqa xmmword ptr [r8+64], xmm2

; Perform packed GREATER THAN OR EQUAL compare
vcmpdpd xmm2, xmm0, xmm1, CMP_GE
vmovdqa xmmword ptr [r8+80], xmm2

; Perform packed ORDERED compare
vcmpdpd xmm2, xmm0, xmm1, CMP_ORD
vmovdqa xmmword ptr [r8+96], xmm2

; Perform packed UNORDERED compare
vcmpdpd xmm2, xmm0, xmm1, CMP_UNORD
vmovdqa xmmword ptr [r8+112], xmm2
ret
AvxPackedCompareF64_ endp
end

;-----
;##### Ch06_03.asm
;-----
; extern "C" bool AvxPackedConvertFP_(const XmmVal& a, XmmVal& b, CvtOp cvt_op);
; Note: This function requires linker option /LARGEADDRESSAWARE:NO
; to be explicitly set.

.code
AvxPackedConvertFP_ proc

; Make sure cvt_op is valid
mov r9d, r8d                ;r9 = cvt_op (zero extended)
cmp r9, CvtOpTableCount     ;is cvt_op valid?
jae InvalidCvtOp            ;jump if cvt_op is invalid

mov eax, 1                   ;set valid cvt_op return code
jmp [CvtOpTable+r9*8]       ;jump to specified conversion

; Convert packed signed doubleword integers to packed SPFP values
I32_F32:
vmovdqa xmm0, xmmword ptr [rcx]
vcvtdq2ps xmm1, xmm0
vmovaps xmmword ptr [rdx], xmm1
ret

; Convert packed SPFP values to packed signed doubleword integers
F32_I32:
vmovaps xmm0, xmmword ptr [rcx]
vcvtps2dq xmm1, xmm0
vmovdqa xmmword ptr [rdx], xmm1
ret

```

```

; Convert packed signed doubleword integers to packed DPFP values
I32_F64:
vmovdqa xmm0, xmmword ptr [rcx]
vcvtdq2pd xmm1, xmm0
vmovapd xmmword ptr [rdx], xmm1
ret

; Convert packed DPFP values to packed signed doubleword integers
F64_I32:
vmovapd xmm0, xmmword ptr [rcx]
vcvtpd2dq xmm1, xmm0
vmovdqa xmmword ptr [rdx], xmm1
ret

; Convert packed SPFP to packed DPFP
F32_F64:
vmovaps xmm0, xmmword ptr [rcx]
vcvtps2pd xmm1, xmm0
vmovapd xmmword ptr [rdx], xmm1
ret

; Convert packed DPFP to packed SPFP
F64_F32:
vmovapd xmm0, xmmword ptr [rcx]
vcvtpd2ps xmm1, xmm0
vmovaps xmmword ptr [rdx], xmm1
ret

InvalidCvtOp:
xor eax, eax                ;set invalid cvt_op return code
ret

; The order of values in the following table must match the enum CvtOp
; that's defined in Ch06_03.cpp.

```

```

        align 8
CvtOpTable qword I32_F32, F32_I32
           qword I32_F64, F64_I32
           qword F32_F64, F64_F32
CvtOpTableCount equ ($ - CvtOpTable) / size qword

AvxPackedConvertFP_ endp
end

;-----
;##### Ch06_04.asm
;-----
; extern "C" bool AvxCalcSqrts_(float* y, const float* x, size_t n);

.code
AvxCalcSqrts_ proc
xor eax, eax                ;set error return code (also array offset)
test r8, r8
jz Done                     ;jump if n is zero

test rcx, 0fh               ;jump if 'y' is not aligned
jnz Done

test rdx, 0fh               ;jump if 'x' is not aligned
jnz Done

; Calculate packed square roots
cmp r8, 4
jb FinalVals                ;jump if n < 4
@@: vsqrtps xmm0, xmmword ptr [rdx+rax] ;calculate 4 square roots x[i+3:i]
vmovaps xmmword ptr [rcx+rax], xmm0 ;save results to y[i+3:i]

```

```

    add rax,16                ;update offset to next set of values
    sub r8,4
    cmp r8,4                  ;are there 4 or more elements remaining?
    jae @B                    ;jump if yes

; Calculate square roots of final 1 - 3 values, note switch to scalar instructions
FinalVals:
    test r8,r8                ;more elements to process?
    jz SetRC                  ;jump if no more elements

    vsqrtss xmm0,xmm0,real4 ptr [rdx+rax] ;calculate sqrt(x[i])
    vmovss real4 ptr [rcx+rax],xmm0      ;save result to y[i]
    add rax,4
    dec r8
    jz SetRC

    vsqrtss xmm0,xmm0,real4 ptr [rdx+rax]
    vmovss real4 ptr [rcx+rax],xmm0
    add rax,4
    dec r8
    jz SetRC

    vsqrtss xmm0,xmm0,real4 ptr [rdx+rax]
    vmovss real4 ptr [rcx+rax],xmm0

SetRC: mov eax,1                ;set success return code

Done:  ret
AvxCalcSqrts_ endp
;-----
;##### Ch06_05.asm
;-----

extern g_MinValInit:real4
extern g_MaxValInit:real4

; extern "C" bool CalcArrayMinMaxF32(float* min_val, float* max_val, const float* x, size_t n)
.code
CalcArrayMinMaxF32_ proc
; Validate arguments
xor eax,eax                ;set error return code

    test r8,0fh              ;is x aligned to 16-byte boundary?
    jnz Done                  ;jump if no

    vbroadcastss xmm4,real4 ptr [g_MinValInit] ;xmm4 = min values
    vbroadcastss xmm5,real4 ptr [g_MaxValInit] ;xmm5 = max values

    cmp r9,4
    jb FinalVals              ;jump if n < 4

; Main processing loop
@@: vmovaps xmm0,xmmword ptr [r8] ;load next set of array values
    vminps xmm4,xmm4,xmm0      ;update packed min values
    vmaxps xmm5,xmm5,xmm0      ;update packed max values

    add r8,16
    sub r9,4
    cmp r9,4
    jae @B

; Process the final 1 - 3 values of the input array
FinalVals:
    test r9,r9
    jz SaveResults

    vminss xmm4,xmm4,real4 ptr [r8] ;update packed min values

```

```

    vmaxss xmm5,xmm5,real4 ptr [r8] ;update packed max values
    dec r9
    jz SaveResults

    vminss xmm4,xmm4,real4 ptr [r8+4]
    vmaxss xmm5,xmm5,real4 ptr [r8+4]
    dec r9
    jz SaveResults

    vminss xmm4,xmm4,real4 ptr [r8+8]
    vmaxss xmm5,xmm5,real4 ptr [r8+8]

; Calculate and save final min & max values
SaveResults:
    vshufps xmm0,xmm4,xmm4,00001110b ;xmm0[63:0] = xmm4[128:64]
    vminps xmm1,xmm0,xmm4              ;xmm1[63:0] contains final 2 values
    vshufps xmm2,xmm1,xmm1,00000001b ;xmm2[31:0] = xmm1[63:32]
    vminps xmm3,xmm2,xmm1              ;xmm3[31:0] contains final value
    vmovss real4 ptr [rcx],xmm3        ;save array min value

    vshufps xmm0,xmm5,xmm5,00001110b
    vmxpps xmm1,xmm0,xmm5
    vshufps xmm2,xmm1,xmm1,00000001b
    vmxpps xmm3,xmm2,xmm1
    vmovss real4 ptr [rdx],xmm3        ;save array max value

    mov eax,1                        ;set success return code
Done:  ret
CalcArrayMinMaxF32_ endp
;-----
;##### Ch06_06.asm
;-----
include <MacrosX86-64-AVX.asmh>

extern LsEpsilon:real8          ;global value defined in C++ file

; extern "C" bool AvxCalcLeastSquares(const double* x, const double* y, int n, double* m, double* b);
; Returns 0 = error (invalid n or improperly aligned array), 1 = success

.const
align 16
AbsMaskF64 qword 7fffffffffffffffh, 7fffffffffffffffh ;mask for DFPF absolute value

.code
AvxCalcLeastSquares_ proc frame
    _CreateFrame LS_,0,48,rbx
    _SaveXmmRegs xmm6,xmm7,xmm8
    _EndProlog

; Validate arguments
xor eax,eax                ;set error return code
cmp r8d,2
jl Done                    ;jump if n < 2
test rcx,0fh
jnz Done                    ;jump if x not aligned to 16-byte boundary
test rdx,0fh
jnz Done                    ;jump if y not aligned to 16-byte boundary

; Perform required initializations
vcvttsi2sd xmm3,xmm3,r8d    ;xmm3 = n
mov eax,r8d
and r8d,0fffffffh           ;rd8 = n / 2 * 2
and eax,1                   ;eax = n % 2

vxorpd xmm4,xmm4,xmm4       ;sum_x (both qwords)
vxorpd xmm5,xmm5,xmm5       ;sum_y (both qwords)
vxorpd xmm6,xmm6,xmm6       ;sum_xx (both qwords)

```

```

    vxorpd xmm7, xmm7, xmm7      ;sum_xy (both qwords)
    xor ebx, ebx                 ;rbx = array offset
    mov r10, [rbp+LS_OffsetStackArgs] ;r10 = b

; Calculate sum variables. Note that two values are processed each iteration.
@@:  vmovapd xmm0, xmmword ptr [rcx+rbx] ;load next two x values
     vmovapd xmm1, xmmword ptr [rdx+rbx] ;load next two y values

     vaddpd xmm4, xmm4, xmm0        ;update sum_x
     vaddpd xmm5, xmm5, xmm1        ;update sum_y

     vmulpd xmm2, xmm0, xmm0        ;calc x * x
     vaddpd xmm6, xmm6, xmm2        ;update sum_xx

     vmulpd xmm2, xmm0, xmm1        ;calc x * y
     vaddpd xmm7, xmm7, xmm2        ;update sum_xy

     add rbx, 16                   ;rbx = next offset
     sub r8d, 2                    ;adjust counter
     jnz @B                        ;repeat until done

; Update sum variables with the final x, y values if 'n' is odd
or eax, eax
jz CalcFinalSums                ;jump if n is even
vmovsd xmm0, real8 ptr [rcx+rbx] ;load final x
vmovsd xmm1, real8 ptr [rdx+rbx] ;load final y

vaddsd xmm4, xmm4, xmm0          ;update sum_x
vaddsd xmm5, xmm5, xmm1          ;update sum_y

vmulsd xmm2, xmm0, xmm0          ;calc x * x
vaddsd xmm6, xmm6, xmm2          ;update sum_xx

vmulsd xmm2, xmm0, xmm1          ;calc x * y
vaddsd xmm7, xmm7, xmm2          ;update sum_xy

; Calculate final sum_x, sum_y, sum_xx, sum_xy
CalcFinalSums:
    vhaddpd xmm4, xmm4, xmm4      ;xmm4[63:0] = final sum_x
    vhaddpd xmm5, xmm5, xmm5      ;xmm5[63:0] = final sum_y
    vhaddpd xmm6, xmm6, xmm6      ;xmm6[63:0] = final sum_xx
    vhaddpd xmm7, xmm7, xmm7      ;xmm7[63:0] = final sum_xy

; Compute denominator and make sure it's valid
; denom = n * sum_xx - sum_x * sum_x
    vmulsd xmm0, xmm3, xmm6       ;n * sum_xx
    vmulsd xmm1, xmm4, xmm4       ;sum_x * sum_x
    vsubsd xmm2, xmm0, xmm1       ;denom
    vandpd xmm8, xmm2, xmmword ptr [AbsMaskF64] ;fabs(denom)
    vcomisd xmm8, real8 ptr [LsEpsilon]
    jb BadDen                    ;jump if denom < fabs(denom)

; Compute and save slope
; slope = (n * sum_xy - sum_x * sum_y) / denom
    vmulsd xmm0, xmm3, xmm7       ;n * sum_xy
    vmulsd xmm1, xmm4, xmm5       ;sum_x * sum_y
    vsubsd xmm2, xmm0, xmm1       ;slope numerator
    vdivsd xmm3, xmm2, xmm8       ;final slope
    vmovsd real8 ptr [r9], xmm3    ;save slope

; Compute and save intercept
; intercept = (sum_xx * sum_y - sum_x * sum_xy) / denom
    vmulsd xmm0, xmm6, xmm5       ;sum_xx * sum_y
    vmulsd xmm1, xmm4, xmm7       ;sum_x * sum_xy
    vsubsd xmm2, xmm0, xmm1       ;intercept numerator
    vdivsd xmm3, xmm2, xmm8       ;final intercept
    vmovsd real8 ptr [r10], xmm3   ;save intercept

```

```

    mov eax, 1                   ;success return code
    jmp Done

; Bad denominator detected, set m and b to 0.0
BadDen: vxorpd xmm0, xmm0, xmm0
        vmovsd real8 ptr [r9], xmm0 ;*m = 0.0
        vmovsd real8 ptr [r10], xmm0 ;*b = 0.0
        xor eax, eax              ;set error code

Done:   _RestoreXmmRegs xmm6, xmm7, xmm8
        _DeleteFrame rbx
        ret
AvxCalcLeastSquares_ endp
end

;-----
;##### Ch06_07.asm
;include <MacrosX86-64-AVX.asmh>

; _Mat4x4TransposeF32 macro
; Description: This macro transposes a 4x4 matrix of single-precision
;              floating-point values.
;
;   Input Matrix           Output Matrix
;-----
;   xmm0    a3 a2 a1 a0      xmm4    d0 c0 b0 a0
;   xmm1    b3 b2 b1 b0      xmm5    d1 c1 b1 a1
;   xmm2    c3 c2 c1 c0      xmm6    d2 c2 b2 a2
;   xmm3    d3 d2 d1 d0      xmm7    d3 c3 b3 a3

;_Mat4x4TransposeF32 macro
;vunpcklps xmm6, xmm0, xmm1      ;xmm6 = b1 a1 b0 a0
;vunpckhps xmm0, xmm0, xmm1      ;xmm0 = b3 a3 b2 a2
;vunpcklps xmm7, xmm2, xmm3      ;xmm7 = d1 c1 d0 c0
;vunpckhps xmm1, xmm2, xmm3      ;xmm1 = d3 c3 d2 c2

;vmovhlps xmm4, xmm6, xmm7      ;xmm4 = d0 c0 b0 a0
;vmovhlps xmm5, xmm7, xmm6      ;xmm5 = d1 c1 b1 a1
;vmovhlps xmm6, xmm0, xmm1      ;xmm6 = d2 c2 b2 a2
;vmovhlps xmm7, xmm1, xmm0      ;xmm7 = d3 c3 b2 a3
;endm

; extern "C" void AvxMat4x4TransposeF32_(float* m_des, const float* m_src)
; .code
AvxMat4x4TransposeF32_ proc frame
    _CreateFrame MT_, 0, 32
    _SaveXmmRegs xmm6, xmm7
    _EndProlog

; Transpose matrix m_src1
    vmovaps xmm0, [rdx]          ;xmm0 = m_src.row_0
    vmovaps xmm1, [rdx+16]       ;xmm1 = m_src.row_1
    vmovaps xmm2, [rdx+32]       ;xmm2 = m_src.row_2
    vmovaps xmm3, [rdx+48]       ;xmm3 = m_src.row_3

    _Mat4x4TransposeF32

    vmovaps [rcx], xmm4          ;save m_des.row_0
    vmovaps [rcx+16], xmm5       ;save m_des.row_1
    vmovaps [rcx+32], xmm6       ;save m_des.row_2
    vmovaps [rcx+48], xmm7       ;save m_des.row_3

Done:   _RestoreXmmRegs xmm6, xmm7
        _DeleteFrame
        ret
AvxMat4x4TransposeF32_ endp
end

;-----
;##### Ch06_08.asm

```



```

include <MacrosX86-64-AVX.asmh>

; _Mat4x4MulCalcRowF32 macro
; Description: This macro is used to compute one row of a 4x4 matrix
; multiply.
; Registers: xmm0 = m_src2.row0
;            xmm1 = m_src2.row1
;            xmm2 = m_src2.row2
;            xmm3 = m_src2.row3
;            rcx = m_des.ptr
;            rdx = m_src1.ptr
;            xmm4 - xmm7 = scratch registers

_Mat4x4MulCalcRowF32 macro disp
vbroadcastss xmm4, real4 ptr [rdx+disp] :broadcast m_src1[i][0]
vbroadcastss xmm5, real4 ptr [rdx+disp+4] :broadcast m_src1[i][1]
vbroadcastss xmm6, real4 ptr [rdx+disp+8] :broadcast m_src1[i][2]
vbroadcastss xmm7, real4 ptr [rdx+disp+12] :broadcast m_src1[i][3]

vmulps xmm4, xmm4, xmm0 :m_src1[i][0] * m_src2.row_0
vmulps xmm5, xmm5, xmm1 :m_src1[i][1] * m_src2.row_1
vmulps xmm6, xmm6, xmm2 :m_src1[i][2] * m_src2.row_2
vmulps xmm7, xmm7, xmm3 :m_src1[i][3] * m_src2.row_3

vaddps xmm4, xmm4, xmm5 :calc m_des.row_i
vaddps xmm6, xmm6, xmm7
vaddps xmm4, xmm4, xmm6

vmovaps [rcx+disp], xmm4 :save m_des.row_i
endm

; extern "C" void AvxMat4x4MulF32(float* m_des, const float* m_src1, const float* m_src2)
; Description: The following function computes the product of two
; single-precision floating-point 4x4 matrices.

.code
AvxMat4x4MulF32 proc frame
_CreateFrame MM_, 0, 32
_SaveXmmRegs xmm6, xmm7
_EndProlog

; Compute matrix product m_des = m_src1 * m_src2
vmovaps xmm0, [r8] :xmm0 = m_src2.row_0
vmovaps xmm1, [r8+16] :xmm1 = m_src2.row_1
vmovaps xmm2, [r8+32] :xmm2 = m_src2.row_2
vmovaps xmm3, [r8+48] :xmm3 = m_src2.row_3

_Mat4x4MulCalcRowF32 0 :calculate m_des.row_0
_Mat4x4MulCalcRowF32 16 :calculate m_des.row_1
_Mat4x4MulCalcRowF32 32 :calculate m_des.row_2
_Mat4x4MulCalcRowF32 48 :calculate m_des.row_3

Done: _RestoreXmmRegs xmm6, xmm7
_DeleteFrame
ret
AvxMat4x4MulF32 endp
end

;-----
;##### Ch07_01.asm
;-----
; extern "C" void AvxPackedAddI16(const XmmVal& a, const XmmVal& b, XmmVal c[2])
.code
AvxPackedAddI16 proc

; Packed signed word addition
vmovdqa xmm0, xmmword ptr [rcx] :xmm0 = a
vmovdqa xmm1, xmmword ptr [rdx] :xmm1 = b

vpaddw xmm2, xmm0, xmm1 :packed add - wraparound
vpaddsw xmm3, xmm0, xmm1 :packed add - saturated

vmovdqa xmmword ptr [r8], xmm2 :save c[0]
vmovdqa xmmword ptr [r8+16], xmm3 :save c[1]
ret
AvxPackedAddI16 endp

;-----
;##### Ch07_02.asm
;-----
; extern "C" bool AvxPackedIntegerShift_(XmmVal& b, const XmmVal& a, ShiftOp shift_op, unsigned int count)
; Returns: 0 = invalid shift_op argument, 1 = success
; Note: This module requires linker option /LARGEADDRESSAWARE:NO
; to be explicitly set.

.code
AvxPackedIntegerShift_ proc
; Make sure 'shift_op' is valid
mov r8d, r8d :zero extend shift_op
cmp r8, ShiftOpTableCount :compare against table count

```

```

vpaddw xmm2, xmm0, xmm1 :packed add - wraparound
vpaddsw xmm3, xmm0, xmm1 :packed add - saturated

vmovdqa xmmword ptr [r8], xmm2 :save c[0]
vmovdqa xmmword ptr [r8+16], xmm3 :save c[1]
ret
AvxPackedAddI16_ endp

; extern "C" void AvxPackedSubI16_(const XmmVal& a, const XmmVal& b, XmmVal c[2])
AvxPackedSubI16_ proc

; Packed signed word subtraction
vmovdqa xmm0, xmmword ptr [rcx] :xmm0 = a
vmovdqa xmm1, xmmword ptr [rdx] :xmm1 = b

vpsubw xmm2, xmm0, xmm1 :packed sub - wraparound
vpsubsw xmm3, xmm0, xmm1 :packed sub - saturated

vmovdqa xmmword ptr [r8], xmm2 :save c[0]
vmovdqa xmmword ptr [r8+16], xmm3 :save c[1]
ret
AvxPackedSubI16_ endp

; extern "C" void AvxPackedAddU16_(const XmmVal& a, const XmmVal& b, XmmVal c[2])
AvxPackedAddU16_ proc

; Packed unsigned word addition
vmovdqu xmm0, xmmword ptr [rcx] :xmm0 = a
vmovdqu xmm1, xmmword ptr [rdx] :xmm1 = b

vpaddw xmm2, xmm0, xmm1 :packed add - wraparound
vpaddusw xmm3, xmm0, xmm1 :packed add - saturated

vmovdqu xmmword ptr [r8], xmm2 :save c[0]
vmovdqu xmmword ptr [r8+16], xmm3 :save c[1]
ret
AvxPackedAddU16_ endp

; extern "C" void AvxPackedSubU16_(const XmmVal& a, const XmmVal& b, XmmVal c[2])
AvxPackedSubU16_ proc

; Packed unsigned word subtraction
vmovdqu xmm0, xmmword ptr [rcx] :xmm0 = a
vmovdqu xmm1, xmmword ptr [rdx] :xmm1 = b

vpsubw xmm2, xmm0, xmm1 :packed sub - wraparound
vpsubusw xmm3, xmm0, xmm1 :packed sub - saturated

vmovdqu xmmword ptr [r8], xmm2 :save c[0]
vmovdqu xmmword ptr [r8+16], xmm3 :save c[1]
ret
AvxPackedSubU16_ endp
end

;-----
;##### Ch07_02.asm
;-----
; extern "C" bool AvxPackedIntegerShift_(XmmVal& b, const XmmVal& a, ShiftOp shift_op, unsigned int count)
; Returns: 0 = invalid shift_op argument, 1 = success
; Note: This module requires linker option /LARGEADDRESSAWARE:NO
; to be explicitly set.

.code
AvxPackedIntegerShift_ proc
; Make sure 'shift_op' is valid
mov r8d, r8d :zero extend shift_op
cmp r8, ShiftOpTableCount :compare against table count

```

```
jae Error ;jump if shift_op is invalid
```

```
; Jump to the operation specified by shift_op
vmovdqa xmm0,xmmword ptr [rdx] ;xmm0 = a
vmovd xmm1,r9d ;xmm1[31:0] = shift count
mov eax,1 ;set success return code
jmp [ShiftOpTable+r8*8]
```

```
; Packed shift left logical - word
U16_LL: vpsllw xmm2,xmm0,xmm1
vmovdqa xmmword ptr [rcx],xmm2
ret
```

```
; Packed shift right logical - word
U16_RL: vpsrlw xmm2,xmm0,xmm1
vmovdqa xmmword ptr [rcx],xmm2
ret
```

```
; Packed shift right arithmetic - word
U16_RA: vpsraw xmm2,xmm0,xmm1
vmovdqa xmmword ptr [rcx],xmm2
ret
```

```
; Packed shift left logical - doubleword
U32_LL: vpslld xmm2,xmm0,xmm1
vmovdqa xmmword ptr [rcx],xmm2
ret
```

```
; Packed shift right logical - doubleword
U32_RL: vpsrld xmm2,xmm0,xmm1
vmovdqa xmmword ptr [rcx],xmm2
ret
```

```
; Packed shift right arithmetic - doubleword
U32_RA: vpsrad xmm2,xmm0,xmm1
vmovdqa xmmword ptr [rcx],xmm2
ret
```

```
Error: xor eax,eax ;set error code
vpxor xmm0,xmm0,xmm0
vmovdqa xmmword ptr [rcx],xmm0 ;set result to zero
ret
```

```
; The order of the labels in the following table must correspond
; to the enums that are defined in .cpp file.
```

```
align 8
ShiftOpTable qword U16_LL, U16_RL, U16_RA
qword U32_LL, U32_RL, U32_RA
ShiftOpTableCount equ ($ - ShiftOpTable) / size qword
```

```
AvxPackedIntegerShift_ endp
end
```

```
-----
:##### Ch07_03.asm
:-----
```

```
; extern "C" void AvxPackedMulI16_(XmmVal c[2], const XmmVal* a, const XmmVal* b)
```

```
.code
AvxPackedMulI16_ proc
vmovdqa xmm0,xmmword ptr [rdx] ;xmm0 = a
vmovdqa xmm1,xmmword ptr [r8] ;xmm1 = b

vpmullw xmm2,xmm0,xmm1 ;xmm2 = packed a * b low result
vpmulhw xmm3,xmm0,xmm1 ;xmm3 = packed a * b high result

vpunpcklwd xmm4,xmm2,xmm3 ;merge low and high results
vpunpckhwd xmm5,xmm2,xmm3 ;into final signed dwords
```

```
vmovdqa xmmword ptr [rcx],xmm4 ;save final results
vmovdqa xmmword ptr [rcx+16],xmm5
ret
```

```
AvxPackedMulI16_ endp
```

```
; extern "C" void AvxPackedMulI32A_(XmmVal c[2], const XmmVal* a, const XmmVal* b)
AvxPackedMulI32A_ proc
```

```
; Perform packed signed dword multiplication. Note that vpmuldq
; performs following operations:
; xmm2[63:0] = xmm0[31:0] * xmm1[31:0]
; xmm2[127:64] = xmm0[95:64] * xmm1[95:64]
```

```
vmovdqa xmm0,xmmword ptr [rdx] ;xmm0 = a
vmovdqa xmm1,xmmword ptr [r8] ;xmm1 = b
vpmuldq xmm2,xmm0,xmm1
```

```
; Shift source operands right by 4 bytes and repeat vpmuldq
vpsrldq xmm0,xmm0,4
vpsrldq xmm1,xmm1,4
vpmuldq xmm3,xmm0,xmm1
```

```
; Save results
vpextrq qword ptr [rcx],xmm2,0 ;save xmm2[63:0]
vpextrq qword ptr [rcx+8],xmm3,0 ;save xmm3[63:0]
vpextrq qword ptr [rcx+16],xmm2,1 ;save xmm2[127:63]
vpextrq qword ptr [rcx+24],xmm3,1 ;save xmm3[127:63]
ret
```

```
AvxPackedMulI32A_ endp
```

```
; extern "C" void AvxPackedMulI32B_(XmmVal*, const XmmVal* a, const XmmVal* b)
AvxPackedMulI32B_ proc
```

```
; Perform packed signed integer multiplication and save low packed dword result
vmovdqa xmm0,xmmword ptr [rdx] ;xmm0 = a
vpmulld xmm1,xmm0,xmmword ptr [r8] ;xmm1 = packed a * b
vmovdqa xmmword ptr [rcx],xmm1 ;save packed dword result
ret
```

```
AvxPackedMulI32B_ endp
end
```

```
-----
:##### Ch07_04_.asm
:-----
```

```
; extern "C" bool AvxCalcMinMaxU8_(uint8_t* x, size_t n, uint8_t* x_min, uint8_t* x_max)
; Returns: 0 = invalid n or unaligned array, 1 = success
```

```
.const
align 16
StartMinVal qword 0fffffffffffffffh ;Initial packed min values
qword 0fffffffffffffffh
```

```
StartMaxVal qword 0000000000000000h ;Initial packed max values
qword 0000000000000000h
```

```
.code
AvxCalcMinMaxU8_ proc
```

```
; Make sure 'n' is valid
xor eax,eax ;set error return code
or rdx,rdx ;is n == 0?
jz Done ;jump if yes

test rdx,3fh ;is n a multiple of 64?
jnz Done ;jump if no

test rcx,0fh ;is x properly aligned?
```

```

        jnz Done                :jump if no

; Initialize packed min-max values
        vmovdqa xmm2,xmmword ptr [StartMinVal]
        vmovdqa xmm3,xmm2      :xmm3:xmm2 = packed min values
        vmovdqa xmm4,xmmword ptr [StartMaxVal]
        vmovdqa xmm5,xmm4      :xmm5:xmm4 = packed max values

; Scan array for min & max values
@@:     vmovdqa xmm0,xmmword ptr [rcx]      :xmm0 = x[i + 15] : x[i]
        vmovdqa xmm1,xmmword ptr [rcx+16]  :xmm1 = x[i + 31] : x[i + 16]
        vpmiub xmm2,xmm2,xmm0
        vpmiub xmm3,xmm3,xmm1      :xmm3:xmm2 = updated min values
        vpmiub xmm4,xmm4,xmm0
        vpmiub xmm5,xmm5,xmm1      :xmm5:xmm4 = updated max values

        vmovdqa xmm0,xmmword ptr [rcx+32]  :xmm0 = x[i + 47] : x[i + 32]
        vmovdqa xmm1,xmmword ptr [rcx+48]  :xmm1 = x[i + 63] : x[i + 48]
        vpmiub xmm2,xmm2,xmm0
        vpmiub xmm3,xmm3,xmm1      :xmm3:xmm2 = updated min values
        vpmiub xmm4,xmm4,xmm0
        vpmiub xmm5,xmm5,xmm1      :xmm5:xmm4 = updated max values

        add rcx,64
        sub rdx,64
        jnz @B

; Determine final minimum value
        vpmiub xmm0,xmm2,xmm3      :xmm0[127:0] = final 16 min vals
        vpsrldq xmm1,xmm0,8        :xmm1[63:0] = xmm0[127:64]
        vpmiub xmm2,xmm1,xmm0      :xmm2[63:0] = final 8 min vals
        vpsrldq xmm3,xmm2,4        :xmm3[31:0] = xmm2[63:32]
        vpmiub xmm0,xmm3,xmm2      :xmm0[31:0] = final 4 min vals
        vpsrldq xmm1,xmm0,2        :xmm1[15:0] = xmm0[31:16]
        vpmiub xmm2,xmm1,xmm0      :xmm2[15:0] = final 2 min vals
        vpextrw eax,xmm2,0         :ax = final 2 min vals
        cmp al,ah
        jbe @F                    :jump if al <= ah
        mov al,ah                 :al = final min value
@@:     mov [r8],al               :save final min

; Determine final maximum value
        vpmiub xmm0,xmm4,xmm5      :xmm0[127:0] = final 16 max vals
        vpsrldq xmm1,xmm0,8        :xmm1[63:0] = xmm0[127:64]
        vpmiub xmm2,xmm1,xmm0      :xmm2[63:0] = final 8 max vals
        vpsrldq xmm3,xmm2,4        :xmm3[31:0] = xmm2[63:32]
        vpmiub xmm0,xmm3,xmm2      :xmm0[31:0] = final 4 max vals
        vpsrldq xmm1,xmm0,2        :xmm1[15:0] = xmm0[31:16]
        vpmiub xmm2,xmm1,xmm0      :xmm2[15:0] = final 2 max vals
        vpextrw eax,xmm2,0         :ax = final 2 min vals
        cmp al,ah
        jae @F                    :jump if al >= ah
        mov al,ah                 :al = final max value
@@:     mov [r9],al               :save final max

        mov eax,1                 :set success return code
Done:   ret
AvxCalcMinMaxU8_endp
end

;-----
;##### Ch07_05.asm
;include <MacrosX86-64-AVX.asmh>
;extern g_NumElementsMax:qword

; extern "C" bool AvxCalcMeanU8_(const UInt8* x, size_t n, int64_t* sum_x, double* mean);
; Returns      0 = invalid n or unaligned array, 1 = success

```

```

        .code
AvxCalcMeanU8_ proc frame
        _CreateFrame CM_,0,64
        _SaveXmmRegs xmm6,xmm7,xmm8,xmm9
        _EndProlog

; Verify function arguments
        xor eax,eax                :set error return code
        or rdx,rdx
        jz Done                    :jump if n == 0

        cmp rdx,[g_NumElementsMax]
        jae Done                    :jump if n > NumElementsMax

        test rdx,3fh
        jnz Done                    :jump if (n % 64) != 0

        test rcx,0fh
        jnz Done                    :jump if x is not properly aligned

; Perform required initializations
        mov r10,rdx                :save n for later use
        add rdx,rcx                 :rdx = end of array
        vpxor xmm8,xmm8,xmm8       :xmm8 = packed intermediate sums (4 dwords)
        vpxor xmm9,xmm9,xmm9       :xmm9 = packed zero for promotions

; Promote 32 pixel values from bytes to words, then sum the words
@@:     vmovdqa xmm0,xmmword ptr [rcx]
        vmovdqa xmm1,xmmword ptr [rcx+16] :xmm1:xmm0 = 32 pixels
        vpunpcklbw xmm2,xmm0,xmm9    :xmm2 = 8 words
        vpunpckhbw xmm3,xmm0,xmm9    :xmm3 = 8 words
        vpunpcklbw xmm4,xmm1,xmm9    :xmm4 = 8 words
        vpunpckhbw xmm5,xmm1,xmm9    :xmm5 = 8 words
        vpaddw xmm0,xmm2,xmm3
        vpaddw xmm1,xmm4,xmm5
        vpaddw xmm6,xmm0,xmm1        :xmm6 = packed sums (8 words)

; Promote another 32 pixel values from bytes to words, then sum the words
        vmovdqa xmm0,xmmword ptr [rcx+32]
        vmovdqa xmm1,xmmword ptr [rcx+48] :xmm1:xmm0 = 32 pixels
        vpunpcklbw xmm2,xmm0,xmm9    :xmm2 = 8 words
        vpunpckhbw xmm3,xmm0,xmm9    :xmm3 = 8 words
        vpunpcklbw xmm4,xmm1,xmm9    :xmm4 = 8 words
        vpunpckhbw xmm5,xmm1,xmm9    :xmm5 = 8 words
        vpaddw xmm0,xmm2,xmm3
        vpaddw xmm1,xmm4,xmm5
        vpaddw xmm7,xmm0,xmm1        :xmm7 = packed sums (8 words)

; Promote packed sums to dwords, then update dword sums
        vpaddw xmm0,xmm6,xmm7        :xmm0 = packed sums (8 words)
        vpunpcklwd xmm1,xmm0,xmm9    :xmm1 = packed sums (4 dwords)
        vpunpckhwd xmm2,xmm0,xmm9    :xmm2 = packed sums (4 dwords)
        vpaddq xmm8,xmm8,xmm1
        vpaddq xmm8,xmm8,xmm2

        add rcx,64                  :rcx = next 64 byte block
        cmp rcx,rdx
        jne @B                      :repeat loop if not done

; Compute final sum_x (note vpextrd zero extends extracted dword to 64 bits)
        vpextrd eax,xmm8,0           :rax = partial sum 0
        vpextrd edx,xmm8,1           :rdx = partial sum 1
        add rax,rdx
        vpextrd ecx,xmm8,2           :rcx = partial sum 2
        vpextrd edx,xmm8,3           :rdx = partial sum 3
        add rax,rcx
        add rax,rdx
        mov [r8],rax                :save sum_x

```

```

; Compute mean value
vcvttsi2sd xmm0, xmm0, rax      ;xmm0 = sum_x (DPFP)
vcvttsi2sd xmm1, xmm1, r10      ;xmm1 = n (DPFP)
vdivsd xmm2, xmm0, xmm1        ;calc mean = sum_x / n
vmovsd real8 ptr [r9], xmm2     ;save mean

mov eax, 1                      ;set success return code

Done: _RestoreXmmRegs xmm6, xmm7, xmm8, xmm9
      _DeleteFrame
      ret
AvxCalcMeanU8_ endp
end

;-----
##### Ch07_06.asm
include <MacrosX86-64-AVX.asmh>      include <cmpequ.asmh>

        .const
        align 16
UInt8ToFloat      real4 255.0, 255.0, 255.0, 255.0
FloatToUInt8Min   real4 0.0, 0.0, 0.0, 0.0
FloatToUInt8Max   real4 1.0, 1.0, 1.0, 1.0
FloatToUInt8Scale real4 255.0, 255.0, 255.0, 255.0

extern c_NumPixelsMax:dword

; extern "C" bool ConvertImgU8ToF32_(float* des, const uint8_t* src, uint32_t num_pixels)
.code
ConvertImgU8ToF32_ proc frame
    _CreateFrame U2F_, 0, 160
    _SaveXmmRegs xmm6, xmm7, xmm8, xmm9, xmm10, xmm11, xmm12, xmm13, xmm14, xmm15
    _EndProlog

; Make sure num_pixels is valid and pixel buffers are properly aligned
xor eax, eax                ;set error return code
or r8d, r8d
jz Done                    ;jump if num_pixels is zero
cmp r8d, [c_NumPixelsMax]
ja Done                    ;jump if num_pixels too big
test r8d, 1fh
jnz Done                   ;jump if num_pixels % 32 != 0
test rcx, 0fh
jnz Done                   ;jump if des not aligned
test rdx, 0fh
jnz Done                   ;jump if src not aligned

; Initialize processing loop registers
shr r8d, 5                  ;number of pixel blocks
vmovaps xmm6, xmmword ptr [UInt8ToFloat] ;xmm6 = packed 255.0f
vpxor xmm7, xmm7, xmm7     ;xmm7 = packed 0

; Load the next block of 32 pixels
@@: vmovdq xmm0, xmmword ptr [rdx] ;xmm0 = 16 pixels (x[i+15]:x[i])
    vmovdq xmm1, xmmword ptr [rdx+16] ;xmm8 = 16 pixels (x[i+31]:x[i+16])
; Promote the pixel values in xmm0 from unsigned bytes to unsigned dwords
vpunpcklbw xmm2, xmm0, xmm7
vpunpckhbw xmm3, xmm0, xmm7
vpunpcklwd xmm8, xmm2, xmm7
vpunpckhwd xmm9, xmm2, xmm7
vpunpcklwd xmm10, xmm3, xmm7
vpunpckhwd xmm11, xmm3, xmm7      ;xmm11:xmm8 = 16 dword pixels

; Promote the pixel values in xmm1 from unsigned bytes to unsigned dwords
vpunpcklbw xmm2, xmm1, xmm7
vpunpckhbw xmm3, xmm1, xmm7
vpunpcklwd xmm12, xmm2, xmm7
vpunpckhwd xmm13, xmm2, xmm7

```

```

vpunpcklwd xmm14, xmm3, xmm7
vpunpckhwd xmm15, xmm3, xmm7      ;xmm15:xmm12 = 16 dword pixels

; Convert pixel values from dwords to SPFP
vcvtqdq2ps xmm8, xmm8
vcvtqdq2ps xmm9, xmm9
vcvtqdq2ps xmm10, xmm10
vcvtqdq2ps xmm11, xmm11          ;xmm11:xmm8 = 16 SPFP pixels

vcvtqdq2ps xmm12, xmm12
vcvtqdq2ps xmm13, xmm13
vcvtqdq2ps xmm14, xmm14
vcvtqdq2ps xmm15, xmm15          ;xmm15:xmm12 = 16 SPFP pixels

; Normalize all pixel values to [0.0, 1.0] and save the results
vdivps xmm0, xmm8, xmm6
vmovaps xmmword ptr [rcx], xmm0 ;save pixels 0 - 3
vdivps xmm1, xmm9, xmm6
vmovaps xmmword ptr [rcx+16], xmm1 ;save pixels 4 - 7
vdivps xmm2, xmm10, xmm6
vmovaps xmmword ptr [rcx+32], xmm2 ;save pixels 8 - 11
vdivps xmm3, xmm11, xmm6
vmovaps xmmword ptr [rcx+48], xmm3 ;save pixels 12 - 15

vdivps xmm0, xmm12, xmm6
vmovaps xmmword ptr [rcx+64], xmm0 ;save pixels 16 - 19
vdivps xmm1, xmm13, xmm6
vmovaps xmmword ptr [rcx+80], xmm1 ;save pixels 20 - 23
vdivps xmm2, xmm14, xmm6
vmovaps xmmword ptr [rcx+96], xmm2 ;save pixels 24 - 27
vdivps xmm3, xmm15, xmm6
vmovaps xmmword ptr [rcx+112], xmm3 ;save pixels 28 - 31

add rdx, 32                  ;update src ptr
add rcx, 128                 ;update des ptr
sub r8d, 1
jnz @B                      ;repeat until done
mov eax, 1                   ;set success return code

Done: _RestoreXmmRegs xmm6, xmm7, xmm8, xmm9, xmm10, xmm11, xmm12, xmm13, xmm14, xmm15
      _DeleteFrame
      ret

ConvertImgU8ToF32_ endp

; extern "C" bool ConvertImgF32ToU8_(uint8_t* des, const float* src, uint32_t num_pixels)
ConvertImgF32ToU8_ proc frame
    _CreateFrame F2U_, 0, 96
    _SaveXmmRegs xmm6, xmm7, xmm12, xmm13, xmm14, xmm15
    _EndProlog

; Make sure num_pixels is valid and pixel buffers are properly aligned
xor eax, eax                ;set error return code
or r8d, r8d
jz Done                    ;jump if num_pixels is zero
cmp r8d, [c_NumPixelsMax]
ja Done                    ;jump if num_pixels too big
test r8d, 1fh
jnz Done                   ;jump if num_pixels % 32 != 0
test rcx, 0fh
jnz Done                   ;jump if des not aligned
test rdx, 0fh
jnz Done                   ;jump if src not aligned

; Load required packed constants into registers
vmovaps xmm13, xmmword ptr [FloatToUInt8Scale] ;xmm13 = packed 255.0
vmovaps xmm14, xmmword ptr [FloatToUInt8Min]   ;xmm14 = packed 0.0

```

```

        vmovaps xmm15,xmmword ptr [FloatToUint8Max]    ;xmm15 = packed 1.0

LP1:    shr r8d,4          ;number of pixel blocks
        mov r9d,4          ;num pixel quartets per block

; Convert 16 float pixels to uint8_t
LP2:    vmovaps xmm0,xmmword ptr [rdx]                ;xmm0 = next pixel quartet
        vcmpps xmm1,xmm0,xmm14,CMP_LT                 ;compare pixels to 0.0
        vandnps xmm2,xmm1,xmm0                       ;clip pixels < 0.0 to 0.0

        vcmpps xmm3,xmm2,xmm15,CMP_GT                 ;compare pixels to 1.0
        vandps  xmm4,xmm3,xmm15                      ;clip pixels > 1.0 to 1.0
        vandnps xmm5,xmm3,xmm2                       ;xmm5 = pixels <= 1.0
        vorps   xmm6,xmm5,xmm4                       ;xmm6 = final clipped pixels
        vmulps  xmm7,xmm6,xmm13                      ;xmm7 = FP pixels [0.0, 255.0]

        vcvtps2dq xmm0,xmm7                          ;xmm0 = dword pixels [0, 255]
        vpackusdw xmm1,xmm0,xmm0                    ;xmm1[63:0] = word pixels
        vpackuswb xmm2,xmm1,xmm1                    ;xmm2[31:0] = bytes pixels

; Save the current byte pixel quartet
        vpextrd eax,xmm2,0                          ;eax = new pixel quartet
        vpsrldq xmm12,xmm12,4                       ;adjust xmm12 for new quartet
        vpinsrd xmm12,xmm12,eax,3                   ;xmm12[127:96] = new quartet

        add rdx,16                                  ;update src ptr
        sub r9d,1
        jnz LP2                                     ;repeat until done

; Save the current byte pixel block (16 pixels)
        vmovdqa xmmword ptr [rcx],xmm12             ;save current pixel block
        add rcx,16                                  ;update des ptr
        sub r8d,1
        jnz LP1                                     ;repeat until done
        mov eax,1                                  ;set success return code

Done:    _RestoreXmmRegs xmm6,xmm7,xmm12,xmm13,xmm14,xmm15
        _DeleteFrame
        ret

ConvertImgF32ToU8_ endp
end
;-----
;##### Ch07_07.asm
;include <MacrosX86-64-AVX.asmh>

; extern bool AvxBuildImageHistogram_(uint32_t* histo, const uint8_t* pixel_buff, uint32_t
num_pixels)
; Returns:      0 = invalid argument value, 1 = success

        .code
        extern c_NumPixelsMax:dword

AvxBuildImageHistogram_ proc frame
        _CreateFrame BIH_,1024,0,rbx,rsi,rdi
        _EndProlog

; Make sure num_pixels is valid
        xor eax,eax                                ;set error code
        test r8d,r8d
        jz Done
        cmp r8d,[c_NumPixelsMax]
        ja Done
        test r8d,1fh
        jnz Done                                  ;jump if num_pixels % 32 != 0

; Make sure histo & pixel_buff are properly aligned
        mov rsi,rcx                                ;rsi = ptr to histo

```

```

        test rsi,0fh
        jnz Done                                  ;jump if histo misaligned
        mov r9,rdx
        test r9,0fh
        jnz Done                                  ;jump if pixel_buff misaligned

; Initialize local histogram buffers (set all entries to zero)
        xor eax,eax
        mov rdi,rsi                                ;rdi = ptr to histo
        mov rcx,128                                ;rcx = size in qwords
        rep stosq                                   ;zero histo
        mov rdi,rbp                                ;rdi = ptr to histo2
        mov rcx,128                                ;rcx = size in qwords
        rep stosq                                   ;zero histo2

; Perform processing loop initializations
        shr r8d,5                                  ;number of pixel blocks (32 pixels/block)
        mov rdi,rbp                                ;ptr to histo2

; Build the histograms
        align 16                                    ;align jump target
@@:      vmovdqa xmm0,xmmword ptr [r9]              ;load pixel block
        vmovdqa xmm1,xmmword ptr [r9+16]           ;load pixel block

; Process pixels 0 - 3
        vpextrb rax,xmm0,0
        add dword ptr [rsi+rax*4],1                ;count pixel 0
        vpextrb rbx,xmm0,1
        add dword ptr [rdi+rbx*4],1                ;count pixel 1
        vpextrb rcx,xmm0,2
        add dword ptr [rsi+rcx*4],1                ;count pixel 2
        vpextrb rdx,xmm0,3
        add dword ptr [rdi+rdx*4],1                ;count pixel 3

; Process pixels 4 - 7
        vpextrb rax,xmm0,4
        add dword ptr [rsi+rax*4],1                ;count pixel 4
        vpextrb rbx,xmm0,5
        add dword ptr [rdi+rbx*4],1                ;count pixel 5
        vpextrb rcx,xmm0,6
        add dword ptr [rsi+rcx*4],1                ;count pixel 6
        vpextrb rdx,xmm0,7
        add dword ptr [rdi+rdx*4],1                ;count pixel 7

; Process pixels 8 - 11
        vpextrb rax,xmm0,8
        add dword ptr [rsi+rax*4],1                ;count pixel 8
        vpextrb rbx,xmm0,9
        add dword ptr [rdi+rbx*4],1                ;count pixel 9
        vpextrb rcx,xmm0,10
        add dword ptr [rsi+rcx*4],1                ;count pixel 10
        vpextrb rdx,xmm0,11
        add dword ptr [rdi+rdx*4],1                ;count pixel 11

; Process pixels 12 - 15
        vpextrb rax,xmm0,12
        add dword ptr [rsi+rax*4],1                ;count pixel 12
        vpextrb rbx,xmm0,13
        add dword ptr [rdi+rbx*4],1                ;count pixel 13
        vpextrb rcx,xmm0,14
        add dword ptr [rsi+rcx*4],1                ;count pixel 14
        vpextrb rdx,xmm0,15
        add dword ptr [rdi+rdx*4],1                ;count pixel 15

; Process pixels 16 - 19
        vpextrb rax,xmm1,0
        add dword ptr [rsi+rax*4],1                ;count pixel 16

```



```

vpextrb rbx, xmm1, 1
add dword ptr [rdi+rbx*4], 1      ;count pixel 17
vpextrb rcx, xmm1, 2
add dword ptr [rsi+rcx*4], 1      ;count pixel 18
vpextrb rdx, xmm1, 3
add dword ptr [rdi+rdx*4], 1      ;count pixel 19

; Process pixels 20 - 23
vpextrb rax, xmm1, 4
add dword ptr [rsi+rax*4], 1      ;count pixel 20
vpextrb rbx, xmm1, 5
add dword ptr [rdi+rbx*4], 1      ;count pixel 21
vpextrb rcx, xmm1, 6
add dword ptr [rsi+rcx*4], 1      ;count pixel 22
vpextrb rdx, xmm1, 7
add dword ptr [rdi+rdx*4], 1      ;count pixel 23

; Process pixels 24 - 27
vpextrb rax, xmm1, 8
add dword ptr [rsi+rax*4], 1      ;count pixel 24
vpextrb rbx, xmm1, 9
add dword ptr [rdi+rbx*4], 1      ;count pixel 25
vpextrb rcx, xmm1, 10
add dword ptr [rsi+rcx*4], 1      ;count pixel 26
vpextrb rdx, xmm1, 11
add dword ptr [rdi+rdx*4], 1      ;count pixel 27

; Process pixels 28 - 31
vpextrb rax, xmm1, 12
add dword ptr [rsi+rax*4], 1      ;count pixel 28
vpextrb rbx, xmm1, 13
add dword ptr [rdi+rbx*4], 1      ;count pixel 29
vpextrb rcx, xmm1, 14
add dword ptr [rsi+rcx*4], 1      ;count pixel 30
vpextrb rdx, xmm1, 15
add dword ptr [rdi+rdx*4], 1      ;count pixel 31

add r9, 32                        ;r9 = next pixel block
sub r8d, 1
jnz @B                            ;repeat loop if not done

; Merge intermediate histograms into final histogram
mov ecx, 32                      ;ecx = num iterations
xor eax, eax                     ;rax = common offset

@@: vmovdqa xmm0, xmmword ptr [rsi+rax]      ;load histo counts
    vmovdqa xmm1, xmmword ptr [rsi+rax+16]
    vpaddq xmm0, xmm0, xmmword ptr [rdi+rax] ;add counts from histo2
    vpaddq xmm1, xmm1, xmmword ptr [rdi+rax+16]
    vmovdqa xmmword ptr [rsi+rax], xmm0      ;save final result
    vmovdqa xmmword ptr [rsi+rax+16], xmm1

add rax, 32
sub ecx, 1
jnz @B
mov eax, 1                        ;set success return code

```

```

Done: _DeleteFrame rbx, rsi, rdi
ret
AvxBuildImageHistogram_ endp
end

```

```

;-----
;##### Ch07_08.asm
include <MacrosX86-64-AVX.asm>

```

```

; Image threshold data structure (see Ch07_08.h)
ITD struct

```

```

PbSrc          qword ?
PbMask         qword ?
NumPixels       dword ?
NumMaskedPixels dword ?
SumMaskedPixels dword ?
Threshold       byte ?
Pad            byte 3 dup(?)
MeanMaskedPixels real8 ?
ITD            ends

```

```

.const
align 16
PixelScale      byte 16 dup(80h)      ;uint8 to int8 scale value
CountPixelsMask byte 16 dup(01h)      ;mask to count pixels
R8_MinusOne     real8 -1.0            ;invalid mean value

```

```

.code
extern IsValid:proc

; extern "C" bool AvxThresholdImage_(ITD* itd);
; Returns:      0 = invalid size or unaligned image buffer, 1 = success

```

```

AvxThresholdImage_ proc frame
    _CreateFrame TI_, 0, 0, rbx
    _EndProlog

```

```

; Verify the arguments in the ITD structure
mov rbx, rcx                    ;copy itd ptr to non-volatile register
mov ecx, [rbx+ITD.NumPixels]    ;ecx = num_pixels
mov rdx, [rbx+ITD.PbSrc]        ;rdx = pb_src
mov r8, [rbx+ITD.PbMask]       ;r8 = pb_mask
sub rsp, 32                     ;allocate home area for IsValid
call IsValid                    ;validate args
or al, al
jz Done                         ;jump if invalid

```

```

; Initialize registers for processing loop
mov ecx, [rbx+ITD.NumPixels]    ;ecx = num_pixels
shr ecx, 6                      ;ecx = number of 64b pixel blocks
mov rdx, [rbx+ITD.PbSrc]        ;rdx = pb_src
mov r8, [rbx+ITD.PbMask]       ;r8 = pb_mask

```

```

movzx r9d, byte ptr [rbx+ITD.Threshold] ;r9d = threshold
vmovd xmm1, r9d                ;xmm1[7:0] = threshold
vpxor xmm0, xmm0, xmm0         ;mask for vpslufb
vpslufb xmm1, xmm1, xmm0       ;xmm1 = packed threshold

```

```

vmovdqa xmm4, xmmword ptr [PixelScale] ;packed pixel scale factor
vpsubb xmm5, xmm1, xmm4        ;scaled threshold

```

```

; Create the mask image
@@: vmovdqa xmm0, xmmword ptr [rdx]      ;original image pixels
    vpsubb xmm1, xmm0, xmm4              ;scaled image pixels
    vpcmpgtb xmm2, xmm1, xmm5            ;mask pixels
    vmovdqa xmmword ptr [r8], xmm2      ;save mask result

```

```

vmovdqa xmm0, xmmword ptr [rdx+16]
vpsubb xmm1, xmm0, xmm4
vpcmpgtb xmm2, xmm1, xmm5
vmovdqa xmmword ptr [r8+16], xmm2

```

```

vmovdqa xmm0, xmmword ptr [rdx+32]
vpsubb xmm1, xmm0, xmm4
vpcmpgtb xmm2, xmm1, xmm5
vmovdqa xmmword ptr [r8+32], xmm2

```

```

vmovdqa xmm0, xmmword ptr [rdx+48]

```

```

vpsubb xmm1, xmm0, xmm4
vpcmpgtb xmm2, xmm1, xmm5
vmovdqa xmmword ptr [r8+48], xmm2

add rdx, 64
add r8, 64           ;update pointers
sub ecx, 1           ;update counter
jnz @B              ;repeat until done

mov eax, 1           ;set success return code

Done: _DeleteFrame rbx
ret
AvxThresholdImage_ endp

; Macro _UpdateBlockSums

_UpdateBlockSums macro disp
vmovdqa xmm0, xmmword ptr [rdx+disp] ;xmm0 = 16 image pixels
vmovdqa xmm1, xmmword ptr [r8+disp] ;xmm1 = 16 mask pixels
vpand xmm2, xmm1, xmm8 ;xmm2 = 16 mask pixels (0x00 or 0x01)
vpaddb xmm6, xmm6, xmm2 ;update block num_masked_pixels
vpand xmm2, xmm0, xmm1 ;zero out unmasked image pixel
vpunpcklbw xmm3, xmm2, xmm9 ;promote image pixels from byte to word
vpunpckhbw xmm4, xmm2, xmm9
vpaddw xmm4, xmm4, xmm3
vpaddw xmm7, xmm7, xmm4 ;update block sum_mask_pixels
endm

; extern "C" bool AvxCalcImageMean_(ITD* itd);
; Returns: 0 = invalid image size or unaligned image buffer, 1 = success

AvxCalcImageMean_ proc frame
_CreateFrame C1M_, 0, 64, rbx
_SaveXmmRegs xmm6, xmm7, xmm8, xmm9
_EndProlog

; Verify the arguments in the ITD structure
mov rbx, rcx ;rbx = itd ptr
mov ecx, [rbx+ITD.NumPixels] ;ecx = num_pixels
mov rdx, [rbx+ITD.PbSrc] ;rdx = pb_src
mov r8, [rbx+ITD.PbMask] ;r8 = pb_mask
sub rsp, 32 ;allocate home area for IsValid
call IsValid ;validate args
or al, al
jz Done ;jump if invalid

; Initialize registers for processing loop
mov ecx, [rbx+ITD.NumPixels] ;ecx = num_pixels
shr ecx, 6 ;ecx = number of 64b pixel blocks
mov rdx, [rbx+ITD.PbSrc] ;rdx = pb_src
mov r8, [rbx+ITD.PbMask] ;r8 = pb_mask

vmovdqa xmm8, xmmword ptr [CountPixelsMask] ;mask for counting pixels
vpxor xmm9, xmm9, xmm9 ;xmm9 = packed zero

xor r10d, r10d ;r10d = num_masked_pixels (1 dword)
vpxor xmm5, xmm5, xmm5 ;sum_masked_pixels (4 dwords)
;Calculate num_mask_pixels and sum_mask_pixels
LP1: vpxor xmm6, xmm6, xmm6 ;num_masked_pixels_tmp (16 byte values)
vpxor xmm7, xmm7, xmm7 ;sum_masked_pixels_tmp (8 word values)
_UpdateBlockSums 0
_UpdateBlockSums 16
_UpdateBlockSums 32
_UpdateBlockSums 48

; Update num_masked_pixels
vpsrldq xmm0, xmm6, 8
vpaddb xmm6, xmm6, xmm0 ;num_mask_pixels_tmp (8 byte vals)
vpsrldq xmm0, xmm6, 4
vpaddb xmm6, xmm6, xmm0 ;num_mask_pixels_tmp (4 byte vals)
vpsrldq xmm0, xmm6, 2
vpaddb xmm6, xmm6, xmm0 ;num_mask_pixels_tmp (2 byte vals)
vpsrldq xmm0, xmm6, 1
vpaddb xmm6, xmm6, xmm0 ;num_mask_pixels_tmp (1 byte val)
vpextrb eax, xmm6, 0
add r10d, eax ;num_mask_pixels += num_mask_pixels_tmp

; Update sum_masked_pixels
vpunpcklwd xmm0, xmm7, xmm9 ;promote sum_mask_pixels_tmp to dwords
vpunpckhwd xmm1, xmm7, xmm9
vpaddd xmm5, xmm5, xmm0 ;sum_mask_pixels += sum_masked_pixels_tmp
vpaddd xmm5, xmm5, xmm1

add rdx, 64 ;update pb_src pointer
add r8, 64 ;update pb_mask pointer

sub rcx, 1 ;update loop counter
jnz LP1 ;repeat if not done

; Compute mean of masked pixels
vphaddq xmm0, xmm5, xmm5
vphaddq xmm1, xmm0, xmm0
vmovd eax, xmm1 ;eax = final sum_mask_pixels

test r10d, r10d ;is num_mask_pixels zero?
jz NoMean ;if yes, skip calc of mean
vcvtsi2sd xmm0, xmm0, eax ;xmm0 = sum_masked_pixels
vcvtsi2sd xmm1, xmm1, r10d ;xmm1 = num_masked_pixels
vdivsd xmm2, xmm0, xmm1 ;xmm2 = mean_masked_pixels
jmp @F

NoMean: vmovsd xmm2, [R8_MinusOne] ;use -1.0 for no mean

@@: mov [rbx+ITD.SumMaskedPixels], eax ;save sum_masked_pixels
mov [rbx+ITD.NumMaskedPixels], r10d ;save num_masked_pixels
vmovsd [rbx+ITD.MeanMaskedPixels], xmm2 ;save mean
mov eax, 1 ;set success return code

Done: _RestoreXmmRegs xmm6, xmm7, xmm8, xmm9
_DeleteFrame rbx
ret
AvxCalcImageMean_ endp
end

;-----
;##### Ch09_01.asm
;-----
; Mask values used to calculate floating-point absolute values
.const
AbsMaskF32 dword 8 dup(7fffffffh)
AbsMaskF64 qword 4 dup(7fffffffffffffffh)
; extern "C" void AvxPackedMathF32_(const YmmVal& a, const YmmVal& b, YmmVal c[8]);

.code
AvxPackedMathF32_ proc

; Load packed SP floating-point values
vmovaps ymm0, ymmword ptr [rcx] ;ymm0 = *a
vmovaps ymm1, ymmword ptr [rdx] ;ymm1 = *b

; Packed SP floating-point addition
vaddps ymm2, ymm0, ymm1
vmovaps ymmword ptr [r8], ymm2

```

```

; Packed SP floating-point subtraction
vsubps ymm2, ymm0, ymm1
vmovaps ymmword ptr [r8+32], ymm2

; Packed SP floating-point multiplication
vmulps ymm2, ymm0, ymm1
vmovaps ymmword ptr [r8+64], ymm2

; Packed SP floating-point division
vdivps ymm2, ymm0, ymm1
vmovaps ymmword ptr [r8+96], ymm2

; Packed SP floating-point absolute value (b)
vandps ymm2, ymm1, ymmword ptr [AbsMaskF32]
vmovaps ymmword ptr [r8+128], ymm2

; Packed SP floating-point square root (a)
vsqrtps ymm2, ymm0
vmovaps ymmword ptr [r8+160], ymm2

; Packed SP floating-point minimum
vminps ymm2, ymm0, ymm1
vmovaps ymmword ptr [r8+192], ymm2

; Packed SP floating-point maximum
vmaxps ymm2, ymm0, ymm1
vmovaps ymmword ptr [r8+224], ymm2

vzeroupper
ret
AvxPackedMathF32_ endp

; extern "C" void AvxPackedMathF64_(const YmmVal& a, const YmmVal& b, YmmVal c[8]);

```

AvxPackedMathF64\_ proc

```

; Load packed DP floating-point values
vmovapd ymm0, ymmword ptr [rcx] ;ymm0 = *a
vmovapd ymm1, ymmword ptr [rdx] ;ymm1 = *b

; Packed DP floating-point addition
vaddpd ymm2, ymm0, ymm1
vmovapd ymmword ptr [r8], ymm2

; Packed DP floating-point subtraction
vsubpd ymm2, ymm0, ymm1
vmovapd ymmword ptr [r8+32], ymm2

; Packed DP floating-point multiplication
vmulpd ymm2, ymm0, ymm1
vmovapd ymmword ptr [r8+64], ymm2

; Packed DP floating-point division
vdivpd ymm2, ymm0, ymm1
vmovapd ymmword ptr [r8+96], ymm2

; Packed DP floating-point absolute value (b)
vandpd ymm2, ymm1, ymmword ptr [AbsMaskF64]
vmovapd ymmword ptr [r8+128], ymm2

; Packed DP floating-point square root (a)
vsqrtpd ymm2, ymm0
vmovapd ymmword ptr [r8+160], ymm2

; Packed DP floating-point minimum
vminpd ymm2, ymm0, ymm1
vmovapd ymmword ptr [r8+192], ymm2

```

```

; Packed DP floating-point maximum
vmaxpd ymm2, ymm0, ymm1
vmovapd ymmword ptr [r8+224], ymm2

```

```

vzeroupper
ret
AvxPackedMathF64_ endp
end

```

```

;-----
##### Ch09_02.asm
include <cmpequ.asmh> include <MacrosX86-64-AVX.asmh>

```

```

.const
r4_3p0 real4 3.0
r4_4p0 real4 4.0

```

```

extern c_PI_F32:real4
extern c_QNaN_F32:real4

```

```

; extern "C" void AvxCalcSphereAreaVolume_(float* sa, float* vol, const float* r, size_t n);

```

```

.code
AvxCalcSphereAreaVolume_ proc frame
.CreateFrame CC_0, 64
.SaveXmmRegs xmm6, xmm7, xmm8, xmm9
.EndProlog

```

```

; Initialize
vbroadcastss ymm0, real4 ptr [r4_4p0] ;packed 4.0
vbroadcastss ymm1, real4 ptr [c_PI_F32] ;packed PI
vmulps ymm6, ymm0, ymm1 ;packed 4.0 * PI
vbroadcastss ymm7, real4 ptr [r4_3p0] ;packed 3.0
vbroadcastss ymm8, real4 ptr [c_QNaN_F32] ;packed QNaN
vxorps ymm9, ymm9, ymm9 ;packed 0.0

xor eax, eax ;common offset for arrays

cmp r9, 8
jb FinalR ;skip main loop if n < 8

```

```

; Calculate surface area and volume values using packed arithmetic
@@:

```

```

vmovdqa ymm0, ymmword ptr [r8+rax] ;load next 8 radii
vmulps ymm2, ymm6, ymm0 ;4.0 * PI * r
vmulps ymm3, ymm2, ymm0 ;4.0 * PI * r * r

vcmps ymm1, ymm0, ymm9, CMP_LT ;ymm1 = mask of radii < 0.0

vandps ymm4, ymm1, ymm8 ;set surface area to QNaN for radii < 0.0
vandnps ymm5, ymm1, ymm3 ;keep surface area for radii >= 0.0
vorps ymm5, ymm4, ymm5 ;final packed surface area
vmovaps ymmword ptr [rcx+rax], ymm5 ;save packed surface area

vmulps ymm2, ymm3, ymm0 ;4.0 * PI * r * r * r
vdivps ymm3, ymm2, ymm7 ;4.0 * PI * r * r * r / 3.0
vandps ymm4, ymm1, ymm8 ;set volume to QNaN for radii < 0.0
vandnps ymm5, ymm1, ymm3 ;keep volume for radii >= 0.0
vorps ymm5, ymm4, ymm5 ;final packed volume
vmovaps ymmword ptr [rdx+rax], ymm5 ;save packed volume

add rax, 32 ;rax = offset to next set of radii
sub r9, 8
cmp r9, 8
jae @B ;repeat until n < 8

```

```

; Perform final calculations using scalar arithmetic
FinalR: test r9, r9

```

```

jz Done                :skip loop of no more elements

@@:  vmovss xmm0,real4 ptr [r8+rax]
      vmulss xmm2,xmm6,xmm0      :4.0 * PI * r
      vmulss xmm3,xmm2,xmm0      :4.0 * PI * r * r

      vcmpss xmm1,xmm0,xmm9,CMP_LT

      vandps xmm4,xmm1,xmm8
      vandps xmm5,xmm1,xmm3
      vorps  xmm5,xmm4,xmm5
      vmovss real4 ptr[rcx+rax],xmm5      :save surface area

      vmulss xmm2,xmm3,xmm0      :4.0 * PI * r * r * r
      vdivss xmm3,xmm2,xmm7      :4.0 * PI * r * r * r / 3.0
      vandps xmm4,xmm1,xmm8
      vandps xmm5,xmm1,xmm3
      vorps  xmm5,xmm4,xmm5
      vmovss real4 ptr[rdx+rax],xmm5      :save volume

      add rax,4
      dec r9
      jnz @B                :repeat until done

Done:  vzeroupper

      _RestoreXmmRegs xmm6,xmm7,xmm8,xmm9
      _DeleteFrame
      ret

AvxCalcSphereAreaVolume_ endp
end

;-----
;##### Ch09_03.asm
;-----
; extern "C" bool AvxCalcColMeans_(const double* x, size_t nrows, size_t ncols, double* col_means)
; extern c_NumRowsMax:qword
; extern c_NumColsMax:qword

.code
AvxCalcColumnMeans_ proc

; Validate nrows and ncols
xor eax,eax                :error return code (also col_mean index)
test rdx,rdx
jz Done                    :jump if nrows is zero
cmp rdx,[c_NumRowsMax]
ja Done                    :jump if nrows is too large
test r8,r8
jz Done                    :jump if ncols is zero
cmp r8,[c_NumColsMax]
ja Done                    :jump if ncols is too large

; Initialize elements of col_means to zero
vxorpd xmm0,xmm0,xmm0      :xmm0[63:0] = 0.0
@@:  vmovsd real8 ptr[r9+rax*8],xmm0 :col_means[i] = 0.0
      inc rax
      cmp rax,r8
      jb @B                :repeat until done

      vcvtsi2sd xmm2,xmm2,rdx      :convert nrows for later use

; Compute the sum of each column in x
LP1:  mov r11,r9                :r11 = ptr to col_means
      xor r10,r10              :r10 = col_index

LP2:  mov rax,r10                :rax = col_index
      add rax,4

```

```

      cmp rax,r8                :4 or more columns remaining?
      ja @F                    :jump if no (col_index + 4 > ncols)

; Update col_means using next four columns
vmovupd ymm0,ymmword ptr [rcx] :load next 4 columns of current row
vaddpd ymm1,ymm0,ymmword ptr [r11] :add to col_means
vmovupd ymmword ptr [r11],ymm1 :save updated col_means
add r10,4                        :col_index += 4
add rcx,32                      :update x ptr
add r11,32                      :update col_means ptr
jmp NextColSet

@@:  sub rax,2
      cmp rax,r8                :2 or more columns remaining?
      ja @F                    :jump if no (col_index + 2 > ncols)

; Update col_means using next two columns
vmovupd xmm0,xmmword ptr [rcx] :load next 2 columns of current row
vaddpd xmm1,xmm0,xmmword ptr [r11] :add to col_means
vmovupd xmmword ptr [r11],xmm1 :save updated col_means
add r10,2                        :col_index += 2
add rcx,16                      :update x ptr
add r11,16                      :update col_means ptr
jmp NextColSet

; Update col_means using next column (or last column in the current row)
@@:  vmovsd xmm0,real8 ptr [rcx] :load x from last column
      vaddsd xmm1,xmm0,real8 ptr [r11] :add to col_means
      vmovsd real8 ptr [r11],xmm1 :save updated col_means
      inc r10                      :col_index += 1
      add rcx,8                    :update x ptr

NextColSet:
      cmp r10,r8                :more columns in current row?
      jb LP2                    :jump if yes
      dec rdx                    :nrows -= 1
      jnz LP1                    :jump if more rows

; Compute the final col_means
@@:  vmovsd xmm0,real8 ptr [r9]      :xmm0 = col_means[i]
      vdivsd xmm1,xmm0,xmm2          :compute final mean
      vmovsd real8 ptr [r9],xmm1     :save col_mean[i]
      add r9,8                      :update col_means ptr
      dec r8                        :ncols -= 1
      jnz @B                        :repeat until done

      mov eax,1                    :set success return code

Done:  vzeroupper
      ret

AvxCalcColumnMeans_ endp
end

;-----
;##### Ch09_04.asm
; include <MacrosX86-64-AVX.asmh>

;extern "C" bool AvxCalcCorrCoef(const double* x,const double* y,size_t n, double sums[5], double epsilon, double* rho)
; Returns 0 = error, 1 = success

.code
AvxCalcCorrCoef_ proc frame
      _CreateFrame CC_,0,32
      _SaveXmmRegs xmm6,xmm7
      _EndProlog

; Validate arguments
or r8,r8
jz BadArg                    :jump if n == 0

```

```

test rcx,1fh
jnz BadArg          ;jump if x is not aligned
test rdx,1fh
jnz BadArg          ;jump if y is not aligned

```

; Initialize sum variables to zero

```

vxorpd ymm3,ymm3,ymm3 ;ymm3 = packed sum_x
vxorpd ymm4,ymm4,ymm4 ;ymm4 = packed sum_y
vxorpd ymm5,ymm5,ymm5 ;ymm5 = packed sum_xx
vxorpd ymm6,ymm6,ymm6 ;ymm6 = packed sum_yy
vxorpd ymm7,ymm7,ymm7 ;ymm7 = packed sum_xy
mov r10,r8
; r10 = n

```

```

cmp r8,4
jb LP2          ;jump if n >= 1 && n <= 3

```

; Calculate intermediate packed sum variables

```

LP1: vmovapd ymm0,ymmword ptr [rcx] ;ymm0 = packed x values
vmovapd ymm1,ymmword ptr [rdx] ;ymm1 = packed y values

```

```

vaddpd ymm3,ymm3,ymm0 ;update packed sum_x
vaddpd ymm4,ymm4,ymm1 ;update packed sum_y

```

```

vmulpd ymm2,ymm0,ymm1 ;ymm2 = packed xy values
vaddpd ymm7,ymm7,ymm2 ;update packed sum_xy

```

```

vmulpd ymm0,ymm0,ymm0 ;ymm0 = packed xx values
vmulpd ymm1,ymm1,ymm1 ;ymm1 = packed yy values
vaddpd ymm5,ymm5,ymm0 ;update packed sum_xx
vaddpd ymm6,ymm6,ymm1 ;update packed sum_yy

```

```

add rcx,32 ;update x ptr
add rdx,32 ;update y ptr
sub r8,4 ;n -- 4
cmp r8,4 ;is n >= 4?
jae LP1 ;jump if yes

```

```

or r8,r8 ;is n == 0?
jz FSV ;jump if yes

```

; Update sum variables with final x & y values

```

LP2: vmovsd xmm0,real8 ptr [rcx] ;xmm0[63:0] = x[i], ymm0[255:64] = 0
vmovsd xmm1,real8 ptr [rdx] ;xmm1[63:0] = y[i], ymm1[255:64] = 0

```

```

vaddpd ymm3,ymm3,ymm0 ;update packed sum_x
vaddpd ymm4,ymm4,ymm1 ;update packed sum_y

```

```

vmulpd ymm2,ymm0,ymm1 ;ymm2 = packed xy values
vaddpd ymm7,ymm7,ymm2 ;update packed sum_xy

```

```

vmulpd ymm0,ymm0,ymm0 ;ymm0 = packed xx values
vmulpd ymm1,ymm1,ymm1 ;ymm1 = packed yy values
vaddpd ymm5,ymm5,ymm0 ;update packed sum_xx
vaddpd ymm6,ymm6,ymm1 ;update packed sum_yy

```

```

add rcx,8 ;update x ptr
add rdx,8 ;update y ptr
sub r8,1 ;n -- 1
jnz LP2 ;repeat until done

```

; Calculate final sum variables

```

FSV: vextractf128 xmm0,ymm3,1
vaddpd xmm1,xmm0,xmm3
vhaddpd xmm3,xmm1,xmm1 ;xmm3[63:0] = sum_x

```

```

vextractf128 xmm0,ymm4,1
vaddpd xmm1,xmm0,xmm4

```

```

vhaddpd xmm4,xmm1,xmm1 ;xmm4[63:0] = sum_y

```

```

vextractf128 xmm0,ymm5,1
vaddpd xmm1,xmm0,xmm5
vhaddpd xmm5,xmm1,xmm1 ;xmm5[63:0] = sum_xx

```

```

vextractf128 xmm0,ymm6,1
vaddpd xmm1,xmm0,xmm6
vhaddpd xmm6,xmm1,xmm1 ;xmm6[63:0] = sum_yy

```

```

vextractf128 xmm0,ymm7,1
vaddpd xmm1,xmm0,xmm7
vhaddpd xmm7,xmm1,xmm1 ;xmm7[63:0] = sum_xy

```

; Save final sum variables

```

vmovsd real8 ptr [r9],xmm3 ;save sum_x
vmovsd real8 ptr [r9+8],xmm4 ;save sum_y
vmovsd real8 ptr [r9+16],xmm5 ;save sum_xx
vmovsd real8 ptr [r9+24],xmm6 ;save sum_yy
vmovsd real8 ptr [r9+32],xmm7 ;save sum_xy

```

; Calculate rho numerator

```

; rho_num = n * sum_xy - sum_x * sum_y;
vcvtsi2sd xmm2,xmm2,r10 ;xmm2 = n
vmulsd xmm0,xmm2,xmm7 ;xmm0 = n * sum_xy
vmulsd xmm1,xmm3,xmm4 ;xmm1 = sum_x * sum_y
vsubsd xmm7,xmm0,xmm1 ;xmm7 = rho_num

```

; Calculate rho denominator

```

; t1 = sqrt(n * sum_xx - sum_x * sum_x)
; t2 = sqrt(n * sum_yy - sum_y * sum_y)
; rho_den = t1 * t2
vmulsd xmm0,xmm2,xmm5 ;xmm0 = n * sum_xx
vmulsd xmm3,xmm3,xmm3 ;xmm3 = sum_x * sum_x
vsubsd xmm3,xmm0,xmm3 ;xmm3 = n * sum_xx - sum_x * sum_x
vsqrtsd xmm3,xmm3,xmm3 ;xmm3 = t1

```

```

vmulsd xmm0,xmm2,xmm6 ;xmm0 = n * sum_yy
vmulsd xmm4,xmm4,xmm4 ;xmm4 = sum_y * sum_y
vsubsd xmm4,xmm0,xmm4 ;xmm4 = n * sum_yy - sum_y * sum_y
vsqrtsd xmm4,xmm4,xmm4 ;xmm4 = t2

```

```

vmulsd xmm0,xmm3,xmm4 ;xmm0 = rho_den

```

; Calculate and save final rho

```

xor eax,eax
vcomisd xmm0,real8 ptr [rbp+CC_OffsetStackArgs] ;rho_den < epsilon?
setae al ;set return code
jb BadRho ;jump if rho_den < epsilon
vdivsd xmm1,xmm7,xmm0 ;xmm1 = rho

```

```

SavRho: mov rdx,[rbp+CC_OffsetStackArgs+8] ;rdx = ptr to rho
vmovsd real8 ptr [rdx],xmm1 ;save rho

```

```

Done: vzeroupper
_RestoreXmmRegs xmm6,xmm7
_DeleteFrame
ret

```

; Error handling code

```

BadRho: vxorpd xmm1,xmm1,xmm1 ;rho = 0
jmp SavRho

```

```

BadArg: xor eax,eax ;eax = invalid arg ret code
jmp Done

```

AvxCalcCorrCoef\_ endp

```

end
:-----
:##### Ch09_05.asm
:include <MacrosX86-64-AVX.asmh>

;_Mat4x4TransposeF64 macro
;Description: This macro computes the transpose of a 4x4
:double-precision floating-point matrix.
;Input Matrix      Output Matrtix
:-----
:ymm0  a3 a2 a1 a0      ymm0  d0 c0 b0 a0
:ymm1  b3 b2 b1 b0      ymm1  d1 c1 b1 a1
:ymm2  c3 c2 c1 c0      ymm2  d2 c2 b2 a2
:ymm3  d3 d2 d1 d0      ymm3  d3 c3 b3 a3
:-----

;_Mat4x4TransposeF64 macro
:      vunpcklpd ymm4, ymm0, ymm1      ;ymm4 = b2 a2 b0 a0
:      vunpckhpd ymm5, ymm0, ymm1      ;ymm5 = b3 a3 b1 a1
:      vunpcklpd ymm6, ymm2, ymm3      ;ymm6 = d2 c2 d0 c0
:      vunpckhpd ymm7, ymm2, ymm3      ;ymm7 = d3 c3 d1 c1
:
:      vperm2f128 ymm0, ymm4, ymm6, 20h ;ymm0 = d0 c0 b0 a0
:      vperm2f128 ymm1, ymm5, ymm7, 20h ;ymm1 = d1 c1 b1 a1
:      vperm2f128 ymm2, ymm4, ymm6, 31h ;ymm2 = d2 c2 b2 a2
:      vperm2f128 ymm3, ymm5, ymm7, 31h ;ymm3 = d3 c3 b3 a3
:      endm

; extern "C" void AvxMat4x4TransposeF64_(double* m_des, const double* m_src1)
:code
AvxMat4x4TransposeF64_ proc frame
    _CreateFrame MT_, 0, 32
    _SaveXmmRegs xmm6, xmm7
    _EndProlog

; Transpose matrix m_src1
    vmovaps ymm0, [rdx]      ;ymm0 = m_src1.row_0
    vmovaps ymm1, [rdx+32]   ;ymm1 = m_src2.row_1
    vmovaps ymm2, [rdx+64]   ;ymm2 = m_src3.row_2
    vmovaps ymm3, [rdx+96]   ;ymm3 = m_src4.row_3

;_Mat4x4TransposeF64

    vmovaps [rcx], ymm0      ;save m_des.row_0
    vmovaps [rcx+32], ymm1   ;save m_des.row_1
    vmovaps [rcx+64], ymm2   ;save m_des.row_2
    vmovaps [rcx+96], ymm3   ;save m_des.row_3

    vzeroupper
Done: _RestoreXmmRegs xmm6, xmm7
    _DeleteFrame
    ret
AvxMat4x4TransposeF64_ endp

;_Mat4x4MulCalcRowF64 macro
;Description: This macro computes one row of a 4x4 matrix multiplication.
;Registers:   ymm0 = m_src2.row0
:             ymm1 = m_src2.row1
:             ymm2 = m_src2.row2
:             ymm3 = m_src2.row3
:             rcx = m_des ptr
:             rdx = m_src1 ptr
:             ymm4 - ymm4 = scratch registers

;_Mat4x4MulCalcRowF64 macro disp
    vbroadcastsd ymm4, real8 ptr [rdx+disp] ;broadcast m_src1[i][0]
    vbroadcastsd ymm5, real8 ptr [rdx+disp+8] ;broadcast m_src1[i][1]
    vbroadcastsd ymm6, real8 ptr [rdx+disp+16] ;broadcast m_src1[i][2]

```

```

    vbroadcastsd ymm7, real8 ptr [rdx+disp+24] ;broadcast m_src1[i][3]

    vmulpd ymm4, ymm4, ymm0      ;m_src1[i][0] * m_src2.row_0
    vmulpd ymm5, ymm5, ymm1      ;m_src1[i][1] * m_src2.row_1
    vmulpd ymm6, ymm6, ymm2      ;m_src1[i][2] * m_src2.row_2
    vmulpd ymm7, ymm7, ymm3      ;m_src1[i][3] * m_src2.row_3

    vaddpd ymm4, ymm4, ymm5      ;calc m_des.row_i
    vaddpd ymm6, ymm6, ymm7
    vaddpd ymm4, ymm4, ymm6

    vmovapd [rcx+disp], ymm4      ;save m_des.row_i
    endm

; extern "C" void AvxMat4x4MulF64_(double* m_des, const double* m_src1, const double* m_src2)
AvxMat4x4MulF64_ proc frame
    _CreateFrame MM_, 0, 32
    _SaveXmmRegs xmm6, xmm7
    _EndProlog

; Load m_src2 into YMM3:YMM0
    vmovapd ymm0, [r8]           ;ymm0 = m_src2.row_0
    vmovapd ymm1, [r8+32]        ;ymm1 = m_src2.row_1
    vmovapd ymm2, [r8+64]        ;ymm2 = m_src2.row_2
    vmovapd ymm3, [r8+96]        ;ymm3 = m_src2.row_3

; Compute matrix product
    _Mat4x4MulCalcRowF64 0      ;calculate m_des.row_0
    _Mat4x4MulCalcRowF64 32     ;calculate m_des.row_1
    _Mat4x4MulCalcRowF64 64     ;calculate m_des.row_2
    _Mat4x4MulCalcRowF64 96     ;calculate m_des.row_3

    vzeroupper
Done: _RestoreXmmRegs xmm6, xmm7
    _DeleteFrame
    ret
AvxMat4x4MulF64_ endp

:-----
:##### Ch09_06.asm
:include <MacrosX86-64-AVX.asmh>

; Custom segment for constants
ConstVals segment readonly align(32) 'const'
Mat4x4I real8 1.0, 0.0, 0.0, 0.0
        real8 0.0, 1.0, 0.0, 0.0
        real8 0.0, 0.0, 1.0, 0.0
        real8 0.0, 0.0, 0.0, 1.0

r8_SignBitMask qword 4 dup (8000000000000000h)
r8_AbsMask qword 4 dup (7fffffffffffffffh)
r8_1p0 real8 1.0
r8_N1p0 real8 -1.0
r8_N0p5 real8 -0.5
r8_N0p3333 real8 -0.3333333333333333
r8_N0p25 real8 -0.25
ConstVals ends
:code

;_Mat4x4TraceF64 macro
;Description: This macro contains instructions that compute the trace
:of the 4x4 double-precision floating-point matrix in ymm3:ymm0.

;_Max4x4TraceF64 macro
    vblendpd ymm0, ymm0, ymm1, 00000010b ;ymm0[127:0] = row 1,0 diag vals
    vblendpd ymm1, ymm2, ymm3, 00001000b ;ymm1[255:128] = row 3,2 diag vals
    vperm2f128 ymm2, ymm1, ymm1, 00000001b ;ymm2[127:0] = row 3,2 diag vals

```

```

vaddpd ymm3, ymm0, ymm2
vhaddpd ymm0, ymm3, ymm3          ;xmm0[63:0] = trace
endm

; extern "C" double Avx2Mat4x4TraceF64_(const double* m_src1)
; Description: The following function computes the trace of a
;              4x4 double-precision floating-point array.

Avx2Mat4x4TraceF64_ proc
    vmovapd ymm0, [rcx]             ;ymm0 = m_src1.row_0
    vmovapd ymm1, [rcx+32]          ;ymm1 = m_src1.row_1
    vmovapd ymm2, [rcx+64]          ;ymm2 = m_src1.row_2
    vmovapd ymm3, [rcx+96]          ;ymm3 = m_src1.row_3

    _Max4x4TraceF64                 ;xmm0[63:0] = m_src1.trace()
    vzeroupper
    ret
Avx2Mat4x4TraceF64_ endp

; _Mat4x4MulCalcRowF64 macro
; Description: This macro is used to compute one row of a 4x4 matrix
;              multiply.
; Registers:   ymm0 = m_src2.row0
;              ymm1 = m_src2.row1
;              ymm2 = m_src2.row2
;              ymm3 = m_src2.row3
;              ymm4 - ymm7 = scratch registers

_Mat4x4MulCalcRowF64 macro dreg, sreg, disp
    vbroadcastsd ymm4, real8 ptr [sreg+disp] ;broadcast m_src1[i][0]
    vbroadcastsd ymm5, real8 ptr [sreg+disp+8] ;broadcast m_src1[i][1]
    vbroadcastsd ymm6, real8 ptr [sreg+disp+16] ;broadcast m_src1[i][2]
    vbroadcastsd ymm7, real8 ptr [sreg+disp+24] ;broadcast m_src1[i][3]

    vmulpd ymm4, ymm4, ymm0          ;m_src1[i][0] * m_src2.row_0
    vmulpd ymm5, ymm5, ymm1          ;m_src1[i][1] * m_src2.row_1
    vmulpd ymm6, ymm6, ymm2          ;m_src1[i][2] * m_src2.row_2
    vmulpd ymm7, ymm7, ymm3          ;m_src1[i][3] * m_src2.row_3

    vaddpd ymm4, ymm4, ymm5          ;calc m_des.row_i
    vaddpd ymm6, ymm6, ymm7
    vaddpd ymm4, ymm4, ymm6
    vmovapd[dreg+disp], ymm4         ;save m_des.row_i
endm

; extern "C" void Avx2Mat4x4MulF64_(double* m_des, const double* m_src1, const double* m_src2)
Avx2Mat4x4MulF64_ proc frame
    _CreateFrame MM_, 0, 32
    _SaveXmmRegs xmm6, xmm7
    _EndProlog

    vmovapd ymm0, [r8]              ;ymm0 = m_src2.row_0
    vmovapd ymm1, [r8+32]            ;ymm1 = m_src2.row_1
    vmovapd ymm2, [r8+64]            ;ymm2 = m_src2.row_2
    vmovapd ymm3, [r8+96]            ;ymm3 = m_src2.row_3

    _Mat4x4MulCalcRowF64 rcx, rdx, 0 ;calculate m_des.row_0
    _Mat4x4MulCalcRowF64 rcx, rdx, 32 ;calculate m_des.row_1
    _Mat4x4MulCalcRowF64 rcx, rdx, 64 ;calculate m_des.row_2
    _Mat4x4MulCalcRowF64 rcx, rdx, 96 ;calculate m_des.row_3

    vzeroupper
    _RestoreXmmRegs xmm6, xmm7
    _DeleteFrame
    ret
Avx2Mat4x4MulF64_ endp

```

```

; extern "C" bool Avx2Mat4x4InvF64_(double* m_inv, const double* m, double epsilon, bool* is_singular);

; Offsets of intermediate matrices on stack relative to rsp
OffsetM2 equ 32
OffsetM3 equ 160
OffsetM4 equ 288

Avx2Mat4x4InvF64_ proc frame
    _CreateFrame M1_, 0, 160
    _SaveXmmRegs xmm6, xmm7, xmm8, xmm9, xmm10, xmm11, xmm12, xmm13, xmm14, xmm15
    _EndProlog

; Save args to home area for later use
    mov qword ptr [rbp+M1_OffsetHomeRCX], rcx ;save m_inv ptr
    mov qword ptr [rbp+M1_OffsetHomeRDX], rdx ;save m ptr
    vmovsd real8 ptr [rbp+M1_OffsetHomeR8], xmm2 ;save epsilon
    mov qword ptr [rbp+M1_OffsetHomeR9], r9 ;save is_singular ptr

; Allocate 384 bytes of stack space for temp matrices + 32 bytes for function calls
    and rsp, 0ffffff0h ;align rsp to 32-byte boundary
    sub rsp, 416 ;alloc stack space

; Calculate m2
    lea rcx, [rsp+OffsetM2] ;rcx = m2 ptr
    mov r8, rdx ;rdx, r8 = m ptr
    call Avx2Mat4x4MulF64_ ;calculate and save m2

; Calculate m3
    lea rcx, [rsp+OffsetM3] ;rcx = m3 ptr
    lea rdx, [rsp+OffsetM2] ;rdx = m2 ptr
    mov r8, [rbp+M1_OffsetHomeRDX] ;r8 = m
    call Avx2Mat4x4MulF64_ ;calculate and save m3

; Calculate m4
    lea rcx, [rsp+OffsetM4] ;rcx = m4 ptr
    lea rdx, [rsp+OffsetM3] ;rdx = m3 ptr
    mov r8, [rbp+M1_OffsetHomeRDX] ;r8 = m
    call Avx2Mat4x4MulF64_ ;calculate and save m4

; Calculate trace of m, m2, m3, and m4
    mov rcx, [rbp+M1_OffsetHomeRDX]
    call Avx2Mat4x4TraceF64_
    vmovsd xmm8, xmm0 ;xmm8 = t1

    lea rcx, [rsp+OffsetM2]
    call Avx2Mat4x4TraceF64_
    vmovsd xmm9, xmm9, xmm0 ;xmm9 = t2

    lea rcx, [rsp+OffsetM3]
    call Avx2Mat4x4TraceF64_
    vmovsd xmm10, xmm10, xmm0 ;xmm10 = t3

    lea rcx, [rsp+OffsetM4]
    call Avx2Mat4x4TraceF64_
    vmovsd xmm11, xmm11, xmm0 ;xmm10 = t4

; Calculate the required coefficients
; c1 = -t1;
; c2 = -1.0f / 2.0f * (c1 * t1 + t2);
; c3 = -1.0f / 3.0f * (c2 * t1 + c1 * t2 + t3);
; c4 = -1.0f / 4.0f * (c3 * t1 + c2 * t2 + c1 * t3 + t4);
; Registers used:
;   t1-t4 = xmm8-xmm11
;   c1-c4 = xmm12-xmm15

vxorpd xmm12, xmm8, real8 ptr [r8_SignBitMask] ;xmm12 = c1

```



```

vmulsd xmm13, xmm12, xmm8      :c1 * t1
vaddsd xmm13, xmm13, xmm9      :c1 * t1 + t2
vmulsd xmm13, xmm13, [r8_N0p5] :c2

vmulsd xmm14, xmm13, xmm8      :c2 * t1
vmulsd xmm0, xmm12, xmm9       :c1 * t2
vaddsd xmm14, xmm14, xmm0       :c2 * t1 + c1 * t2
vaddsd xmm14, xmm14, xmm10      :c2 * t1 + c1 * t2 + t3
vmulsd xmm14, xmm14, [r8_N0p3333] :c3

vmulsd xmm15, xmm14, xmm8      :c3 * t1
vmulsd xmm0, xmm13, xmm9       :c2 * t2
vmulsd xmm1, xmm12, xmm10      :c1 * t3
vaddsd xmm2, xmm0, xmm1        :c2 * t2 + c1 * t3
vaddsd xmm15, xmm15, xmm2       :c3 * t1 + c2 * t2 + c1 * t3
vaddsd xmm15, xmm15, xmm11      :c3 * t1 + c2 * t2 + c1 * t3 + t4
vmulsd xmm15, xmm15, [r8_N0p25] :c4

; Make sure matrix is not singular
vandpd xmm0, xmm15, [r8_AbsMask] ;compute fabs(c4)
vmovsd xmm1, real8 ptr [rbp+M1_OffsetHomeR8]
vcomisd xmm0, real8 ptr [rbp+M1_OffsetHomeR8] ;compare against epsilon
setp al ;set al = if unordered
setb ah ;set ah = if fabs(c4) < epsilon
or al, ah ;al = is_singular
mov rcx, [rbp+M1_OffsetHomeR9] ;rax = is_singular ptr
mov [rcx], al ;save is_singular state
jnz Error ;jump if singular

; Calculate m_inv = -1.0 / c4 * (m3 + c1 * m2 + c2 * m1 + c3 * I)
vbroadcastsd ymm14, xmm14 ;ymm14 = packed c3
lea rcx, [Mat4x4I] ;rcx = I ptr
vmulpd ymm0, ymm14, ymmword ptr [rcx]
vmulpd ymm1, ymm14, ymmword ptr [rcx+32]
vmulpd ymm2, ymm14, ymmword ptr [rcx+64]
vmulpd ymm3, ymm14, ymmword ptr [rcx+96] ;c3 * I

vbroadcastsd ymm13, xmm13 ;ymm13 = packed c2
mov rcx, [rbp+M1_OffsetHomeRDX] ;rcx = m ptr
vmulpd ymm4, ymm13, ymmword ptr [rcx]
vmulpd ymm5, ymm13, ymmword ptr [rcx+32]
vmulpd ymm6, ymm13, ymmword ptr [rcx+64]
vmulpd ymm7, ymm13, ymmword ptr [rcx+96] ;c2 * m1

vaddpd ymm0, ymm0, ymm4
vaddpd ymm1, ymm1, ymm5
vaddpd ymm2, ymm2, ymm6
vaddpd ymm3, ymm3, ymm7 ;c2 * m1 + c3 * I

vbroadcastsd ymm12, xmm12 ;ymm12 = packed c1
lea rcx, [rsp+OffsetM2] ;rcx = m2 ptr
vmulpd ymm4, ymm12, ymmword ptr [rcx]
vmulpd ymm5, ymm12, ymmword ptr [rcx+32]
vmulpd ymm6, ymm12, ymmword ptr [rcx+64]
vmulpd ymm7, ymm12, ymmword ptr [rcx+96] ;c1 * m2

vaddpd ymm0, ymm0, ymm4
vaddpd ymm1, ymm1, ymm5
vaddpd ymm2, ymm2, ymm6
vaddpd ymm3, ymm3, ymm7 ;c1 * m2 + c2 * m1 + c3 * I

lea rcx, [rsp+OffsetM3] ;rcx = m3 ptr
vaddpd ymm0, ymm0, ymmword ptr [rcx]
vaddpd ymm1, ymm1, ymmword ptr [rcx+32]
vaddpd ymm2, ymm2, ymmword ptr [rcx+64]
vaddpd ymm3, ymm3, ymmword ptr [rcx+96] ;m3 + c1 * m2 + c2 * m1 + c3 * I

vmovsd xmm4, [r8_N1p0]
vdvsd xmm4, xmm4, xmm15 ;xmm4 = -1.0 / c4

```

```

vbroadcastsd ymm4, xmm4
vmulpd ymm0, ymm0, ymm4
vmulpd ymm1, ymm1, ymm4
vmulpd ymm2, ymm2, ymm4
vmulpd ymm3, ymm3, ymm4 ;ymm3:ymm0 = m_inv

; Save m_inv
mov rcx, [rbp+M1_OffsetHomeRCX]
vmovapd ymmword ptr [rcx], ymm0
vmovapd ymmword ptr [rcx+32], ymm1
vmovapd ymmword ptr [rcx+64], ymm2
vmovapd ymmword ptr [rcx+96], ymm3
mov eax, 1 ;set success return code

Done: vzeroupper
;RestoreXmmRegs xmm6, xmm7, xmm8, xmm9, xmm10, xmm11, xmm12, xmm13, xmm14, xmm15
_DeleteFrame
ret

Error: xor eax, eax
jmp Done

Avx2Mat4x4InvF64_ endp
end
;-----
;##### Ch09_07.asm
;-----
; extern "C" void AvxBleedF32_(YmmVal* des1, YmmVal* src1, YmmVal* src2, YmmVal* idx1)
;code
AvxBleedF32_ proc
vmovaps ymm0, ymmword ptr [rdx] ;ymm0 = src1
vmovaps ymm1, ymmword ptr [r8] ;ymm1 = src2
vmovdqa ymm2, ymmword ptr [r9] ;ymm2 = idx1
vblendvps ymm3, ymm0, ymm1, ymm2 ;blend ymm0 & ymm1, ymm2 "indices"
vmovaps ymmword ptr [rcx], ymm3 ;Save result to des1

vzeroupper
ret
AvxBleedF32_ endp

; extern "C" void Avx2PermuteF32_(YmmVal* des1, YmmVal* src1, YmmVal* idx1, YmmVal* des2, YmmVal* src2, YmmVal* idx2)
Avx2PermuteF32_ proc
; Perform vpermpps permutation
vmovaps ymm0, ymmword ptr [rdx] ;ymm0 = src1
vmovdqa ymm1, ymmword ptr [r8] ;ymm1 = idx1
vpermpps ymm2, ymm1, ymm0 ;permute ymm0 using ymm1 indices
vmovaps ymmword ptr [rcx], ymm2 ;save result to des1

; Perform vpermilps permutation
mov rdx, [rsp+40] ;rdx = src2 ptr
mov r8, [rsp+48] ;r8 = idx2 ptr
vmovaps ymm3, ymmword ptr [rdx] ;ymm3 = src2
vmovdqa ymm4, ymmword ptr [r8] ;ymm4 = idx1
vpermilps ymm5, ymm3, ymm4 ;permute ymm3 using ymm4 indices
vmovaps ymmword ptr [r9], ymm5 ;save result to des2

vzeroupper
ret
Avx2PermuteF32_ endp
end
;-----
;##### Ch09_08.asm
;-----
; For each of the following functions, the contents of y are loaded
; into ymm0 prior to execution of the vgatherXXX instruction in order to
; demonstrate the effects of conditional merging.

```

```

.code
; extern "C" void Avx2Gather8xF32_I32_(float* y, const float* x, const int32_t* indices, const
int32_t* merge)
Avx2Gather8xF32_I32_ proc
    vmovups ymm0,ymmword ptr [rcx]    ;ymm0 = y[7]:y[0]
    vmovdqu ymm1,ymmword ptr [r8]     ;ymm1 = indices[7]:indices[0]
    vmovdqu ymm2,ymmword ptr [r9]     ;ymm2 = merge[7]:merge[0]
    vpslld ymm2,ymm2,31               ;shift merge vals to high-order bits
    vgatherdps ymm0,[rdx+ymm1*4],ymm2  ;ymm0 = gathered elements
    vmovups ymmword ptr [rcx],ymm0    ;save gathered elements

    vzeroupper
    ret
Avx2Gather8xF32_I32_ endp

; extern "C" void Avx2Gather8xF32_I64_(float* y, const float* x, const int64_t* indices, const
int32_t* merge)
Avx2Gather8xF32_I64_ proc
    vmovups xmm0,xmmword ptr [rcx]    ;xmm0 = y[3]:y[0]
    vmovdqu ymm1,ymmword ptr [r8]     ;ymm1 = indices[3]:indices[0]
    vmovdqu xmm2,xmmword ptr [r9]     ;xmm2 = merge[3]:merge[0]
    vpslld xmm2,xmm2,31               ;shift merge vals to high-order bits
    vgatherqps xmm0,[rdx+ymm1*4],xmm2  ;xmm0 = gathered elements
    vmovups xmmword ptr [rcx],xmm0    ;save gathered elements

    vmovups xmm3,xmmword ptr [rcx+16]  ;xmm0 = des[7]:des[4]
    vmovdqu ymm1,ymmword ptr [r8+32]  ;ymm1 = indices[7]:indices[4]
    vmovdqu xmm2,xmmword ptr [r9+16]  ;xmm2 = merge[7]:merge[4]
    vpslld xmm2,xmm2,31               ;shift merge vals to high-order bits
    vgatherqps xmm3,[rdx+ymm1*4],xmm2  ;xmm0 = gathered elements
    vmovups xmmword ptr [rcx+16],xmm3  ;save gathered elements

    vzeroupper
    ret
Avx2Gather8xF32_I64_ endp

; extern "C" void Avx2Gather8xF64_I32_(double* y, const double* x, const int32_t* indices, const
int64_t* merge)
Avx2Gather8xF64_I32_ proc
    vmovupd ymm0,ymmword ptr [rcx]    ;ymm0 = y[3]:y[0]
    vmovdqu xmm1,xmmword ptr [r8]     ;xmm1 = indices[3]:indices[0]
    vmovdqu ymm2,ymmword ptr [r9]     ;ymm2 = merge[3]:merge[0]
    vpsllq ymm2,ymm2,63               ;shift merge vals to high-order bits
    vgatherdpd ymm0,[rdx+xmm1*8],ymm2  ;ymm0 = gathered elements
    vmovupd ymmword ptr [rcx],ymm0    ;save gathered elements

    vmovupd ymm0,ymmword ptr [rcx+32]  ;ymm0 = y[7]:y[4]
    vmovdqu xmm1,xmmword ptr [r8+16]  ;xmm1 = indices[7]:indices[4]
    vmovdqu ymm2,ymmword ptr [r9+32]  ;ymm2 = merge[7]:merge[4]
    vpsllq ymm2,ymm2,63               ;shift merge vals to high-order bits
    vgatherdpd ymm0,[rdx+xmm1*8],ymm2  ;ymm0 = gathered elements
    vmovupd ymmword ptr [rcx+32],ymm0  ;save gathered elements

    vzeroupper
    ret
Avx2Gather8xF64_I32_ endp

; extern "C" void Avx2Gather8xF64_I64_(double* y, const double* x, const int64_t* indices, const
int64_t* merge)
Avx2Gather8xF64_I64_ proc
    vmovupd ymm0,ymmword ptr [rcx]    ;ymm0 = y[3]:y[0]
    vmovdqu xmm1,xmmword ptr [r8]     ;ymm1 = indices[3]:indices[0]
    vmovdqu ymm2,ymmword ptr [r9]     ;ymm2 = merge[3]:merge[0]
    vpsllq ymm2,ymm2,63               ;shift merge vals to high-order bits
    vgatherdpd ymm0,[rdx+ymm1*8],ymm2  ;ymm0 = gathered elements
    vmovupd ymmword ptr [rcx],ymm0    ;save gathered elements

    vzeroupper
    ret
Avx2Gather8xF64_I64_ endp

```

```

vmovupd ymm0,ymmword ptr [rcx+32]    ;ymm0 = y[7]:y[4]
vmovdqu ymm1,ymmword ptr [r8+32]     ;ymm1 = indices[7]:indices[4]
vmovdqu ymm2,ymmword ptr [r9+32]     ;ymm2 = merge[7]:merge[4]
vpsllq ymm2,ymm2,63                  ;shift merge vals to high-order bits
vgatherqpd ymm0,[rdx+ymm1*8],ymm2    ;ymm0 = gathered elements
vmovupd ymmword ptr [rcx+32],ymm0    ;save gathered elements

```

```

vzeroupper
ret
Avx2Gather8xF64_I64_ endp
end

```

```

;-----
;##### Ch10_01.asm
;-----

```

```

; extern "C" void Avx2PackedMathI16_(const YmmVal& a, const YmmVal& b, YmmVal c[6])
.code

```

```

Avx2PackedMathI16_ proc

```

```

; Load values a and b, which must be properly aligned

```

```

    vmovdqa ymm0,ymmword ptr [rcx]    ;ymm0 = a
    vmovdqa ymm1,ymmword ptr [rdx]    ;ymm1 = b

```

```

; Perform packed arithmetic operations

```

```

    vpaddw ymm2,ymm0,ymm1             ;add
    vmovdqa ymmword ptr [r8],ymm2     ;save vpaddw result

```

```

    vpaddsw ymm2,ymm0,ymm1            ;add with signed saturation
    vmovdqa ymmword ptr [r8+32],ymm2  ;save vpaddsw result

```

```

    vpsubw ymm2,ymm0,ymm1             ;sub
    vmovdqa ymmword ptr [r8+64],ymm2  ;save vpsubw result

```

```

    vpsubsw ymm2,ymm0,ymm1            ;sub with signed saturation
    vmovdqa ymmword ptr [r8+96],ymm2  ;save vpsubsw result

```

```

    vpminsw ymm2,ymm0,ymm1            ;signed minimums
    vmovdqa ymmword ptr [r8+128],ymm2 ;save vpminsw result

```

```

    vpmasw ymm2,ymm0,ymm1             ;signed maximums
    vmovdqa ymmword ptr [r8+160],ymm2 ;save vpmasw result

```

```

    vzeroupper
    ret
Avx2PackedMathI16_ endp

```

```

; extern "C" void Avx2PackedMathI32_(const YmmVal& a, const YmmVal& b, YmmVal c[6])

```

```

Avx2PackedMathI32_ proc

```

```

; Load values a and b, which must be properly aligned

```

```

    vmovdqa ymm0,ymmword ptr [rcx]    ;ymm0 = a
    vmovdqa ymm1,ymmword ptr [rdx]    ;ymm1 = b

```

```

; Perform packed arithmetic operations

```

```

    vpaddd ymm2,ymm0,ymm1             ;add
    vmovdqa ymmword ptr [r8],ymm2     ;save vpaddd result

```

```

    vpsubd ymm2,ymm0,ymm1             ;sub
    vmovdqa ymmword ptr [r8+32],ymm2  ;save vpsubd result

```

```

    vpmulld ymm2,ymm0,ymm1            ;signed mul (low 32 bits)
    vmovdqa ymmword ptr [r8+64],ymm2  ;save vpmulld result

```

```

    vpslld ymm2,ymm0,ymm1             ;shift left logical
    vmovdqa ymmword ptr [r8+96],ymm2  ;save vpslld result

```

```

    vpsravd ymm2,ymm0,ymm1            ;shift right arithmetic
    vmovdqa ymmword ptr [r8+128],ymm2 ;save vpsravd result

```

```

    vpbabsd ymm2,ymm0                ;absolute value

```

```

        vmovdqa ymmword ptr [r8+160],ymm2    ;save vpabsd result

        vzeroupper
        ret
Avx2PackedMathI32_ endp
end;

-----
;##### Ch10_02.asm
;-----
; extern "C" YmmValI2 Avx2UnpackU32_U64_(const YmmVal& a, const YmmVal& b);

        .code
Avx2UnpackU32_U64_ proc

; Load argument values
        vmovdqa ymm0,ymmword ptr [rdx]    ;ymm0 = a
        vmovdqa ymm1,ymmword ptr [r8]     ;ymm1 = b

; Perform dword to qword unpacks
        vpunpckldq ymm2,ymm0,ymm1        ;unpack low doublewords
        vpunpckhdq ymm3,ymm0,ymm1        ;unpack high doublewords

; Save result to YmmVal2 buffer
        vmovdqa ymmword ptr [rcx],ymm2    ;save low result
        vmovdqa ymmword ptr [rcx+32],ymm3 ;save high result

        mov rax,rcx                      ;rax = ptr to YmmVal2

        vzeroupper
        ret
Avx2UnpackU32_U64_ endp

; extern "C" void Avx2PackI32_I16_(const YmmVal& a, const YmmVal& b, YmmVal* c);

Avx2PackI32_I16_ proc

; Load argument values
        vmovdqa ymm0,ymmword ptr [rcx]    ;ymm0 = a
        vmovdqa ymm1,ymmword ptr [rdx]    ;ymm1 = b

; Perform pack dword to word with signed saturation
        vpackssdw ymm2,ymm0,ymm1          ;ymm2 = packed words
        vmovdqa ymmword ptr [r8],ymm2     ;save result

        vzeroupper
        ret
Avx2PackI32_I16_ endp

Foo1_ proc
        ret
Foo1_ endp
end

;-----
;##### Ch10_03.asm
;-----
; extern "C" void Avx2ZeroExtU8_U16_(YmmVal*a, YmmVal b[2]);

        .code
Avx2ZeroExtU8_U16_ proc
        vpmovzxbw ymm0,xmmword ptr [rcx]    ;zero extend a[0] - a[15]
        vpmovzxbw ymm1,xmmword ptr [rcx+16] ;zero extend a[16] - a[31]

        vmovdqa ymmword ptr [rdx],ymm0     ;save results
        vmovdqa ymmword ptr [rdx+32],ymm1

        vzeroupper
        ret

```

```

Avx2ZeroExtU8_U16_ endp

; extern "C" void Avx2ZeroExtU8_U32_(YmmVal*a, YmmVal b[4]);

Avx2ZeroExtU8_U32_ proc
        vpmovzxbd ymm0,qword ptr [rcx]      ;zero extend a[0] - a[7]
        vpmovzxbd ymm1,qword ptr [rcx+8]    ;zero extend a[8] - a[15]
        vpmovzxbd ymm2,qword ptr [rcx+16]   ;zero extend a[16] - a[23]
        vpmovzxbd ymm3,qword ptr [rcx+24]   ;zero extend a[24] - a[31]

        vmovdqa ymmword ptr [rdx],ymm0     ;save results
        vmovdqa ymmword ptr [rdx+32],ymm1
        vmovdqa ymmword ptr [rdx+64],ymm2
        vmovdqa ymmword ptr [rdx+96],ymm3

        vzeroupper
        ret
Avx2ZeroExtU8_U32_ endp

; extern "C" void Avx2SignExtI16_I32_(YmmVal*a, YmmVal b[2])
Avx2SignExtI16_I32_ proc
        vpmovsxdw ymm0,xmmword ptr [rcx]    ;sign extend a[0] - a[7]
        vpmovsxdw ymm1,xmmword ptr [rcx+16] ;sign extend a[8] - a[15]

        vmovdqa ymmword ptr [rdx],ymm0     ;save results
        vmovdqa ymmword ptr [rdx+32],ymm1

        vzeroupper
        ret
Avx2SignExtI16_I32_ endp

; extern "C" void Avx2SignExtI16_I64_(YmmVal*a, YmmVal b[4])
Avx2SignExtI16_I64_ proc
        vpmovsxwq ymm0,qword ptr [rcx]      ;sign extend a[0] - a[3]
        vpmovsxwq ymm1,qword ptr [rcx+8]    ;sign extend a[4] - a[7]
        vpmovsxwq ymm2,qword ptr [rcx+16]   ;sign extend a[8] - a[11]
        vpmovsxwq ymm3,qword ptr [rcx+24]   ;sign extend a[12] - a[15]

        vmovdqa ymmword ptr [rdx],ymm0     ;save results
        vmovdqa ymmword ptr [rdx+32],ymm1
        vmovdqa ymmword ptr [rdx+64],ymm2
        vmovdqa ymmword ptr [rdx+96],ymm3

        vzeroupper
        ret
Avx2SignExtI16_I64_ endp
end

;-----
;##### Ch10_04.asm
;-----
; The following structure must match the structure that's declared in the file .h file
ClipData struct
Src          qword ?      ;source buffer pointer
Des          qword ?      ;destination buffer pointer
NumPixels    qword ?      ;number of pixels
NumClippedPixels qword ?  ;number of clipped pixels
ThreshLo     byte ?       ;low threshold
ThreshHi     byte ?       ;high threshold
ClipData     ends

; extern "C" bool Avx2ClipPixels_(ClipData* cd)
        .code
Avx2ClipPixels_ proc

; Load and validate arguments
        xor eax,eax                ;set error return code
        xor r8d,r8d                ;r8 = number of clipped pixels

```

```

mov rdx, [rcx+ClipData.NumPixels]    ;rdx = num_pixels
or rdx, rdx
jz Done                               ;jump of num_pixels is zero
test rdx, 1fh
jnz Done                               ;jump if num_pixels % 32 != 0

mov r10, [rcx+ClipData.Src]          ;r10 = Src
test r10, 1fh
jnz Done                               ;jump if Src is misaligned

mov r11, [rcx+ClipData.Des]          ;r11 = Des
test r11, 1fh
jnz Done                               ;jump if Des is misaligned

; Create packed thresh_lo and thresh_hi data values
vpbroadcastb ymm4, [rcx+ClipData.ThreshLo] ;ymm4 = packed thresh_lo
vpbroadcastb ymm5, [rcx+ClipData.ThreshHi] ;ymm5 = packed thresh_hi

; Clip pixels to threshold values
@@: vmovdqa ymm0, ymmword ptr [r10] ;ymm0 = 32 pixels
vpmmaxub ymm1, ymm0, ymm4           ;clip to thresh_lo
vpminub ymm2, ymm1, ymm5           ;clip to thresh_hi
vmovdqa ymmword ptr [r11], ymm2    ;save clipped pixels

; Count number of clipped pixels
vpcmpeqb ymm3, ymm2, ymm0           ;compare clipped pixels to original
vpmovmskb eax, ymm3                 ;eax = mask of non-clipped pixels
not eax                             ;eax = mask of clipped pixels
popcnt eax, eax                     ;eax = number of clipped pixels
add r8, rax                          ;update clipped pixel count

; Update pointers and loop counter
add r10, 32                         ;update Src ptr
add r11, 32                         ;update Des ptr
sub rdx, 32                         ;update loop counter
jnz @B                               ;repeat if not done

mov eax, 1                          ;set success return code

; Save num_clipped_pixels
Done: mov [rcx+ClipData.NumClippedPixels], r8 ;save num_clipped_pixels
vzeroupper
ret

Avx2ClipPixels_ endp
end
;-----
;##### Ch10_05.asm
include <MacrosX86-64-AVX.asm>

; 256-bit wide constants
ConstVals segment readonly align(32) 'const'
InitialPminVal db 32 dup(0ffh)
InitialPmaxVal db 32 dup(00h)
ConstVals ends

; Macro _YmmVpextrMinub
; This macro generates code that extracts the smallest unsigned byte from register YmmSrc.

_YmmVpextrMinub macro GprDes, YmmSrc, YmmTmp

; Make sure YmmSrc and YmmTmp are different
erridni <YmmSrc>, <YmmTmp>, <Invalid registers>

; Construct text strings for the corresponding XMM registers

```

```

YmmSrcSuffix SUBSTR <YmmSrc>, 2
XmmSrc CATSTR <X>, YmmSrcSuffix

YmmTmpSuffix SUBSTR <YmmTmp>, 2
XmmTmp CATSTR <X>, YmmTmpSuffix

; Reduce the 32 byte values in YmmSrc to the smallest value
vextracti128 XmmTmp, YmmSrc, 1
vpminub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 16 min values

vpsrldq XmmTmp, XmmSrc, 8
vpminub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 8 min values

vpsrldq XmmTmp, XmmSrc, 4
vpminub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 4 min values

vpsrldq XmmTmp, XmmSrc, 2
vpminub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 2 min values

vpsrldq XmmTmp, XmmSrc, 1
vpminub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 1 min value

vpextrb GprDes, XmmSrc, 0           ;mov final min value to Gpr
endm

; Macro _YmmVpextrMaxub
; This macro generates code that extracts the largest unsigned byte from register YmmSrc.

_YmmVpextrMaxub macro GprDes, YmmSrc, YmmTmp

; Make sure YmmSrc and YmmTmp are different
erridni <YmmSrc>, <YmmTmp>, <Invalid registers>

; Construct text strings for the corresponding XMM registers
YmmSrcSuffix SUBSTR <YmmSrc>, 2
XmmSrc CATSTR <X>, YmmSrcSuffix

YmmTmpSuffix SUBSTR <YmmTmp>, 2
XmmTmp CATSTR <X>, YmmTmpSuffix

; Reduce the 32 byte values in YmmSrc to the largest value
vextracti128 XmmTmp, YmmSrc, 1
vpmaxub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 16 max values

vpsrldq XmmTmp, XmmSrc, 8
vpmaxub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 8 max values

vpsrldq XmmTmp, XmmSrc, 4
vpmaxub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 4 max values

vpsrldq XmmTmp, XmmSrc, 2
vpmaxub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 2 max values

vpsrldq XmmTmp, XmmSrc, 1
vpmaxub XmmSrc, XmmSrc, XmmTmp      ;XmmSrc = final 1 max value

vpextrb GprDes, XmmSrc, 0           ;mov final max value to Gpr
endm

; extern "C" bool Avx2CalcRgbMinMax_(uint8_t* rgb[3], size_t num_pixels, uint8_t min_vals[3],
uint8_t max_vals[3])
.code
Avx2CalcRgbMinMax_ proc frame
_CreateFrame CalcMinMax_, 0, 48, r12
_SaveXmmRegs xmm6, xmm7, xmm8
_EndProlog

```

```

; Make sure num_pixels and the color plane arrays are valid
xor eax, eax                ;set error code

test rdx, rdx
jz Done                    ;jump if num_pixels == 0
test rdx, 01fh
jnz Done                   ;jump if num_pixels % 32 != 0

mov r10, [rcx]              ;r10 = color plane R
test r10, 1fh
jnz Done                   ;jump if color plane R is not aligned

mov r11, [rcx+8]            ;r11 = color plane G
test r11, 1fh
jnz Done                   ;jump if color plane G is not aligned

mov r12, [rcx+16]           ;r12 = color plane B
test r12, 1fh
jnz Done                   ;jump if color plane B is not aligned

; Initialize the processing loop registers
vmovdqa ymm3, ymmword ptr [InitialPminVal] ;ymm3 = R minimums
vmovdqa ymm4, ymm3         ;ymm4 = G minimums
vmovdqa ymm5, ymm3         ;ymm5 = B minimums

vmovdqa ymm6, ymmword ptr [InitialPmaxVal] ;ymm6 = R maximums
vmovdqa ymm7, ymm6         ;ymm7 = G maximums
vmovdqa ymm8, ymm6         ;ymm8 = B maximums

xor rcx, rcx                ;rcx = common array offset

; Scan RGB color plane arrays for packed minimums and maximums
align 16
@@: vmovdqa ymm0, ymmword ptr [r10+rcx] ;ymm0 = R pixels
vmovdqa ymm1, ymmword ptr [r11+rcx] ;ymm1 = G pixels
vmovdqa ymm2, ymmword ptr [r12+rcx] ;ymm2 = B pixels

vpminub ymm3, ymm3, ymm0    ;update R minimums
vpminub ymm4, ymm4, ymm1    ;update G minimums
vpminub ymm5, ymm5, ymm2    ;update B minimums

vpmaxub ymm6, ymm6, ymm0    ;update R maximums
vpmaxub ymm7, ymm7, ymm1    ;update G maximums
vpmaxub ymm8, ymm8, ymm2    ;update B maximums

add rcx, 32
sub rdx, 32
jnz @B

; Calculate the final RGB minimum values
_YmmVpextrMinub rax, ymm3, ymm0
mov byte ptr [r8], al       ;save min R
_YmmVpextrMinub rax, ymm4, ymm0
mov byte ptr [r8+1], al     ;save min G
_YmmVpextrMinub rax, ymm5, ymm0
mov byte ptr [r8+2], al     ;save min B

; Calculate the final RGB maximum values
_YmmVpextrMaxub rax, ymm6, ymm1
mov byte ptr [r9], al       ;save max R
_YmmVpextrMaxub rax, ymm7, ymm1
mov byte ptr [r9+1], al     ;save max G
_YmmVpextrMaxub rax, ymm8, ymm1
mov byte ptr [r9+2], al     ;save max B

mov eax, 1                  ;set success return code

```

```

Done: vzeroupper
      _RestoreXmmRegs xmm6, xmm7, xmm8
      _DeleteFrame r12
      ret
Avx2CalcRgbMinMax_ endp
end

;-----
;##### Ch10_06.asm
include <MacrosX86-64-AVX.asmh>

GsMask      .const
             dword 0fffffffh, 0, 0, 0, 0fffffffh, 0, 0, 0
r4_0p5      real4 0.5
r4_255p0     real4 255.0

             extern c_NumPixelsMin:dword
             extern c_NumPixelsMax:dword

;extern "C" bool Avx2ConvertRgbToGs_(uint8_t* pb_gs, const RGB32* pb_rgb, int num_pixels, const float coef[4])
; Note: Memory pointed to by pb_rgb is ordered as follows:
;       R(0,0), G(0,0), B(0,0), A(0,0), R(0,1), G(0,1), B(0,1), A(0,1), ...

.code
Avx2ConvertRgbToGs_ proc frame
    _CreateFrame RGBGS_, 0, 112
    _SaveXmmRegs xmm6, xmm7, xmm11, xmm12, xmm13, xmm14, xmm15
    _EndProlog

; Validate argument values
xor eax, eax                ;set error return code
cmp r8d, [c_NumPixelsMin]   ;jump if num_pixels < min value
jl Done
cmp r8d, [c_NumPixelsMax]   ;jump if num_pixels > max value
jg Done
test r8d, 7
jnz Done                    ;jump if (num_pixels % 8) != 0

test rcx, 1fh
jnz Done                    ;jump if pb_gs is not aligned
test rdx, 1fh
jnz Done                    ;jump if pb_rgb is not aligned

; Perform required initializations
vbroadcastss ymm11, real4 ptr [r4_255p0] ;ymm11 = packed 255.0
vbroadcastss ymm12, real4 ptr [r4_0p5]   ;ymm12 = packed 0.5
vpxor ymm13, ymm13, ymm13               ;ymm13 = packed zero

vmovups xmm0, xmmword ptr [r9]
vperm2f128 ymm14, ymm0, ymm0, 00000000b ;ymm14 = packed coef

vmovups ymm15, ymmword ptr [GsMask]      ;ymm15 = GsMask (SPFP)
; Load next 8 RGB32 pixel values (P0 - P7)
align 16
@@: vmovdqa ymm0, ymmword ptr [rdx] ;ymm0 = 8 rgb32 pixels (P7 - P0)
; Size-promote RGB32 color components from bytes to dwords
vpunpcklbw ymm1, ymm0, ymm13
vpunpckhbw ymm2, ymm0, ymm13
vpunpcklwd ymm3, ymm1, ymm13 ;ymm3 = P1, P0 (dword)
vpunpckhwd ymm4, ymm1, ymm13 ;ymm4 = P3, P2 (dword)
vpunpcklwd ymm5, ymm2, ymm13 ;ymm5 = P5, P4 (dword)
vpunpckhwd ymm6, ymm2, ymm13 ;ymm6 = P7, P6 (dword)
; Convert color component values to single-precision floating-point
vcvtdq2ps ymm0, ymm3 ;ymm0 = P1, P0 (SPFP)
vcvtdq2ps ymm1, ymm4 ;ymm1 = P3, P2 (SPFP)
vcvtdq2ps ymm2, ymm5 ;ymm2 = P5, P4 (SPFP)
vcvtdq2ps ymm3, ymm6 ;ymm3 = P7, P6 (SPFP)
; Multiply color component values by color conversion coefficients

```

```

vmulps ymm0, ymm0, ymm14
vmulps ymm1, ymm1, ymm14
vmulps ymm2, ymm2, ymm14
vmulps ymm3, ymm3, ymm14

; Sum weighted color components for final grayscale values
vhaddps ymm4, ymm0, ymm0
vhaddps ymm4, ymm4, ymm4           ;ymm4[159:128] = P1, ymm4[31:0] = P0
vhaddps ymm5, ymm1, ymm1
vhaddps ymm5, ymm5, ymm5           ;ymm5[159:128] = P3, ymm4[31:0] = P2
vhaddps ymm6, ymm2, ymm2
vhaddps ymm6, ymm6, ymm6           ;ymm6[159:128] = P5, ymm4[31:0] = P4
vhaddps ymm7, ymm3, ymm3
vhaddps ymm7, ymm7, ymm7           ;ymm7[159:128] = P7, ymm4[31:0] = P6

; Merge SPFP grayscale values into a single YMM register
vandps ymm4, ymm4, ymm15           ;mask out unneeded SPFP values
vandps ymm5, ymm5, ymm15
vandps ymm6, ymm6, ymm15
vandps ymm7, ymm7, ymm15
vpslldq ymm5, ymm5, 4
vpslldq ymm6, ymm6, 8
vpslldq ymm7, ymm7, 12
vorps ymm0, ymm4, ymm5
vorps ymm1, ymm6, ymm7
vorps ymm2, ymm0, ymm1           ;ymm2 = 8 GS pixel values (SPFP)

; Add 0.5 rounding factor and clip to 0.0 - 255.0
vaddps ymm2, ymm2, ymm12           ;add 0.5f rounding factor
vminps ymm3, ymm2, ymm11           ;clip pixels above 255.0
vmaxps ymm4, ymm3, ymm13           ;clip pixels below 0.0

; Convert SPFP values to bytes and save
vcvtps2dq ymm3, ymm2           ;convert GS SPFP to dwords
vpckusdw ymm4, ymm3, ymm13       ;convert GS dwords to words
vpckuswb ymm5, ymm4, ymm13       ;convert GS words to bytes

vperm2i128 ymm6, ymm13, ymm5, 3   ;xmm5 = GS P3:P0, xmm6 = GS P7:P4

vmovd dword ptr [rcx], xmm5       ;save P3 - P0
vmovd dword ptr [rcx+4], xmm6     ;save P7 - P4

add rdx, 32                       ;update pb_rgb to next block
add rcx, 8                         ;update pb_gs to next block
sub r8d, 8                         ;num_pixels -= 8
jnz @B                             ;repeat until done

mov eax, 1                         ;set success return code

Done: vzeroupper
      _RestoreXmmRegs xmm6, xmm7, xmm11, xmm12, xmm13, xmm14, xmm15
      _DeleteFrame
      ret
Avx2ConvertRgbToGs_ endp
end

;-----
;##### Ch11_01_.asm
include <MacrosX86-64-AVX.asmh>
extern c_NumPtsMin:dword
extern c_NumPtsMax:dword
extern c_KernelSizeMin:dword
extern c_KernelSizeMax:dword

; extern "C" bool Convolve1_(float* y, const float* x, int num_pts, const float* kernel, int kernel_size)
.code
Convolve1_ proc frame
      _CreateFrame CV_, 0, 0, rbx, rsi
      _EndProlog

```

```

; Verify argument values
xor eax, eax                      ;set error code (rax is also loop index var)
mov r10d, dword ptr [rbp+CV_OffsetStackArgs]
test r10d, 1
jz Done                           ;jump if kernel_size is even
cmp r10d, [c_KernelSizeMin]
jl Done                           ;jump if kernel_size too small
cmp r10d, [c_KernelSizeMax]
jg Done                           ;jump if kernel_size too big

cmp r8d, [c_NumPtsMin]
jl Done                           ;jump if num_pts too small
cmp r8d, [c_NumPtsMax]
jg Done                           ;jump if num_pts too big

; Perform required initializations
mov r8d, r8d                      ;r8 = num_pts
shr r10d, 1                       ;ks2 = ks / 2
lea rdx, [rdx+r10*4]              ;rdx = x + ks2 (first data point)

; Perform convolution
LP1: vxorps xmm5, xmm5, xmm5       ;sum = 0.0;
mov r11, r10
neg r11                           ;k = -ks2

LP2: mov rbx, rax
sub rbx, r11                      ;rbx = i - k
vmovss xmm0, real4 ptr [rdx+rbx*4] ;xmm0 = x[i - k]
mov rsi, r11
add rsi, r10                      ;rsi = k + ks2
vfmadd231ss xmm5, xmm0, [r9+rsi*4] ;sum += x[i - k] * kernel[k + ks2]

add r11, 1                        ;k++
cmp r11, r10
jle LP2                           ;jump if k <= ks2

vmovss real4 ptr [rcx+rax*4], xmm5 ;y[i] = sum

add rax, 1                        ;i += 1
cmp rax, r8
jl LP1                            ;jump if i < num_pts

mov eax, 1                        ;set success return code

Done: vzeroupper
      _DeleteFrame rbx, rsi
      ret
Convolve1_ endp

; extern "C" bool Convolve1Ks5_(float* y, const float* x, int num_pts, const float* kernel, int kernel_size)
Convolve1Ks5_ proc
; Verify argument values
xor eax, eax                      ;set error code (rax is also loop index var)
cmp dword ptr [rsp+40], 5
jne Done                          ;jump if kernel_size is not 5

cmp r8d, [c_NumPtsMin]
jl Done                          ;jump if num_pts too small
cmp r8d, [c_NumPtsMax]
jg Done                          ;jump if num_pts too big

; Perform required initializations
mov r8d, r8d                      ;r8 = num_pts
add rdx, 8                        ;x += 2

; Perform convolution
@@: vxorps xmm4, xmm4, xmm4       ;initialize sum vars

```

```

vxorps xmm5, xmm5, xmm5
mov r11, rax
add r11, 2                ;j = i + ks2

vmovss xmm0, real4 ptr [rdx+r11*4] ;xmm0 = x[j]
vmadd231ss xmm4, xmm0, [r9]        ;xmm4 += x[j] * kernel[0]

vmovss xmm1, real4 ptr [rdx+r11*4-4] ;xmm1 = x[j - 1]
vmadd231ss xmm5, xmm1, [r9+4]        ;xmm5 += x[j - 1] * kernel[1]

vmovss xmm0, real4 ptr [rdx+r11*4-8] ;xmm0 = x[j - 2]
vmadd231ss xmm4, xmm0, [r9+8]        ;xmm4 += x[j - 2] * kernel[2]

vmovss xmm1, real4 ptr [rdx+r11*4-12] ;xmm1 = x[j - 3]
vmadd231ss xmm5, xmm1, [r9+12]       ;xmm5 += x[j - 3] * kernel[3]

vmovss xmm0, real4 ptr [rdx+r11*4-16] ;xmm0 = x[j - 4]
vmadd231ss xmm4, xmm0, [r9+16]       ;xmm4 += x[j - 4] * kernel[4]

vaddps xmm4, xmm4, xmm5
vmovss real4 ptr [rcx+rax*4], xmm4    ;save y[i]

inc rax                          ;i += 1
cmp rax, r8                      ;jump if i < num_pts
jl @B

mov eax, 1                       ;set success return code

Done:  vzeroupper
ret
Convolve1Ks5_endp
end
;-----
;##### Ch11_02.asm
include <MacrosX86-64-AVX.asmh>
extern c_NumPtsMin:dword
extern c_NumPtsMax:dword
extern c_KernelSizeMin:dword
extern c_KernelSizeMax:dword

; extern bool Convolve2_(float* y, const float* x, int num_pts, const float* kernel, int kernel_size)
.code
Convolve2_ proc frame
    _CreateFrame CV2_, 0, 0, rbx
    _EndProlog

; Validate argument values
xor eax, eax                    ;set error code

mov r10d, dword ptr [rbp+CV2_OffsetStackArgs]
test r10d, 1
jz Done                        ;kernel_size is even
cmp r10d, [c_KernelSizeMin]
jl Done                        ;kernel_size too small
cmp r10d, [c_KernelSizeMax]
jg Done                        ;kernel_size too big

cmp r8d, [c_NumPtsMin]
jl Done                        ;num_pts too small
cmp r8d, [c_NumPtsMax]
jg Done                        ;num_pts too big
test r8d, 7
jnz Done                      ;num_pts not even multiple of 8

test rcx, 1fh
jnz Done                      ;y is not properly aligned

```

```

; Initialize convolution loop variables
shr r10d, 1                    ;r10 = kernel_size / 2 (ks2)
lea rdx, [rdx+r10*4]           ;rdx = x + ks2 (first data point)
xor ebx, ebx                   ;i = 0

; Perform convolution
LP1:  vxorps ymm0, ymm0, ymm0    ;packed sum = 0.0;
      mov r11, r10              ;r11 = ks2
      neg r11                   ;k = -ks2

LP2:  mov rax, rbx               ;rax = i
      sub rax, r11              ;rax = i - k
      vmovups ymm1, ymmword ptr [rdx+rax*4] ;load x[i - k]:x[i - k + 7]

      mov rax, r11
      add rax, r10              ;rax = k + ks2
      vbroadcastss ymm2, real4 ptr [r9+rax*4] ;ymm2 = kernel[k + ks2]
      vmadd231ps ymm0, ymm1, ymm2 ;ymm0 += x[i-k]:x[i-k+7] * kernel[k+ks2]

      add r11, 1                ;k += 1
      cmp r11, r10
      jle LP2                  ;repeat until k > ks2

      vmovaps ymmword ptr [rcx+rbx*4], ymm0 ;save y[i]:y[i + 7]

      add rbx, 8                ;i += 8
      cmp rbx, r8
      jl LP1                   ;repeat until done
      mov eax, 1                ;set success return code

Done:  vzeroupper
      _DeleteFrame rbx
      ret
Convolve2_ endp

; extern bool Convolve2Ks5_(float* y, const float* x, int num_pts, const float* kernel, int kernel_size)
Convolve2Ks5_ proc frame
    _CreateFrame CKS5_, 0, 48
    _SaveXmmRegs xmm6, xmm7, xmm8
    _EndProlog

; Validate argument values
xor eax, eax                    ;set error code (rax is also loop index var)
cmp dword ptr [rbp+CKS5_OffsetStackArgs], 5
jne Done                        ;jump if kernel_size is not 5

cmp r8d, [c_NumPtsMin]
jl Done                        ;jump if num_pts too small
cmp r8d, [c_NumPtsMax]
jg Done                        ;jump if num_pts too big
test r8d, 7
jnz Done                      ;num_pts not even multiple of 8

test rcx, 1fh
jnz Done                      ;y is not properly aligned

; Perform required initializations
vbroadcastss ymm4, real4 ptr [r9] ;kernel[0]
vbroadcastss ymm5, real4 ptr [r9+4] ;kernel[1]
vbroadcastss ymm6, real4 ptr [r9+8] ;kernel[2]
vbroadcastss ymm7, real4 ptr [r9+12] ;kernel[3]
vbroadcastss ymm8, real4 ptr [r9+16] ;kernel[4]
mov r8d, r8d                    ;r8 = num_pts
add rdx, 8                      ;x += 2

; Perform convolution
@@:  vxorps ymm2, ymm2, ymm2    ;initialize sum vars

```



```

vxorps ymm3, ymm3, ymm3
mov r11, rax
add r11, 2                ;j = i + ks2

vmovups ymm0, ymmword ptr [rdx+r11*4] ;ymm0 = x[j]:x[j + 7]
vmadd231ps ymm2, ymm0, ymm4        ;ymm2 += x[j]:x[j + 7] * kernel[0]

vmovups ymm1, ymmword ptr [rdx+r11*4-4] ;ymm1 = x[j - 1]:x[j + 6]
vmadd231ps ymm3, ymm1, ymm5        ;ymm3 += x[j - 1]:x[j + 6] * kernel[1]

vmovups ymm0, ymmword ptr [rdx+r11*4-8] ;ymm0 = x[j - 2]:x[j + 5]
vmadd231ps ymm2, ymm0, ymm6        ;ymm2 += x[j - 2]:x[j + 5] * kernel[2]

vmovups ymm1, ymmword ptr [rdx+r11*4-12] ;ymm1 = x[j - 3]:x[j + 4]
vmadd231ps ymm3, ymm1, ymm7        ;ymm3 += x[j - 3]:x[j + 4] * kernel[3]

vmovups ymm0, ymmword ptr [rdx+r11*4-16] ;ymm0 = x[j - 4]:x[j + 3]
vmadd231ps ymm2, ymm0, ymm8        ;ymm2 += x[j - 4]:x[j + 3] * kernel[4]

vaddps ymm0, ymm2, ymm3            ;final values
vmovaps ymmword ptr [rcx+rax*4], ymm0 ;save y[i]:y[i + 7]

add rax, 8                ;i += 8
cmp rax, r8
jl @B                     ;jump if i < num_pts
mov eax, 1                ;set success return code

Done:  vzeroupper
       _RestoreXmmRegs xmm6, xmm7, xmm8
       _DeleteFrame
       ret
Convolve2Ks5_endp
end

:-----
:##### Ch11_02_Test_.asm
include <MacrosX86-64-AVX.asmh>
extern c_NumPtsMin:dword
extern c_NumPtsMax:dword

; extern bool Convolve2Ks5Test_(float* y, const float* x, int num_pts, const float* kernel, int
kernel_size)
.code
Convolve2Ks5Test_ proc frame
    _CreateFrame CKS5T_, 0, 48
    _SaveXmmRegs xmm6, xmm7, xmm8
    _EndProlog

; Validate argument values
xor eax, eax                ;set error code (rax is also loop index var)
cmp dword ptr [rbp+CKS5T_OffsetStackArgs], 5
jne Done                   ;jump if kernel_size is not 5

cmp r8d, [c_NumPtsMin]
jl Done                   ;jump if num_pts too small
cmp r8d, [c_NumPtsMax]
jg Done                   ;jump if num_pts too big
test r8d, 7
jnz Done                  ;num_pts not even multiple of 8

test rcx, 1fh
jnz Done                  ;y is not properly aligned

; Perform required initializations
vbroadcastss ymm4, real4 ptr [r9] ;kernel[0]
vbroadcastss ymm5, real4 ptr [r9+4] ;kernel[1]
vbroadcastss ymm6, real4 ptr [r9+8] ;kernel[2]
vbroadcastss ymm7, real4 ptr [r9+12] ;kernel[3]

```

```

vbroadcastss ymm8, real4 ptr [r9+16] ;kernel[4]
mov r8d, r8d                ;r8 = num_pts
add rdx, 8                  ;x += 2

; Perform convolution
@@:  vxorps ymm2, ymm2, ymm2        ;initialize sum vars
      vxorps ymm3, ymm3, ymm3
      mov r11, rax
      add r11, 2                ;j = i + ks2

vmovups ymm0, ymmword ptr [rdx+r11*4] ;ymm0 = x[j]:x[j + 7]
vmulps ymm0, ymm0, ymm4        ;ymm2 += x[j]:x[j + 7] * kernel[0]
vaddps ymm2, ymm2, ymm0

vmovups ymm1, ymmword ptr [rdx+r11*4-4] ;ymm1 = x[j - 1]:x[j + 6]
vmulps ymm1, ymm1, ymm5        ;ymm3 += x[j - 1]:x[j + 6] * kernel[1]
vaddps ymm3, ymm3, ymm1

vmovups ymm0, ymmword ptr [rdx+r11*4-8] ;ymm0 = x[j - 2]:x[j + 5]
vmulps ymm0, ymm0, ymm6        ;ymm2 += x[j - 2]:x[j + 5] * kernel[2]
vaddps ymm2, ymm2, ymm0

vmovups ymm1, ymmword ptr [rdx+r11*4-12] ;ymm1 = x[j - 3]:x[j + 4]
vmulps ymm1, ymm1, ymm7        ;ymm3 += x[j - 3]:x[j + 4] * kernel[3]
vaddps ymm3, ymm3, ymm1

vmovups ymm0, ymmword ptr [rdx+r11*4-16] ;ymm0 = x[j - 4]:x[j + 3]
vmulps ymm0, ymm0, ymm8        ;ymm2 += x[j - 4]:x[j + 3] * kernel[4]
vaddps ymm2, ymm2, ymm0

vaddps ymm0, ymm2, ymm3        ;final values
vmovaps ymmword ptr [rcx+rax*4], ymm0 ;save y[i]:y[i + 7]

add rax, 8                ;i += 8
cmp rax, r8
jl @B                     ;jump if i < num_pts
mov eax, 1                ;set success return code

Done:  vzeroupper
       _RestoreXmmRegs xmm6, xmm7, xmm8
       _DeleteFrame
       ret
Convolve2Ks5Test_ endp
end

:-----
:##### Ch11_03_.asm
:
; extern "C" uint64_t GprMulx_(uint32_t a, uint32_t b, uint64_t flags[2]);
; Requires      BMI2

.code
GprMulx_ proc

; Save copy of status flags before mulx
pushfq
pop rax
mov qword ptr [r8], rax        ;save original status flags

; Perform flagless multiplication. The mulx instruction below computes
; the product of explicit source operand ecx (a) and implicit source
; operand edx (b). The 64-bit result is saved to the register pair r11d:r10d.
mulx r11d, r10d, ecx          ;r11d:r10d = a * b

; Save copy of status flags after mulx
pushfq
pop rax
mov qword ptr [r8+8], rax      ;save post mulx status flags

```

```

; Move 64-bit result to rax
mov eax,r10d
shl r11,32
or rax,r11
ret
GprMulx_ endp

; extern "C" void GprShiftx_(uint32_t x, uint32_t count, uint32_t results[3], uint64_t flags[4])
; Requires BMI2

GprShiftx_ proc

; Save copy of status flags before shifts
pushfq
pop rax
mov qword ptr [r9],rax ;save original status flags

; Load argument values and perform shifts. Note that each shift
; instruction requires three operands: DesOp, SrcOp, and CountOp.

sarl eax,ecx,edx ;shift arithmetic right
mov dword ptr [r8],eax
pushfq
pop rax
mov qword ptr [r9+8],rax

shll eax,ecx,edx ;shift logical left
mov dword ptr [r8+4],eax
pushfq
pop rax
mov qword ptr [r9+16],rax

shrr eax,ecx,edx ;shift logical right
mov dword ptr [r8+8],eax
pushfq
pop rax
mov qword ptr [r9+24],rax

ret
GprShiftx_ endp

;-----
;##### Ch11_04_.asm
;-----
; extern "C" void GprCountZeroBits_(uint32_t x, uint32_t* lzcnt, uint32_t* tzcnt);
; Requires BMI1, LZCNT

.code
GprCountZeroBits_ proc
lzcnt eax,ecx ;count leading zeros
mov dword ptr [rdx],eax ;save result

tzcnt eax,ecx ;count trailing zeros
mov dword ptr [r8],eax ;save result
ret
GprCountZeroBits_ endp

; extern "C" uint32_t GprBextr_(uint32_t x, uint8_t start, uint8_t length);
; Requires BMI1

GprBextr_ proc
mov al,r8b
mov ah,al ;ah = length
mov al,dl ;al = start
bextr eax,ecx,eax ;eax = extracted bit field (from x)
ret

```

```

GprBextr_ endp

; extern "C" uint32_t GprAndNot_(uint32_t x, uint32_t y);
; Requires BMI1

GprAndNot_ proc
andn eax,ecx,edx ;eax = ~x & y
ret
GprAndNot_ endp

;-----
;##### Ch11_05_.asm
;-----
; extern "C" void SingleToHalfPrecision_(uint16_t x_hp[8], float x_sp[8], int rc);

.code
SingleToHalfPrecision_ proc

; Convert packed single-precision to packed half-precision
vmovups ymm0,ymmword ptr [rdx] ;ymm0 = 8 SPFP values

cmp r8d,0
jne @F
vcvtps2ph xmm1,ymm0,0 ;round to nearest
jmp SaveResult

@@: cmp r8d,1
jne @F
vcvtps2ph xmm1,ymm0,1 ;round down
jmp SaveResult

@@: cmp r8d,2
jne @F
vcvtps2ph xmm1,ymm0,2 ;round up
jmp SaveResult

@@: cmp r8d,3
jne @F
vcvtps2ph xmm1,ymm0,3 ;truncate
jmp SaveResult

@@: vcvtps2ph xmm1,ymm0,4 ;use MXCSR.RC

SaveResult:
vmovdqu xmmword ptr [rcx],xmm1 ;save 8 HPFP values
vzeroupper
ret

SingleToHalfPrecision_ endp

; extern "C" void HalfToSinglePrecision_(float x_sp[8], uint16_t x_hp[8]);

HalfToSinglePrecision_ proc

; Convert packed half-precision to packed single-precision
vcvtph2ps ymm0,xmmword ptr [rdx]
vmovups ymmword ptr [rcx],ymm0 ;save 8 SPFP values

vzeroupper
ret

HalfToSinglePrecision_ endp

;-----
;##### Ch13_01.asm
;-----
include <cmpequ.asmh>

```

```

        .const
r8_three    real8 3.0
r8_four     real8 4.0

        extern g_PI:real8

; extern "C" bool Avx512CalcSphereAreaVol_(double* sa, double* v, double r, double error_val);
; Returns: false = invalid radius, true = valid radius

        .code
Avx512CalcSphereAreaVol_ proc

; Test radius for value >= 0.0
vmovsd xmm0, xmm0, xmm2    ;xmm0 = radius
vxorpd xmm5, xmm5, xmm5    ;xmm5 = 0.0
vmovsd xmm16, xmm16, xmm3   ;xmm16 = error_val
vcmpsd k1, xmm0, xmm5, CMP_GE ;k1[0] = 1 if radius >= 0.0

; Calculate surface area and volume using mask from compare
vmulsd xmm1[k1], xmm0, xmm0 ;xmm1 = r * r
vmulsd xmm2[k1], xmm1, [r8_four] ;xmm2 = 4 * r * r
vmulsd xmm3[k1], xmm2, [g_PI] ;xmm3 = 4 * PI * r * r (sa)
vmulsd xmm4[k1], xmm3, xmm0 ;xmm4 = 4 * PI * r * r * r
vdivsd xmm5[k1], xmm4, [r8_three] ;xmm5 = 4 * PI * r * r * r / 3 (vol)

; Set surface area and volume to error_val if radius < 0.0 is true
knotw k2, k1 ;k2[0] = 1 if radius < 0.0
vmovsd xmm3[k2], xmm3, xmm16 ;xmm3 = error_val if radius < 0.0
vmovsd xmm5[k2], xmm5, xmm16 ;xmm5 = error_val if radius < 0.0

; Save results
vmovsd real8 ptr [rcx], xmm3 ;save surface area
vmovsd real8 ptr [rdx], xmm5 ;save volume

        kmovw eax, k1 ;eax = return code
        ret
Avx512CalcSphereAreaVol_ endp
;-----
;##### Ch13_02.asm
include <cmpequ.asmh>

; extern "C" bool Avx512CalcValues_(double* c, const double* a, const double* b, size_t n);

        .code
Avx512CalcValues_ proc

; Validate n and initialize array index i
xor eax, eax ;set error return code (also i = 0)
test r9, r9 ;is n == 0?
jz Done ;jump if n is zero

        vxorpd xmm5, xmm5, xmm5 ;xmm5 = 0.0

; Load next a[i] and b[i], calculate val
@@: vmovsd xmm0, real8 ptr [rdx+rax*8] ;xmm0 = a[i];
vmovsd xmm1, real8 ptr [r8+rax*8] ;xmm1 = b[i];
vmulsd xmm2, xmm0, xmm1 ;val = a[i] * b[i]

; Calculate c[i] = (val >= 0.0) ? sqrt(val) : val * val
vcmpsd k1, xmm2, xmm5, CMP_GE ;k1[0] = 1 if val >= 0.0
vsqrtsd xmm3[k1]{z}, xmm3, xmm2 ;xmm3 = (val > 0.0) ? sqrt(val) : 0.0
knotw k2, k1 ;k2[0] = 1 if val < 0.0
vmulsd xmm4[k2]{z}, xmm2, xmm2 ;xmm4 = (val < 0.0) ? val * val : 0.0
voprd xmm0, xmm4, xmm3 ;xmm0 = (val >= 0.0) ? sqrt(val) : val * val
vmovsd real8 ptr [rcx+rax*8], xmm0 ;save result to c[i]

; Update index i and repeat until done

```

```

        inc rax ;i += 1
        cmp rax, r9
        jl @B
        mov eax, 1 ;set success return code

Done: ret
Avx512CalcValues_ endp
;-----
;##### Ch13_03.asm
;-----
; extern "C" void Avx512CvtF32ToU32_(uint32_t val_cvt[4], float val);

        .code
Avx512CvtF32ToU32_ proc
        vcvts2usi eax, xmm1{rn-sae} ;Convert using round to nearest
        mov dword ptr [rcx], eax

        vcvts2usi eax, xmm1{rd-sae} ;Convert using round down
        mov dword ptr [rcx+4], eax

        vcvts2usi eax, xmm1{ru-sae} ;Convert using round up
        mov dword ptr [rcx+8], eax

        vcvts2usi eax, xmm1{rz-sae} ;Convert using round to zero (truncate)
        mov dword ptr [rcx+12], eax
        ret
Avx512CvtF32ToU32_ endp

; extern "C" void Avx512CvtF64ToU64_(uint64_t val_cvt[4], double val);

Avx512CvtF64ToU64_ proc
        vcvtsd2usi rax, xmm1{rn-sae}
        mov qword ptr [rcx], rax

        vcvtsd2usi rax, xmm1{rd-sae}
        mov qword ptr [rcx+8], rax

        vcvtsd2usi rax, xmm1{ru-sae}
        mov qword ptr [rcx+16], rax

        vcvtsd2usi rax, xmm1{rz-sae}
        mov qword ptr [rcx+24], rax
        ret
Avx512CvtF64ToU64_ endp

; extern "C" void Avx512CvtF64ToF32_(float val_cvt[4], double val);

Avx512CvtF64ToF32_ proc
        vcvtsd2ss xmm2, xmm1{rn-sae}
        vmovss real4 ptr [rcx], xmm2

        vcvtsd2ss xmm2, xmm2, xmm1{rd-sae}
        vmovss real4 ptr [rcx+4], xmm2

        vcvtsd2ss xmm2, xmm2, xmm1{ru-sae}
        vmovss real4 ptr [rcx+8], xmm2

        vcvtsd2ss xmm2, xmm2, xmm1{rz-sae}
        vmovss real4 ptr [rcx+12], xmm2
        ret
Avx512CvtF64ToF32_ endp
;-----
;##### Ch13_04.asm
;-----
; Mask values used to calculate floating-point absolute values

```

```

ConstVals    segment readonly align(64) 'const'
AbsMaskF32   dword 16 dup(7fffffffh)
AbsMaskF64   qword 8 dup(7fffffffffffffffh)
ConstVals    ends
; extern "C" void Avx512PackedMathF32_(const ZmmVal* a, const ZmmVal* b, ZmmVal c[8]);
; code
Avx512PackedMathF32_ proc

; Load packed SP floating-point values
vmovaps zmm0, zmmword ptr [rcx] ; zmm0 = *a
vmovaps zmm1, zmmword ptr [rdx] ; zmm1 = *b

; Packed SP floating-point addition
vaddps zmm2, zmm0, zmm1
vmovaps zmmword ptr [r8+0], zmm2

; Packed SP floating-point subtraction
vsubps zmm2, zmm0, zmm1
vmovaps zmmword ptr [r8+64], zmm2

; Packed SP floating-point multiplication
vmulps zmm2, zmm0, zmm1
vmovaps zmmword ptr [r8+128], zmm2

; Packed SP floating-point division
vdivps zmm2, zmm0, zmm1
vmovaps zmmword ptr [r8+192], zmm2

; Packed SP floating-point absolute value (b)
vandps zmm2, zmm1, zmmword ptr [AbsMaskF32]
vmovaps zmmword ptr [r8+256], zmm2

; Packed SP floating-point square root (a)
vsqrtps zmm2, zmm0
vmovaps zmmword ptr [r8+320], zmm2

; Packed SP floating-point minimum
vminps zmm2, zmm0, zmm1
vmovaps zmmword ptr [r8+384], zmm2

; Packed SP floating-point maximum
vmaxps zmm2, zmm0, zmm1
vmovaps zmmword ptr [r8+448], zmm2

vzeroupper
ret
Avx512PackedMathF32_ endp

; extern "C" void Avx512PackedMathF64_(const ZmmVal* a, const ZmmVal* b, ZmmVal c[8]);
; code
Avx512PackedMathF64_ proc

; Load packed DP floating-point values
vmovapd zmm0, zmmword ptr [rcx] ; zmm0 = *a
vmovapd zmm1, zmmword ptr [rdx] ; zmm1 = *b

; Packed DP floating-point addition
vaddpd zmm2, zmm0, zmm1
vmovapd zmmword ptr [r8+0], zmm2

; Packed DP floating-point subtraction
vsubpd zmm2, zmm0, zmm1
vmovapd zmmword ptr [r8+64], zmm2

; Packed DP floating-point multiplication
vmulpd zmm2, zmm0, zmm1
vmovapd zmmword ptr [r8+128], zmm2

```

```

; Packed DP floating-point division
vdivpd zmm2, zmm0, zmm1
vmovapd zmmword ptr [r8+192], zmm2

; Packed DP floating-point absolute value (b)
vandpd zmm2, zmm1, zmmword ptr [AbsMaskF64]
vmovapd zmmword ptr [r8+256], zmm2

; Packed DP floating-point square root (a)
vsqrtpd zmm2, zmm0
vmovapd zmmword ptr [r8+320], zmm2

; Packed DP floating-point minimum
vminpd zmm2, zmm0, zmm1
vmovapd zmmword ptr [r8+384], zmm2

; Packed DP floating-point maximum
vmaxpd zmm2, zmm0, zmm1
vmovapd zmmword ptr [r8+448], zmm2
vzeroupper
ret
Avx512PackedMathF64_ endp
end;

```

---

```

; ##### Ch13_05.asm
; include <cmpequ.asmh>

```

```

; extern "C" void Avx512PackedCompareF32_(const ZmmVal* a, const ZmmVal* b, ZmmVal c[8]);
; code
Avx512PackedCompareF32_ proc
vmovaps zmm0, [rcx] ; zmm0 = a
vmovaps zmm1, [rdx] ; zmm1 = b

; Perform packed EQUAL compare
vcmpsps k1, zmm0, zmm1, CMP_EQ
kmovw word ptr [r8], k1

; Perform packed NOT EQUAL compare
vcmpsps k1, zmm0, zmm1, CMP_NEQ
kmovw word ptr [r8+2], k1

; Perform packed LESS THAN compare
vcmpsps k1, zmm0, zmm1, CMP_LT
kmovw word ptr [r8+4], k1

; Perform packed LESS THAN OR EQUAL compare
vcmpsps k1, zmm0, zmm1, CMP_LE
kmovw word ptr [r8+6], k1

; Perform packed GREATER THAN compare
vcmpsps k1, zmm0, zmm1, CMP_GT
kmovw word ptr [r8+8], k1

; Perform packed GREATER THAN OR EQUAL compare
vcmpsps k1, zmm0, zmm1, CMP_GE
kmovw word ptr [r8+10], k1

; Perform packed ORDERED compare
vcmpsps k1, zmm0, zmm1, CMP_ORD
kmovw word ptr [r8+12], k1

; Perform packed UNORDERED compare
vcmpsps k1, zmm0, zmm1, CMP_UNORD
kmovw word ptr [r8+14], k1

```

```

        vzeroupper
        ret
Avx512PackedCompareF32_ endp
end
;-----
##### Ch13_06.asm
include <cmpequ.asmh> include <MacrosX86-64-AVX.asmh>

extern c_NumRowsMax:qword
extern c_NumColsMax:qword

; extern "C" bool Avx512CalcColumnMeans(const double* x, size_t nrows, size_t ncols, double*
col_means, size_t* col_counts, double x_min);

.code
Avx512CalcColumnMeans_proc frame
    _CreateFrame CCM_, 0, 0, rbx, r12, r13
    _EndProlog

; Validate nrows and ncols
xor eax, eax                ;set error return code
test rdx, rdx
jz Done                    ;jump if nrows is zero
cmp rdx, [c_NumRowsMax]
ja Done                    ;jump if nrows is too large
test r8, r8
jz Done                    ;jump if ncols is zero
cmp r8, [c_NumColsMax]
ja Done                    ;jump if ncols is too large

; Load argument values col_counts and x_min
mov ebx, 1
vpbroadcastq zmm4, rbx      ;zmm4 = 8 qwords of 1
mov rbx, [rbp+CCM_OffsetStackArgs] ;rbx = col_counts ptr
lea r13, [rbp+CCM_OffsetStackArgs+8] ;r13 = ptr to x_min

; Set initial col_means and col_counts to zero
xor r10, r10
vxorpd xmm0, xmm0, xmm0
@@: vmovsd real8 ptr[r9+rax*8], xmm0 ;col_means[i] = 0.0
    mov [rbx+rax*8], r10             ;col_counts[i] = 0
    inc rax
    cmp rax, r8
    jne @B                          ;repeat until done

; Compute the sum of each column in x
LP1: xor r10, r10                ;r10 = col_index
     mov r11, r9                ;r11 = ptr to col_means
     mov r12, rbx               ;r12 = ptr to col_counts

LP2:  mov rax, r10               ;rax = col_index
     add rax, 8
     cmp rax, r8                ;8 or more columns remaining?
     ja @F                     ;jump if col_index + 8 > ncols

; Update col_means and col_counts using next eight columns
vmovupd zmm0, zmmword ptr [rcx] ;load next 8 cols of cur row
vcmpdpd k1, zmm0, real8 bcst [r13], CMP_GE ;k1 = mask of values >= x_min
vmovupd zmm1{k1}{z}, zmm0       ;values >= x_min or 0.0
vaddpd zmm2, zmm1, zmmword ptr [r11] ;add values to col_means
vmovupd zmmword ptr [r11], zmm2 ;save updated col_means

vpmovm2q zmm0, k1               ;convert mask to vector
vpandq zmm1, zmm0, zmm4         ;qword values for add
vpaddq zmm2, zmm1, zmmword ptr [r12] ;update col_counts
vmovdqu64 zmmword ptr [r12], zmm2 ;save updated col_counts

vpmovm2q zmm0, k1               ;convert mask to vector
vpandq zmm1, zmm0, zmm4         ;qword values for add
vpaddq zmm2, zmm1, zmmword ptr [r12] ;update col_counts
vmovdqu64 zmmword ptr [r12], zmm2 ;save updated col_counts

```

```

add r10, 8                    ;col_index += 8
add rcx, 64                   ;x += 8
add r11, 64                   ;col_means += 8
add r12, 64                   ;col_counts += 8
jmp NextColSet

; Update col_means and col_counts using next four columns
@@: sub rax, 4
    cmp rax, r8                ;4 or more columns remaining?
    ja @F                     ;jump if col_index + 4 > ncols

vmovupd ymm0, ymmword ptr [rcx] ;load next 4 cols of cur row
vcmpdpd k1, ymm0, real8 bcst [r13], CMP_GE ;k1 = mask of values >= x_min
vmovupd ymm1{k1}{z}, ymm0       ;values >= x_min or 0.0
vaddpd ymm2, ymm1, ymmword ptr [r11] ;add values to col_means
vmovupd ymmword ptr [r11], ymm2 ;save updated col_means

vpmovm2q ymm0, k1              ;convert mask to vector
vpandq ymm1, ymm0, ymm4        ;qword values for add
vpaddq ymm2, ymm1, ymmword ptr [r12] ;update col_counts
vmovdqu64 ymmword ptr [r12], ymm2 ;save updated col_counts

add r10, 4                    ;col_index += 4
add rcx, 32                   ;x += 4
add r11, 32                   ;col_means += 4
add r12, 32                   ;col_counts += 4
jmp NextColSet

; Update col_means and col_counts using next two columns
@@: sub rax, 2
    cmp rax, r8                ;2 or more columns remaining?
    ja @F                     ;jump if col_index + 2 > ncols

vmovupd xmm0, xmmword ptr [rcx] ;load next 2 cols of cur row
vcmpdpd k1, xmm0, real8 bcst [r13], CMP_GE ;k1 = mask of values >= x_min
vmovupd xmm1{k1}{z}, xmm0       ;values >= x_min or 0.0
vaddpd xmm2, xmm1, xmmword ptr [r11] ;add values to col_means
vmovupd xmmword ptr [r11], xmm2 ;save updated col_means

vpmovm2q xmm0, k1              ;convert mask to vector
vpandq xmm1, xmm0, xmm4        ;qword values for add
vpaddq xmm2, xmm1, xmmword ptr [r12] ;update col_counts
vmovdqu64 xmmword ptr [r12], xmm2 ;save updated col_counts

add r10, 2                    ;col_index += 2
add rcx, 16                   ;x += 2
add r11, 16                   ;col_means += 2
add r12, 16                   ;col_counts += 2
jmp NextColSet

; Update col_means using last column of current row
@@: vmovsd xmm0, real8 ptr [rcx] ;load x from last column
    vcmpps k1, xmm0, real8 ptr [r13], CMP_GE ;k1 = mask of values >= x_min
    vmovsd xmm1{k1}{z}, xmm0       ;value or 0.0
    vaddsd xmm2, xmm1, real8 ptr [r11] ;add to col_means
    vmovsd real8 ptr [r11], xmm2    ;save updated col_means
    kmovb eax, k1                  ;eax = 0 or 1
    add qword ptr [r12], rax        ;update col_counts

add r10, 1                    ;col_index += 1
add rcx, 8                    ;update x ptr

NextColSet:
    cmp r10, r8                ;more columns in current row?
    jb LP2                     ;jump if yes
    dec rdx                     ;nrows -= 1
    jnz LP1                     ;jump if more rows

```

```

; Compute the final col_means
@@: vmovsd xmm0, real8 ptr [r9] ;xmm0 = col_means[i]
    vcvtsi2sd xmm1, xmm1, qword ptr [rbx] ;xmm1 = col_counts[i]
    vdivsd xmm2, xmm0, xmm1 ;compute final mean
    vmovsd real8 ptr [r9], xmm2 ;save col_mean[i]
    add r9, 8 ;update col_means ptr
    add rbx, 8 ;update col_counts ptr
    sub r8, 1 ;ncols -= 1
    jnz @B ;repeat until done

    mov eax, 1 ;set success return code

Done: _DeleteFrame rbx, r12, r13
      vzeroupper
      ret

Avx512CalcColumnMeans_ endp
end
;-----
;##### Ch13_07.asm
;include <MacrosX86-64-AVX.asmh>

; Indices for gather and scatter instructions
ConstVals segment readonly align(64) 'const'
GS_X qword 0, 3, 6, 9, 12, 15, 18, 21
GS_Y qword 1, 4, 7, 10, 13, 16, 19, 22
GS_Z qword 2, 5, 8, 11, 14, 17, 20, 23
ConstVals ends

; extern "C" bool Avx512VcpAos_(Vector* c, const Vector* a, const Vector* b, size_t num_vectors);

.code
Avx512VcpAos_ proc

; Make sure num_vec is valid
xor eax, eax ;set error code (also i = 0)
test r9, r9
jz Done ;jump if num_vec is zero
test r9, 07h
jnz Done ;jump if num_vec % 8 != 0 is true

; Load indices for gather and scatter operations
vmovdqa64 zmm29, zmmword ptr [GS_X] ;zmm29 = X component indices
vmovdqa64 zmm30, zmmword ptr [GS_Y] ;zmm30 = Y component indices
vmovdqa64 zmm31, zmmword ptr [GS_Z] ;zmm31 = Z component indices

; Load next 8 vectors
align 16
@@: kxnorb k1, k1, k1
    vgatherqpd zmm0{k1}, [rdx+zmm29*8] ;zmm0 = A.X values

    kxnorb k2, k2, k2
    vgatherqpd zmm1{k2}, [rdx+zmm30*8] ;zmm1 = A.Y values

    kxnorb k3, k3, k3
    vgatherqpd zmm2{k3}, [rdx+zmm31*8] ;zmm2 = A.Z values

    kxnorb k4, k4, k4
    vgatherqpd zmm3{k4}, [r8+zmm29*8] ;zmm3 = B.X values

    kxnorb k5, k5, k5
    vgatherqpd zmm4{k5}, [r8+zmm30*8] ;zmm4 = B.Y values

    kxnorb k6, k6, k6
    vgatherqpd zmm5{k6}, [r8+zmm31*8] ;zmm5 = B.Z values

```

```

; Calculate 8 vector cross products
vmulpd zmm16, zmm1, zmm5
vmulpd zmm17, zmm2, zmm4
vsubpd zmm18, zmm16, zmm17 ;c.X = a.Y * b.Z - a.Z * b.Y

vmulpd zmm19, zmm2, zmm3
vmulpd zmm20, zmm0, zmm5
vsubpd zmm21, zmm19, zmm20 ;c.Y = a.Z * b.X - a.X * b.Z

vmulpd zmm22, zmm0, zmm4
vmulpd zmm23, zmm1, zmm3
vsubpd zmm24, zmm22, zmm23 ;c.Z = a.X * b.Y - a.Y * b.X

; Save calculated cross products
kxnorb k4, k4, k4
vscatterqpd [rcx+zmm29*8]{k4}, zmm18 ;save C.X components

kxnorb k5, k5, k5
vscatterqpd [rcx+zmm30*8]{k5}, zmm21 ;save C.Y components

kxnorb k6, k6, k6
vscatterqpd [rcx+zmm31*8]{k6}, zmm24 ;save C.Z components

; Update pointers and counters
add rcx, 192 ;c += 8
add rdx, 192 ;a += 8
add r8, 192 ;b += 8
add rax, 8 ;i += 8
cmp rax, r9
jb @B

mov eax, 1 ;set success return code

Done: vzeroupper
      ret
Avx512VcpAos_ endp

; extern "C" bool Avx512VcpSoa_(VectorSoA* c, const VectorSoA* a, const VectorSoA* b, size_t num_vectors);

Avx512VcpSoa_ proc frame
    _CreateFrame CP2_, 0, 0, rbx, rsi, rdi, r12, r13, r14, r15
    _EndProlog

; Make sure num_vec is valid
xor eax, eax
test r9, r9
jz Done ;jump if num_vec is zero
test r9, 07h
jnz Done ;jump if num_vec % 8 != 0 is true

; Load vector array pointers and check for proper alignment
mov r10, [rdx] ;r10 = a.X
or rax, r10
mov r11, [rdx+8] ;r11 = a.Y
or rax, r11
mov r12, [rdx+16] ;r12 = a.Z
or rax, r12

mov r13, [r8] ;r13 = b.X
or rax, r13
mov r14, [r8+8] ;r14 = b.Y
or rax, r14
mov r15, [r8+16] ;r15 = b.Z
or rax, r15

mov rbx, [rcx] ;rbx = c.X

```

```

    or rax,rbx
    mov rsi,[rcx+8]          ;rsi = c.Y
    or rax,rsi
    mov rdi,[rcx+16]         ;rdi = c.Z
    or rax,rdi

    and rax,3fh              ;misaligned component array?
    mov eax,0                ;error return code (also i = 0)
    jnz Done

; Load next block (8 vectors) from a and b
align 16
@@:  vmovapd zmm0,zmmword ptr [r10+rax*8]    ;zmm0 = a.X values
     vmovapd zmm1,zmmword ptr [r11+rax*8]    ;zmm1 = a.Y values
     vmovapd zmm2,zmmword ptr [r12+rax*8]    ;zmm2 = a.Z values
     vmovapd zmm3,zmmword ptr [r13+rax*8]    ;zmm3 = b.X values
     vmovapd zmm4,zmmword ptr [r14+rax*8]    ;zmm4 = b.Y values
     vmovapd zmm5,zmmword ptr [r15+rax*8]    ;zmm5 = b.Z values

; Calculate cross products
    vmulpd zmm16,zmm1,zmm5
    vmulpd zmm17,zmm2,zmm4
    vsubpd zmm18,zmm16,zmm17                ;c.X = a.Y * b.Z - a.Z * b.Y

    vmulpd zmm19,zmm2,zmm3
    vmulpd zmm20,zmm0,zmm5
    vsubpd zmm21,zmm19,zmm20                ;c.Y = a.Z * b.X - a.X * b.Z

    vmulpd zmm22,zmm0,zmm4
    vmulpd zmm23,zmm1,zmm3
    vsubpd zmm24,zmm22,zmm23                ;c.Z = a.X * b.Y - a.Y * b.X

; Save calculated cross products
    vmovapd zmmword ptr [rbx+rax*8],zmm18    ;save C.X values
    vmovapd zmmword ptr [rsi+rax*8],zmm21    ;save C.Y values
    vmovapd zmmword ptr [rdi+rax*8],zmm24    ;save C.Z values

    add rax,8                                ;i += 8
    cmp rax,r9
    jb @B                                     ;repeat until done

Done:  vzeroupper
       _DeleteFrame rbx,rsi,rdi,r12,r13,r14,r15
       ret
Avx512VcpSoa_endp
end

;-----
;##### Ch13_08.asm
;-----
ConstVals  segment readonly align(64) 'const'
; Indices for matrix permutations
MatPerm0   dword 0, 4, 8, 12, 0, 4, 8, 12, 0, 4, 8, 12, 0, 4, 8, 12
MatPerm1   dword 1, 5, 9, 13, 1, 5, 9, 13, 1, 5, 9, 13, 1, 5, 9, 13
MatPerm2   dword 2, 6, 10, 14, 2, 6, 10, 14, 2, 6, 10, 14, 2, 6, 10, 14
MatPerm3   dword 3, 7, 11, 15, 3, 7, 11, 15, 3, 7, 11, 15, 3, 7, 11, 15
; Indices for vector permutations
VecPerm0   dword 0, 0, 0, 0, 4, 4, 4, 4, 8, 8, 8, 8, 12, 12, 12, 12
VecPerm1   dword 1, 1, 1, 1, 5, 5, 5, 5, 9, 9, 9, 9, 13, 13, 13, 13
VecPerm2   dword 2, 2, 2, 2, 6, 6, 6, 6, 10, 10, 10, 10, 14, 14, 14, 14
VecPerm3   dword 3, 3, 3, 3, 7, 7, 7, 7, 11, 11, 11, 11, 15, 15, 15, 15
ConstVals  ends

; extern "C" bool Avx512MatVecMulF32_(Vec4x1_F32* vec_b, float mat[4][4], Vec4x1_F32* vec_a,
; size_t num_vec);

.code
Avx512MatVecMulF32_proc

```

```

xor eax,eax                                ;set error code (also i = 0)
test r9,r9
jz Done                                    ;jump if num_vec is zero
test r9,3
jnz Done                                  ;jump if n % 4 != 0

test rcx,3fh
jnz Done                                  ;jump if vec_b is not properly aligned
test rdx,3fh
jnz Done                                  ;jump if mat is not properly aligned
test r8,3fh
jnz Done                                  ;jump if vec_a is not properly aligned

; Load permutation indices for matrix columns and vector elements
vmovdqa32 zmm16,zmmword ptr [MatPerm0]    ;mat col 0 indices
vmovdqa32 zmm17,zmmword ptr [MatPerm1]    ;mat col 1 indices
vmovdqa32 zmm18,zmmword ptr [MatPerm2]    ;mat col 2 indices
vmovdqa32 zmm19,zmmword ptr [MatPerm3]    ;mat col 3 indices

vmovdqa32 zmm24,zmmword ptr [VecPerm0]    ;W component indices
vmovdqa32 zmm25,zmmword ptr [VecPerm1]    ;X component indices
vmovdqa32 zmm26,zmmword ptr [VecPerm2]    ;Y component indices
vmovdqa32 zmm27,zmmword ptr [VecPerm3]    ;Z component indices

; Load source matrix and duplicate columns
vmovaps zmm0,zmmword ptr [rdx]            ;zmm0 = mat

vpermpps zmm20,zmm16,zmm0                ;zmm20 = mat col 0 (4x)
vpermpps zmm21,zmm17,zmm0                ;zmm21 = mat col 1 (4x)
vpermpps zmm22,zmm18,zmm0                ;zmm22 = mat col 2 (4x)
vpermpps zmm23,zmm19,zmm0                ;zmm23 = mat col 3 (4x)

; Load the next 4 vectors
align 16
@@:  vmovaps zmm4,zmmword ptr [r8+rax]      ;zmm4 = vec_a (4 vectors)
; Permute the vector elements for subsequent calculations
vpermpps zmm0,zmm24,zmm4                ;zmm0 = vec_a W components
vpermpps zmm1,zmm25,zmm4                ;zmm1 = vec_a X components
vpermpps zmm2,zmm26,zmm4                ;zmm2 = vec_a Y components
vpermpps zmm3,zmm27,zmm4                ;zmm3 = vec_a Z components

; Perform matrix-vector multiplications (4 vectors)
vmulps zmm28,zmm20,zmm0
vmulps zmm29,zmm21,zmm1
vmulps zmm30,zmm22,zmm2
vmulps zmm31,zmm23,zmm3
vaddps zmm4,zmm28,zmm29
vaddps zmm5,zmm30,zmm31
vaddps zmm4,zmm4,zmm5                    ;zmm4 = vec_b (4 vectors)
vmovaps zmmword ptr [rcx+rax],zmm4       ;save result

    add rax,64                            ;rax = offset to next block of 4 vectors
    sub r9,4
    jnz @B                                ;repeat until done

    mov eax,1                             ;set success code

Done:  vzeroupper
       ret
Avx512MatVecMulF32_endp
end;

;-----
;##### Ch13_09_.asm
include <MacrosX86-64-AVX.asmh>
extern c_NumPtsMin:dword
extern c_NumPtsMax:dword
extern c_KernelSizeMin:dword
extern c_KernelSizeMax:dword

```



```

; extern bool Avx512Convolve2_(float* y, const float* x, int num_pts, const float* kernel, int
kernel_size)
.code
Avx512Convolve2_proc frame
    _CreateFrame CV2_, 0, 0, rbx
    _EndProlog

; Validate argument values
xor eax, eax                ;set error code

mov r10d, dword ptr [rbp+CV2_OffsetStackArgs]
test r10d, 1
jz Done                    ;kernel_size is even
cmp r10d, [c_KernelSizeMin]
jl Done                    ;kernel_size too small
cmp r10d, [c_KernelSizeMax]
jg Done                    ;kernel_size too big

cmp r8d, [c_NumPtsMin]
jl Done                    ;num_pts too small
cmp r8d, [c_NumPtsMax]
jg Done                    ;num_pts too big
test r8d, 15
jnz Done                    ;num_pts not even multiple of 16

test rcx, 3fh
jnz Done                    ;y is not properly aligned

; Initialize convolution loop variables
shr r10d, 1                ;r10 = kernel_size / 2 (ks2)
lea rdx, [rdx+r10*4]        ;rdx = x + ks2 (first data point)
xor ebx, ebx                ;i = 0

; Perform convolution
LP1: vxorps zmm0, zmm0, zmm0 ;packed sum = 0.0;
mov r11, r10                ;r11 = ks2
neg r11                     ;k = -ks2

LP2: mov rax, rbx            ;rax = i
sub rax, r11                ;rax = i - k
vmovups zmm1, zmmword ptr [rdx+rax*4] ;load x[i - k]:x[i - k + 15]

mov rax, r11
add rax, r10                ;rax = k + ks2
vbroadcastss zmm2, real4 ptr [r9+rax*4] ;zmm2 = kernel[k + ks2]
vmadd231ps zmm0, zmm1, zmm2 ;zmm0 += x[i-k]:x[i-k+15] * kernel[k+ks2]

add r11, 1                  ;k += 1
cmp r11, r10
jle LP2                     ;repeat until k > ks2

vmovaps zmmword ptr [rcx+rbx*4], zmm0 ;save y[i]:y[i + 15]

add rbx, 16                 ;i += 16
cmp rbx, r8
jl LP1                     ;repeat until done
mov eax, 1                  ;set success return code

Done: vzeroupper
_DeleteFrame rbx
ret
Avx512Convolve2_endp

; extern bool Avx512Convolve2Ks5_(float* y, const float* x, int num_pts, const float* kernel, int kernel_size)
Avx512Convolve2Ks5_proc frame
    _CreateFrame CKS5_, 0, 48

```

```

_SaveXmmRegs xmm6, xmm7, xmm8
_EndProlog

; Validate argument values
xor eax, eax                ;set error code (rax is also loop index var)
cmp dword ptr [rbp+CKS5_OffsetStackArgs], 5
jne Done                    ;jump if kernel_size is not 5

cmp r8d, [c_NumPtsMin]
jl Done                    ;jump if num_pts too small
cmp r8d, [c_NumPtsMax]
jg Done                    ;jump if num_pts too big
test r8d, 15
jnz Done                    ;num_pts not even multiple of 15

test rcx, 3fh
jnz Done                    ;y is not properly aligned

; Perform required initializations
vbroadcastss zmm4, real4 ptr [r9] ;kernel[0]
vbroadcastss zmm5, real4 ptr [r9+4] ;kernel[1]
vbroadcastss zmm6, real4 ptr [r9+8] ;kernel[2]
vbroadcastss zmm7, real4 ptr [r9+12] ;kernel[3]
vbroadcastss zmm8, real4 ptr [r9+16] ;kernel[4]

mov r8d, r8d                ;r8 = num_pts
add rdx, 8                  ;x += 2

; Perform convolution
@@: vxorps zmm2, zmm2, zmm2 ;initialize sum vars
vxorps zmm3, zmm3, zmm3

mov r11, rax
add r11, 2                  ;j = i + ks2

vmovups zmm0, zmmword ptr [rdx+r11*4] ;zmm0 = x[j]:x[j + 15]
vmadd231ps zmm2, zmm0, zmm4 ;zmm2 += x[j]:x[j + 15] * kernel[0]

vmovups zmm1, zmmword ptr [rdx+r11*4-4] ;zmm1 = x[j - 1]:x[j + 14]
vmadd231ps zmm3, zmm1, zmm5 ;zmm3 += x[j - 1]:x[j + 14] * kernel[1]

vmovups zmm0, zmmword ptr [rdx+r11*4-8] ;zmm0 = x[j - 2]:x[j + 13]
vmadd231ps zmm2, zmm0, zmm6 ;zmm2 += x[j - 2]:x[j + 13] * kernel[2]

vmovups zmm1, zmmword ptr [rdx+r11*4-12] ;zmm1 = x[j - 3]:x[j + 12]
vmadd231ps zmm3, zmm1, zmm7 ;zmm3 += x[j - 3]:x[j + 12] * kernel[3]

vmovups zmm0, zmmword ptr [rdx+r11*4-16] ;zmm0 = x[j - 4]:x[j + 11]
vmadd231ps zmm2, zmm0, zmm8 ;zmm2 += x[j - 4]:x[j + 11] * kernel[4]

vaddps zmm0, zmm2, zmm3 ;final values
vmovaps zmmword ptr [rcx+rax*4], zmm0 ;save y[i]:y[i + 15]

add rax, 16                 ;i += 16
cmp rax, r8
jl @B                      ;jump if i < num_pts
mov eax, 1                  ;set success return code

Done: vzeroupper
_RestoreXmmRegs xmm6, xmm7, xmm8
_DeleteFrame
ret
Avx512Convolve2Ks5_endp
end

;-----
;##### Ch14_01.asm
;-----

```

```

; extern "C" void Avx512PackedMathI16_(const ZmmVal* a, const ZmmVal* b, ZmmVal c[6])
.code
Avx512PackedMathI16_ proc
    vmovdqu16 zmm0, zmmword ptr [rcx]    ;zmm0 = a
    vmovdqu16 zmm1, zmmword ptr [rdx]    ;zmm1 = b

; Perform packed word operations
    vpaddw zmm2, zmm0, zmm1              ;add
    vmovdqa64 zmmword ptr [r8], zmm2     ;save vpaddw result

    vpaddsw zmm2, zmm0, zmm1             ;add with signed saturation
    vmovdqa64 zmmword ptr [r8+64], zmm2  ;save vpaddsw result

    vpsubw zmm2, zmm0, zmm1             ;sub
    vmovdqa64 zmmword ptr [r8+128], zmm2 ;save vpsubw result

    vpsubsw zmm2, zmm0, zmm1            ;sub with signed saturation
    vmovdqa64 zmmword ptr [r8+192], zmm2 ;save vpsubsw result

    vpminsw zmm2, zmm0, zmm1            ;signed minimums
    vmovdqa64 zmmword ptr [r8+256], zmm2 ;save vpminsw result

    vpmasw zmm2, zmm0, zmm1            ;signed maximums
    vmovdqa64 zmmword ptr [r8+320], zmm2 ;save vpmasw result

    vzeroupper
    ret
Avx512PackedMathI16_ endp

extern "C" void Avx512PackedMathI64_(const ZmmVal* a, const ZmmVal* b, ZmmVal c[5], unsigned int opmask)
Avx512PackedMathI64_ proc
    vmovdqa64 zmm0, zmmword ptr [rcx]    ;zmm0 = a
    vmovdqa64 zmm1, zmmword ptr [rdx]    ;zmm1 = b

    and r9d, 0ffh                        ;r9d = opmask value
    kmovb k1, r9d                        ;k1 = opmask

; Perform packed quadword operations
    vpaddq zmm2{k1}{z}, zmm0, zmm1       ;add
    vmovdqa64 zmmword ptr [r8], zmm2     ;save vpaddq result

    vpsubq zmm2{k1}{z}, zmm0, zmm1       ;sub
    vmovdqa64 zmmword ptr [r8+64], zmm2  ;save vpsubq result

    vpmullq zmm2{k1}{z}, zmm0, zmm1       ;signed mul (low 64 bits)
    vmovdqa64 zmmword ptr [r8+128], zmm2 ;save vpmullq result

    vpsllvq zmm2{k1}{z}, zmm0, zmm1       ;shift left logical
    vmovdqa64 zmmword ptr [r8+192], zmm2 ;save vpsllvq result

    vpsravq zmm2{k1}{z}, zmm0, zmm1       ;shift right arithmetic
    vmovdqa64 zmmword ptr [r8+256], zmm2 ;save vpsravq result

    vpabsq zmm2{k1}{z}, zmm0             ;absolute value
    vmovdqa64 zmmword ptr [r8+320], zmm2 ;save vpabsq result

    vzeroupper
    ret
Avx512PackedMathI64_ endp
end
;-----
;##### Ch14_02.asm
;include <cmpegu.asmh>
;extern c_NumPixelsMax:dword

;-----
;const
r4_1p0    real4 1.0

```

```

r4_255p0    real4 255.0

;extern "C" bool Avx512ConvertImgU8ToF32_(float* des, const uint8_t* src, uint32_t num_pixels)
.code
Avx512ConvertImgU8ToF32_ proc

; Make sure num_pixels is valid and pixel buffers are properly aligned
    xor eax, eax                        ;set error return code
    or r8d, r8d
    jz Done                             ;jump if num_pixels is zero
    cmp r8d, [c_NumPixelsMax]          ;jump if num_pixels too big
    ja Done
    test r8d, 3fh                       ;jump if num_pixels % 64 != 0
    jnz Done
    test rcx, 3fh                       ;jump if des not aligned
    jnz Done
    test rdx, 3fh                       ;jump if src not aligned
    jnz Done

; Perform required initializations
    shr r8d, 6                          ;number of blocks (64 pixels/block)
    vmovss xmm0, real4 ptr [r4_1p0]
    vdivss xmm1, xmm0, real4 ptr [r4_255p0]
    vbroadcastss zmm5, xmm1            ;packed scale factor (1.0 / 255.0)
    align 16

@@:
    vpmovzxbd zmm0, xmmword ptr [rdx]
    vpmovzxbd zmm1, xmmword ptr [rdx+16]
    vpmovzxbd zmm2, xmmword ptr [rdx+32]
    vpmovzxbd zmm3, xmmword ptr [rdx+48] ;zmm3:zmm0 = 64 U32 pixels

; Convert pixels from uint8_t to float [0.0, 255.0]
    vcvtudq2ps zmm16, zmm0
    vcvtudq2ps zmm17, zmm1
    vcvtudq2ps zmm18, zmm2
    vcvtudq2ps zmm19, zmm3            ;zmm19:zmm16 = 64 F32 pixels

; Normalize pixels to [0.0, 1.0]
    vmulps zmm20, zmm16, zmm5
    vmulps zmm21, zmm17, zmm5
    vmulps zmm22, zmm18, zmm5
    vmulps zmm23, zmm19, zmm5        ;zmm23:zmm20 = 64 F32 pixels (normalized)

; Save F32 pixels to des
    vmovaps zmmword ptr [rcx], zmm20
    vmovaps zmmword ptr [rcx+64], zmm21
    vmovaps zmmword ptr [rcx+128], zmm22
    vmovaps zmmword ptr [rcx+192], zmm23

; Update pointers and counters
    add rdx, 64
    add rcx, 256
    sub r8d, 1
    jnz @B

    mov eax, 1                        ;set success return code

Done:    vzeroupper
        ret
Avx512ConvertImgU8ToF32_ endp

; extern "C" bool Avx512ConvertImgF32ToU8_(uint8_t* des, const float* src, uint32_t num_pixels)
Avx512ConvertImgF32ToU8_ proc
; Make sure num_pixels is valid and pixel buffers are properly aligned
    xor eax, eax                        ;set error return code
    or r8d, r8d
    jz Done                             ;jump if num_pixels is zero
    cmp r8d, [c_NumPixelsMax]          ;jump if num_pixels too big
    ja Done

```

```

    test r8d,3fh
    jnz Done                ;jump if num_pixels % 64 != 0
    test rcx,3fh
    jnz Done                ;jump if des not aligned
    test rdx,3fh
    jnz Done                ;jump if src not aligned

; Perform required initializations
    shr r8d,4               ;number of pixel blocks (16 pixels / block)
    vxorps zmm29,zmm29      ;packed 0.0
    vbroadcastss zmm30,[r4_1p0] ;packed 1.0
    vbroadcastss zmm31,[r4_255p0] ;packed 255.0

    align 16
@@:    vmovaps zmm0,zmmword ptr [rdx] ;zmm0 = block of 16 pixels

; Clip pixels in current block to [0.0, 1.0]
    vcmpps k1,zmm0,zmm29,CMP_GE ;k1 = mask of pixels >= 0.0
    vmovaps zmm1{k1}{z},zmm0    ;all pixels >= 0.0

    vcmpps k2,zmm1,zmm30,CMP_GT ;k2 = mask of pixels > 1.0
    vmovaps zmm1{k2},zmm30      ;all pixels clipped to [0.0, 1.0]

; Convert pixels to uint8_t and save to des
    vmulps zmm2,zmm1,zmm31      ;all pixels [0.0, 255.0]
    vcvtps2udq zmm3,zmm2{ru-sae} ;all pixels [0, 255]
    vpmovusdb xmmword ptr [rcx],zmm3 ;save pixels as unsigned bytes

; Update pointers and counters
    add rdx,64
    add rcx,16
    sub r8d,1
    jnz @B

    mov eax,1                  ;set success return code

Done:    vzeroupper
        ret
Avx512ConvertImgF32ToU8_ endp
end

;-----
;##### Ch14_03.asm
include <cmpequ.asmh>
extern c_NumPixelsMax:qword

; Macro CmpPixels
_CmpPixels macro CmpOp
    align 16
@@:    vmovdqa64 zmm0,zmmword ptr [rdx+rax] ;load next block of 64 pixels
        vpcmpub k1,zmm0,zmm4,CmpOp         ;perform compare operation
        vmovdqu8 zmm1{k1}{z},zmm5          ;set mask pixels to 0 or 255 using opmask
        vmovdqa64 zmmword ptr [rcx+rax],zmm1 ;save mask pixels

        add rax,64                          ;update offset
        sub r8,64
        jnz @B                              ;repeat until done
        mov eax,1                          ;set success return code
        vzeroupper
        ret
    endm

; extern "C" bool Avx512ComparePixels_(uint8_t* des, const uint8_t* src,
;   size_t num_pixels, CmpOp cmp_op, uint8_t cmp_val);

.code
Avx512ComparePixels_ proc

```

```

; Make sure num_pixels is valid and pixel buffers are properly aligned
    xor eax,eax              ;set error code (also array offset)
    or r8,r8
    jz Done                 ;jump if num_pixels is zero
    cmp r8,[c_NumPixelsMax]
    ja Done                 ;jump if num_pixels too big
    test r8,3fh
    jnz Done                ;jump if num_pixels % 64 != 0

    test rcx,3fh
    jnz Done                ;jump if des not aligned
    test rdx,3fh
    jnz Done                ;jump if src not aligned

; Perform required initializations
    vpbroadcastb zmm4,byte ptr [rsp+40] ;zmm4 = packed cmp_val
    mov r10d,255
    vpbroadcastb zmm5,r10d      ;zmm5 = packed 255

; Perform specified compare operation
    cmp r9d,0
    jne LB_NE
    _CmpPixels CMP_EQ           ;CmpOp::EQ

LB_NE: cmp r9d,1
       jne LB_LT
       _CmpPixels CMP_NEQ      ;CmpOp::NE

LB_LT: cmp r9d,2
       jne LB_LE
       _CmpPixels CMP_LT      ;CmpOp::LT

LB_LE: cmp r9d,3
       jne LB_GT
       _CmpPixels CMP_LE      ;CmpOp::LE

LB_GT: cmp r9d,4
       jne LB_GE
       _CmpPixels CMP_NLE     ;CmpOp::GT

LB_GE: cmp r9d,5
       jne Done
       _CmpPixels CMP_NLT     ;CmpOp::GE

Done:    vzeroupper
        ret
Avx512ComparePixels_ endp

;-----
;##### Ch14_04.asm
include <cmpequ.asmh>          include <MacrosX86-64-AVX.asmh>
extern c_NumPixelsMax:qword

; This structure must match the structure that's defined in Ch14_04.h
ImageStats struct
PixelBuffer    qword ?
NumPixels      qword ?
PixelValMin    dword ?
PixelValMax    dword ?
NumPixelsInRange qword ?
PixelSum       qword ?
PixelSumOfSquares qword ?
PixelMean      real8 ?
PixelSd        real8 ?
ImageStats     ends

```

```

_UpdateSums macro Disp
    vmovzxbd zmm0, xmmword ptr [rcx+Disp]    ;zmm0 = 16 pixels
    vpcmpud k1, zmm0, zmm31, CMP_GE         ;k1 = mask of pixels >= pixel_val_min
    vpcmpud k2, zmm0, zmm30, CMP_LE         ;k2 = mask of pixels <= pixel_val_max
    kandw k3, k2, k1                        ;k3 = mask of in-range pixels
    vmovdqa32 zmm1{k3}{z}, zmm0             ;zmm1 = in-range pixels
    vpaddq zmm16, zmm16, zmm1               ;update packed pixel_sum
    vpmulld zmm2, zmm1, zmm1               ;update packed pixel_sum_of_squares
    vpaddq zmm17, zmm17, zmm2
    kmovw rax, k3                           ;count number of in-range pixels
    popcnt rax, rax                         ;update num_pixels_in_range
    add r10, rax
endm

; extern "C" bool Avx512CalcImageStats_(ImageStats& im_stats);

.code
Avx512CalcImageStats_ proc frame
    _CreateFrame CIS_, 0, 0, rsi, r12, r13
    _EndProlog

; Make sure num_pixels is valid and pixel_buff is properly aligned
xor eax, eax                               ;set error return code

    mov rsi, rcx                           ;rsi = im_stats ptr
    mov rcx, qword ptr [rsi+ImageStats.PixelBuffer] ;rcx = pixel buffer ptr
    mov rdx, qword ptr [rsi+ImageStats.NumPixels] ;rdx = num_pixels

    test rdx, rdx
    jz Done                                ;jump if num_pixels is zero
    cmp rdx, [c_NumPixelsMax]
    ja Done                                ;jump if num_pixels too big

    test rcx, 3fh
    jnz Done                               ;jump if pixel_buff misaligned

; Perform required initializations
    mov r8d, dword ptr [rsi+ImageStats.PixelValMin]
    mov r9d, dword ptr [rsi+ImageStats.PixelValMax]

    vpbroadcastd zmm31, r8d                ;packed pixel_val_min
    vpbroadcastd zmm30, r9d                ;packed pixel_val_max

    vpxorq zmm29, zmm29, zmm29             ;packed pixel_sum
    vpxorq zmm28, zmm28, zmm28             ;packed pixel_sum_of_squares
    xor r10d, r10d                         ;num_pixels_in_range = 0

; Compute packed versions of pixel_sum and pixel_sum_of_squares
    cmp rdx, 64
    jb LB1                                  ;jump if there are fewer than 64 pixels

@@:
    align 16
    vpxord zmm16, zmm16, zmm16             ;loop packed pixel_sum = 0
    vpxord zmm17, zmm17, zmm17             ;loop packed pixel_sum_of_squares = 0

    _UpdateSums 0                          ;process pixel_buff[i+15]:pixel_buff[i]
    _UpdateSums 16                         ;process pixel_buff[i+31]:pixel_buff[i+16]
    _UpdateSums 32                         ;process pixel_buff[i+47]:pixel_buff[i+32]
    _UpdateSums 48                         ;process pixel_buff[i+63]:pixel_buff[i+48]

    vextracti32x8 ymm0, zmm16, 1           ;extract top 8 pixel_sum (dwords)
    vpaddq ymm1, ymm0, ymm16
    vmovzxdq zmm2, ymm1
    vpaddq zmm29, zmm29, zmm2              ;update packed pixel_sum (qwords)
    vextracti32x8 ymm0, zmm17, 1           ;extract top 8 pixel_sum_of_squares (dwords)
    vpaddq ymm1, ymm0, ymm17
    vmovzxdq zmm2, ymm1

    vpaddq zmm28, zmm28, zmm2              ;update packed pixel_sum_of_squares (qwords)
    vextracti32x8 ymm0, zmm28, 1           ;extract top 8 pixel_sum_of_squares (dwords)
    vpaddq ymm1, ymm0, ymm28
    vmovzxdq zmm2, ymm1

    vpxorq zmm29, zmm29, zmm29             ;packed pixel_sum
    vpxorq zmm28, zmm28, zmm28             ;packed pixel_sum_of_squares
    xor r10d, r10d                         ;num_pixels_in_range = 0

; Save num_pixel_in_range
LB3: mov qword ptr [rsi+ImageStats.NumPixelsInRange], r10

; Reduce packed pixel_sum to single qword
vextracti64x4 ymm0, zmm29, 1
vpaddq ymm1, ymm0, ymm29
vextracti64x2 xmm2, ymm1, 1
vpaddq xmm3, xmm2, xmm1
vpextrq rax, xmm3, 0
vpextrq r11, xmm3, 1
add rax, r11                               ;rax = sum of qwords in zmm29
add r13, rax                               ;add scalar pixel_sum

    mov qword ptr [rsi+ImageStats.PixelSum], r13

; Reduce packed pixel_sum_of_squares to single qword
vextracti64x4 ymm0, zmm28, 1
vpaddq ymm1, ymm0, ymm28
vextracti64x2 xmm2, ymm1, 1
vpaddq xmm3, xmm2, xmm1
vpextrq rax, xmm3, 0
vpextrq r11, xmm3, 1
add rax, r11                               ;rax = sum of qwords in zmm28
add r12, rax                               ;add scalar pixel_sum_of_squares

    mov qword ptr [rsi+ImageStats.PixelSumOfSquares], r12

; Calculate final mean and sd
vcvtusi2sd xmm0, xmm0, r10                ;num_pixels_in_range (DPFP)
sub r10, 1
vcvtusi2sd xmm1, xmm1, r10                ;num_pixels_in_range - 1 (DPFP)
vcvtusi2sd xmm2, xmm2, r13                ;pixel_sum (DPFP)
vcvtusi2sd xmm3, xmm3, r12                ;pixel_sum_of_squares (DPFP)
ddivsd xmm4, xmm2, xmm0                    ;final pixel_mean

    vmovsd real8 ptr [rsi+ImageStats.PixelMean], xmm4

    vmulsd xmm4, xmm0, xmm3                ;num_pixels_in_range * pixel_sum_of_squares

```

```

vmulsd xmm5, xmm2, xmm2      :pixel_sum * pixel_sum
vsubsd xmm2, xmm4, xmm5      :var_num
vmulsd xmm3, xmm0, xmm1      :var_den
vdivsd xmm4, xmm2, xmm3      :calc variance
vsqrtsd xmm0, xmm0, xmm4      :final pixel_sd

movsd real8 ptr [rsi+ImageStats.PixelSd], xmm0

mov eax, 1                    :set success return code

Done:  vzeroupper
       _DeleteFrame rsi, r12, r13
       ret
Avx512CalcImageStats_ endp
end
-----
:##### Ch14_05.asm
include <MacrosX86-64-AVX.asmh>
extern c_NumPixelsMin:dword
extern c_NumPixelsMax:dword

.const
r4_0p5    real4 0.5
r4_255p0  real4 255.0

; extern "C" bool Avx512RgbToGs_(uint8_t* pb_gs, const uint8_t* const* pb_rgb, int num_pixels,
const float coef[3]);

.code
Avx512RgbToGs_ proc frame
    _CreateFrame RGBGS0_, 0, 96, r13, r14, r15
    _SaveXmmRegs xmm10, xmm11, xmm12, xmm13, xmm14, xmm15
    _EndProlog

    xor eax, eax                :error return code (also pixel_buffer offset)
    cmp r8d, [c_NumPixelsMin]
    jl Done                    :jump if num_pixels < min value
    cmp r8d, [c_NumPixelsMax]
    jg Done                    :jump if num_pixels > max value
    test r8d, 3fh
    jnz Done                    :jump if (num_pixels % 64) != 0

    test rcx, 3fh
    jnz Done                    :jump if pb_gs is not aligned

    mov r13, [rdx]
    test r13, 3fh
    jnz Done                    :jump if pb_r is not aligned
    mov r14, [rdx+8]
    test r14, 3fh
    jnz Done                    :jump if pb_g is not aligned
    mov r15, [rdx+16]
    test r15, 3fh
    jnz Done                    :jump if pb_b is not aligned

; Perform required initializations
vbroadcastss zmm10, real4 ptr [r9] :zmm10 = packed coef[0]
vbroadcastss zmm11, real4 ptr [r9+4] :zmm11 = packed coef[1]
vbroadcastss zmm12, real4 ptr [r9+8] :zmm12 = packed coef[2]
vbroadcastss zmm13, real4 ptr [r4_0p5] :zmm13 = packed 0.5
vbroadcastss zmm14, real4 ptr [r4_255p0] :zmm14 = packed 255.0
vxorps zmm15, zmm15, zmm15 :zmm15 = packed 0.0
mov r8d, r8d :r8 = num_pixels
mov r10, 16 :r10 = number of pixels / iteration

; Load next block of pixels
align 16

```

```

@@:  vpmovzxbd zmm0, xmmword ptr [r13+rax] :zmm0 = 16 pixels (r values)
     vpmovzxbd zmm1, xmmword ptr [r14+rax] :zmm1 = 16 pixels (g values)
     vpmovzxbd zmm2, xmmword ptr [r15+rax] :zmm2 = 16 pixels (b values)
; Convert dword values to SPFP and multiply by coefficients
vcvttdq2ps zmm0, zmm0 :zmm0 = 16 pixels SPFP (r values)
vcvttdq2ps zmm1, zmm1 :zmm1 = 16 pixels SPFP (g values)
vcvttdq2ps zmm2, zmm2 :zmm2 = 16 pixels SPFP (b values)
vmulps zmm0, zmm0, zmm10 :zmm0 = r values * coef[0]
vmulps zmm1, zmm1, zmm11 :zmm1 = g values * coef[1]
vmulps zmm2, zmm2, zmm12 :zmm2 = b values * coef[2]

; Sum color components & clip values to [0.0, 255.0]
vaddps zmm3, zmm0, zmm1 :r + g
vaddps zmm4, zmm3, zmm2 :r + g + b
vaddps zmm5, zmm4, zmm13 :r + g + b + 0.5
vminps zmm0, zmm5, zmm14 :clip pixels above 255.0
vmaxps zmm1, zmm0, zmm15 :clip pixels below 0.0

; Convert grayscale values from SPFP to byte, save results
vcvt2ps2dq zmm2, zmm1 :convert SPFP values to dwords

vpmovusdb xmm3, zmm2 :convert to bytes
vmovdqa xmmword ptr [rcx+rax], xmm3 :save grayscale image pixels

add rax, r10
sub r8, r10
jnz @B

mov eax, 1                    :set success return code
Done:  vzeroupper
       _RestoreXmmRegs xmm10, xmm11, xmm12, xmm13, xmm14, xmm15
       _DeleteFrame r13, r14, r15
       ret
Avx512RgbToGs_ endp

; extern "C" bool Avx2RgbToGs_(uint8_t* pb_gs, const uint8_t* const* pb_rgb, int num_pixels,
const float coef[3]);

.code
Avx2RgbToGs_ proc frame
    _CreateFrame RGBGS1_, 0, 96, r13, r14, r15
    _SaveXmmRegs xmm10, xmm11, xmm12, xmm13, xmm14, xmm15
    _EndProlog

    xor eax, eax                :error return code (also pixel_buffer offset)
    cmp r8d, [c_NumPixelsMin]
    jl Done                    :jump if num_pixels < min value
    cmp r8d, [c_NumPixelsMax]
    jg Done                    :jump if num_pixels > max value
    test r8d, 3fh
    jnz Done                    :jump if (num_pixels % 64) != 0

    test rcx, 3fh
    jnz Done                    :jump if pb_gs is not aligned

    mov r13, [rdx]
    test r13, 3fh
    jnz Done                    :jump if pb_r is not aligned
    mov r14, [rdx+8]
    test r14, 3fh
    jnz Done                    :jump if pb_g is not aligned
    mov r15, [rdx+16]
    test r15, 3fh
    jnz Done                    :jump if pb_b is not aligned

; Perform required initializations
vbroadcastss ymm10, real4 ptr [r9] :ymm10 = packed coef[0]

```

```

vbroadcastss ymm11,real4 ptr [r9+4] ;ymm11 = packed coef[1]
vbroadcastss ymm12,real4 ptr [r9+8] ;ymm12 = packed coef[2]
vbroadcastss ymm13,real4 ptr [r4_0p5] ;ymm13 = packed 0.5
vbroadcastss ymm14,real4 ptr [r4_255p0] ;ymm14 = packed 255.0
vxorps ymm15,ymm15,ymm15 ;ymm15 = packed 0.0
mov r8d,r8d ;r8 = num_pixels
mov r10,8 ;r10 = number of pixels / iteration

; Load next block of pixels
align 16
@@: vpmovzxbd ymm0,qword ptr [r13+rax] ;ymm0 = 8 pixels (r values)
vpmovzxbd ymm1,qword ptr [r14+rax] ;ymm1 = 8 pixels (g values)
vpmovzxbd ymm2,qword ptr [r15+rax] ;ymm2 = 8 pixels (b values)
; Convert dword values to SPFP and multiply by coefficients
vcvttdq2ps ymm0,ymm0 ;ymm0 = 8 pixels SPFP (r values)
vcvttdq2ps ymm1,ymm1 ;ymm1 = 8 pixels SPFP (g values)
vcvttdq2ps ymm2,ymm2 ;ymm2 = 8 pixels SPFP (b values)
vmulps ymm0,ymm0,ymm10 ;ymm0 = r values * coef[0]
vmulps ymm1,ymm1,ymm11 ;ymm1 = g values * coef[1]
vmulps ymm2,ymm2,ymm12 ;ymm2 = b values * coef[2]

; Sum color components & clip values to [0.0, 255.0]
vaddps ymm3,ymm0,ymm1 ;r + g
vaddps ymm4,ymm3,ymm2 ;r + g + b
vaddps ymm5,ymm4,ymm13 ;r + g + b + 0.5
vminps ymm0,ymm5,ymm14 ;clip pixels above 255.0
vmaxps ymm1,ymm0,ymm15 ;clip pixels below 0.0

; Convert grayscale components from SPFP to byte, save results
vcvtts2dq ymm2,ymm1 ;convert SPFP values to dwords

vpackusdw ymm3,ymm2,ymm2
vextracti128 xmm4,ymm3,1
vpackuswb xmm5,xmm3,xmm4 ;byte GS pixels in xmm5[31:0] and xmm5[95:64]
vpextrd r11d,xmm5,0 ;r11d = 4 grayscale pixels
mov dword ptr [rcx+rax],r11d ;save grayscale image pixels
vpextrd r11d,xmm5,2 ;r11d = 4 grayscale pixels
mov dword ptr [rcx+rax+4],r11d ;save grayscale image pixels

add rax,r10
sub r8,r10
jnz @B

mov eax,1 ;set success return code
Done: vzeroupper
_RestoreXmmRegs xmm10,xmm11,xmm12,xmm13,xmm14,xmm15
_DeleteFrame r13,r14,r15
ret
Avx2RgbToGs_ endp
end

;-----
;##### Ch15_01.asm
;-----

.const
r8_2p0 real8 2.0

; extern "C" int CalcResult_(double* y, const double* x, size_t n);

.code
CalcResult_ proc

; Forward conditional jumps are used in this code block since
; the fall-through cases are more likely to occur
test r8,r8
jz Done ;jump if n == 0

test r8,7h

```

```

jnz Error ;jump if (n % 8) != 0
test rcx,1fh
jnz Error ;jump if y is not aligned to a 32b boundary
test rdx,1fh
jnz Error ;jump if x is not aligned to a 32b boundary

; Initialize
xor eax,eax ;set array offset to zero
vbroadcastsd ymm5,real8 ptr [r8_2p0] ;packed 2.0

; Simple array processing loop
align 16
@@: vmovapd ymm0,ymmword ptr [rdx+rax] ;load x[i+3]:x[i]
vdivpd ymm1,ymm0,ymm5
vsqrtpd ymm2,ymm1
vmovapd ymmword ptr [rcx+rax],ymm2 ;save y[i+3]:y[i]

vmovapd ymm0,ymmword ptr [rdx+rax+32] ;load x[i+7]:x[i+4]
vdivpd ymm1,ymm0,ymm5
vsqrtpd ymm2,ymm1
vmovapd ymmword ptr [rcx+rax+32],ymm2 ;save y[i+7]:y[i+4]

; A backward conditional jump is used in this code block since
; the fall-through case is less likely to occur
add rax,64
sub r8,8
jnz @B

Done: xor eax,eax ;set success return code
vzeroupper
ret

; Error handling code that's unlikely to execute
Error: mov eax,1 ;set error return code
ret
CalcResult_ endp
end

;-----
;##### Ch16_02.asm
;-----

; _CalcResult Macro
; The following macro contains a simple calculating loop that is used
; to compare performance of the vmovaps and vmovntps instructions.

_CalcResult macro MovInstr

; Load and validate arguments
xor eax,eax ;set error code
test r9,r9
jz Done ;jump if n <= 0
test r9,0fh
jnz Done ;jump if (n % 16) != 0

test rcx,1fh
jnz Done ;jump if c is not aligned
test rdx,1fh
jnz Done ;jump if a is not aligned
test r8,1fh
jnz Done ;jump if b is not aligned

; Calculate c[i] = sqrt(a[i] * a[i] + b[i] * b[i])
align 16
@@: vmovaps ymm0,ymmword ptr [rdx+rax] ;ymm0 = a[i+7]:a[i]
vmovaps ymm1,ymmword ptr [r8+rax] ;ymm1 = b[i+7]:b[i]
vmulps ymm2,ymm0,ymm0 ;ymm2 = a[i] * a[i]
vmulps ymm3,ymm1,ymm1 ;ymm3 = b[i] * b[i]

```

```

vaddps ymm4, ymm2, ymm3          ;ymm4 = sum
vsqrtps ymm5, ymm4              ;ymm5 = final result
MovInstr ymmword ptr [rcx+rax], ymm5 ;save final values to c

vmovaps ymm0, ymmword ptr [rdx+rax+32] ;ymm0 = a[i+15]:a[i+8]
vmovaps ymm1, ymmword ptr [r8+rax+32] ;ymm1 = b[i+15]:b[i+8]
vmulps ymm2, ymm0, ymm0          ;ymm2 = a[i] * a[i]
vmulps ymm3, ymm1, ymm1          ;ymm3 = b[i] * b[i]
vaddps ymm4, ymm2, ymm3          ;ymm4 = sum
vsqrtps ymm5, ymm4              ;ymm5 = final result
MovInstr ymmword ptr [rcx+rax+32], ymm5 ;save final values to c

add rax, 64                      ;update offset
sub r9, 16                      ;update counter
jnz @B

mov eax, 1                      ;set success return code

Done: vzeroupper
ret
endm

; extern bool CalcResultA_(float* c, const float* a, const float* b, size_t n)
.code
CalcResultA_proc
    _CalcResult vmovaps
CalcResultA_endp

; extern bool CalcResultB_(float* c, const float* a, const float* b, int n)
CalcResultB_proc
    _CalcResult vmovntps
CalcResultB_endp
end
;-----
;##### Ch16_03_.asm
include <Ch16_03_.asmh>

; Macro _LITraverse
; The following macro generates linked list traversal code using the
; prefetchnta instruction if UsePrefetch is equal to 'Y'.

_LITraverse macro UsePrefetch
    mov rax, rcx                ;rax = ptr to 1st node
    test rax, rax
    jz Done                    ;jump if empty list

    align 16
@@:: mov rcx, [rax+LInode.Link] ;rcx = next node
    vmovapd ymm0, ymmword ptr [rax+LInode.ValA] ;ymm0 = ValA
    vmovapd ymm1, ymmword ptr [rax+LInode.ValB] ;ymm1 = ValB

IFIDNI <UsePrefetch>, <Y>
    mov rdx, rcx
    test rdx, rdx              ;is there another node?
    cmovz rdx, rax             ;avoid prefetch of nullptr
    prefetchnta [rdx]          ;prefetch start of next node
ENDIF

; Calculate ValC[i] = sqrt(ValA[i] * ValA[i] + ValB[i] * ValB[i])
vmulpd ymm2, ymm0, ymm0        ;ymm2 = ValA * ValA
vmulpd ymm3, ymm1, ymm1        ;ymm3 = ValB * ValB
vaddpd ymm4, ymm2, ymm3        ;ymm4 = sums
vsqrtpd ymm5, ymm4             ;ymm5 = square roots

vmovntpd ymmword ptr [rax+LInode.ValC], ymm5 ;save result

; Calculate ValD[i] = sqrt(ValA[i] / ValB[i] + ValB[i] / ValA[i]);

```

```

vdivpd ymm2, ymm0, ymm1        ;ymm2 = ValA / ValB
vdivpd ymm3, ymm1, ymm0        ;ymm3 = ValB / ValA
vaddpd ymm4, ymm2, ymm3        ;ymm4 = sums
vsqrtpd ymm5, ymm4             ;ymm5 = square roots

vmovntpd ymmword ptr [rax+LInode.ValD], ymm5 ;save result

mov rax, rcx                    ;rax = ptr to next node
test rax, rax
jnz @B

Done: vzeroupper
ret
endm

; extern "C" void LITraverseA_(LInode* first);

.code
LITraverseA_proc
    _LITraverse n
LITraverseA_endp

; extern "C" void LITraverseB_(LInode* first);

LITraverseB_proc
    _LITraverse y
LITraverseB_endp
end;-----
;##### Ch16_04_.asm
include <MacrosX86-64-AVX.asmh>

CalcInfo struct
    X1 qword ?
    X2 qword ?
    Y1 qword ?
    Y2 qword ?
    Z1 qword ?
    Z2 qword ?
    Result qword ?
    Index0 qword ?
    Index1 qword ?
    Status dword ?
CalcInfo ends

r8_1p0 .const
    real8 1.0

; extern "C" void CalcResult_(CalcInfo* ci)
.code
CalcResult_proc frame
    _CreateFrame CR, 0, 16, r12, r13, r14, r15
    _SaveXmmRegs xmm6
    _EndProlog

    mov dword ptr [rcx+CalcInfo.Status], 0

; Make sure num_elements is valid
mov rax, [rcx+CalcInfo.Index0] ;rax = start index
mov rdx, [rcx+CalcInfo.Index1] ;rdx = stop index
sub rdx, rax
add rdx, 1                      ;rdx = num_elements
test rdx, rdx
jz Done                        ;jump if num_elements == 0
test rdx, 7
jnz Done                      ;jump if num_elements % 8 != 0

; Make sure all arrays are properly aligned

```

```

mov r8d,1fh
mov r9,[rcx+CalcInfo.Result]
test r9,r8
jnz Done

```

```

mov r10,[rcx+CalcInfo.X1]
test r10,r8
jnz Done
mov r11,[rcx+CalcInfo.X2]
test r11,r8
jnz Done

```

```

mov r12,[rcx+CalcInfo.Y1]
test r12,r8
jnz Done
mov r13,[rcx+CalcInfo.Y2]
test r13,r8
jnz Done

```

```

mov r14,[rcx+CalcInfo.Z1]
test r14,r8
jnz Done
mov r15,[rcx+CalcInfo.Z2]
test r15,r8
jnz Done

```

```

; vbroadcastsd ymm6,real8 ptr [r8_1p0] ;ymm6 = packed 1.0 (DPFP)
; Perform simulated calculation

```

```

align 16
LP1: vmovapd ymm0,ymmword ptr [r10+rax*8]
      vmovapd ymm1,ymmword ptr [r12+rax*8]
      vmovapd ymm2,ymmword ptr [r14+rax*8]
      vsubpd ymm0,ymm0,ymmword ptr [r11+rax*8]
      vsubpd ymm1,ymm1,ymmword ptr [r13+rax*8]
      vsubpd ymm2,ymm2,ymmword ptr [r15+rax*8]
      vmulpd ymm3,ymm0,ymm0
      vmulpd ymm4,ymm1,ymm1
      vmulpd ymm5,ymm2,ymm2
      vaddpd ymm0,ymm3,ymm4
      vaddpd ymm1,ymm0,ymm5
      vsqrtpd ymm2,ymm1
      vdivpd ymm3,ymm6,ymm2
      vsqrtpd ymm4,ymm3
      vmovntpd ymmword ptr [r9+rax*8],ymm4

```

```

add rax,4

```

```

vmovapd ymm0,ymmword ptr [r10+rax*8]
vmovapd ymm1,ymmword ptr [r12+rax*8]
vmovapd ymm2,ymmword ptr [r14+rax*8]
vsubpd ymm0,ymm0,ymmword ptr [r11+rax*8]
vsubpd ymm1,ymm1,ymmword ptr [r13+rax*8]
vsubpd ymm2,ymm2,ymmword ptr [r15+rax*8]
vmulpd ymm3,ymm0,ymm0
vmulpd ymm4,ymm1,ymm1
vmulpd ymm5,ymm2,ymm2
vaddpd ymm0,ymm3,ymm4
vaddpd ymm1,ymm0,ymm5
vsqrtpd ymm2,ymm1
vdivpd ymm3,ymm6,ymm2
vsqrtpd ymm4,ymm3
vmovntpd ymmword ptr [r9+rax*8],ymm4

```

```

add rax,4
sub rdx,8
jnz LP1

```

```

mov dword ptr [rcx+CalcInfo.Status],1

```

```

Done: vzeroupper
      _RestoreXmmRegs xmm6
      _DeleteFrame r12,r13,r14,r15
      ret
CalcResult_ endp
end

```